

# Deep Learning with Nontrivial Constraints

**Ju Sun**

McKnight Land-Grant Professor  
Computer Science & Engineering

Dec 10, 2025

**Industrial and Systems Engineering Graduate Seminar**



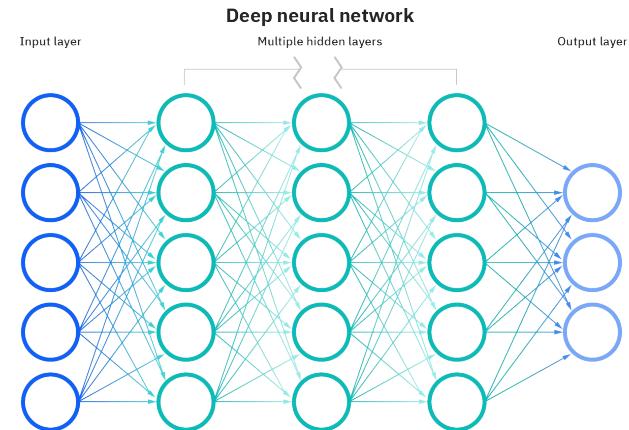
UNIVERSITY OF MINNESOTA  
*Driven to Discover*<sup>SM</sup>

Research in the group



<https://glovex.umn.edu/>

(Machine) **L**earning, (Numerical) **O**ptimization, (Computer) **V**sion, healthcar**E**, + **X**

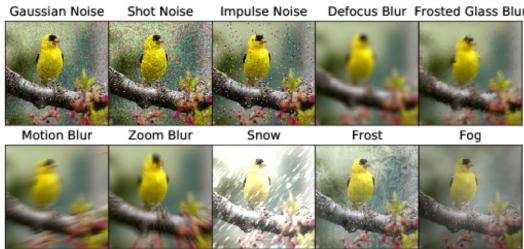
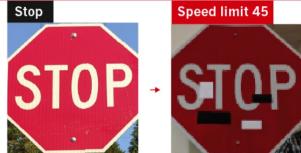


# Our research themes

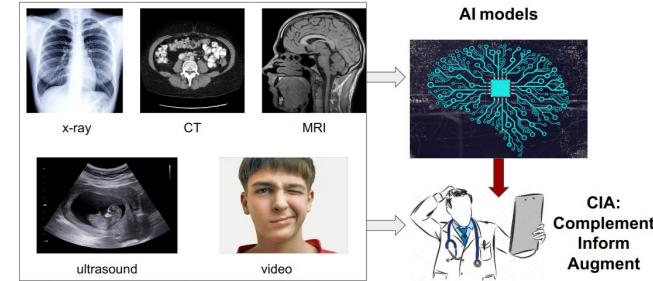
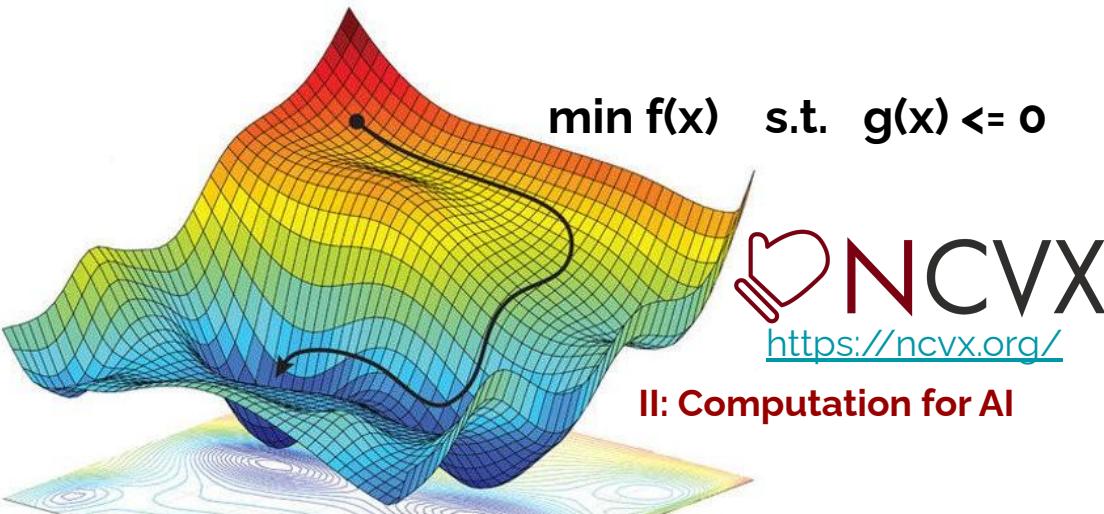
## FOOLING THE AI

Deep neural networks (DNNs) are brilliant at image recognition — but they can be easily hacked.

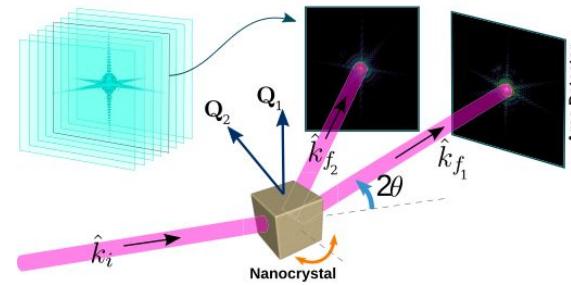
These stickers made an artificial-intelligence system read this stop sign as 'speed limit 45'.



## I: Trustworthy AI



## III: AI for Healthcare



## IV: AI for Science and Engineering

# Robustness & Deployability of Deep Learning Models

First general-purpose, reliable solver for robust evaluation

$$\begin{aligned} \max_{\mathbf{x}'} \ell(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}')) \\ \text{s. t. } d(\mathbf{x}, \mathbf{x}') \leq \varepsilon, \quad \mathbf{x}' \in [0, 1]^n \end{aligned}$$

$$\begin{aligned} \min_{\mathbf{x}'} d(\mathbf{x}, \mathbf{x}') \\ \text{s. t. } \max_{i \neq y} f_{\boldsymbol{\theta}}^i(\mathbf{x}') \geq f_{\boldsymbol{\theta}}^y(\mathbf{x}'), \quad \mathbf{x}' \in [0, 1]^n \end{aligned}$$

## Reliability

- SOTA methods  ROBUSTBENCH  
A standardized benchmark for adversarial robustness  
No stopping criterion (only use maxit); step size scheduler
- PWCF (ours)  
Principled line-search criterion and termination condition

## Generality

- SOTA methods  
Can mostly only handle several lp metrics ( $\ell_1, \ell_2, \ell_\infty$ )
- PWCF (ours)  
Any differentiable metrics and both min and max forms  
E.g., perceptual distance  $d(\mathbf{x}, \mathbf{x}') \doteq \|\phi(\mathbf{x}) - \phi(\mathbf{x}')\|_2$   
where  $\phi(\mathbf{x}) \doteq [\hat{g}_1(\mathbf{x}), \dots, \hat{g}_L(\mathbf{x})]$

## Selective predictions for safety and deployability

predictor  $f : \mathcal{X} \rightarrow \mathbb{R}^K$

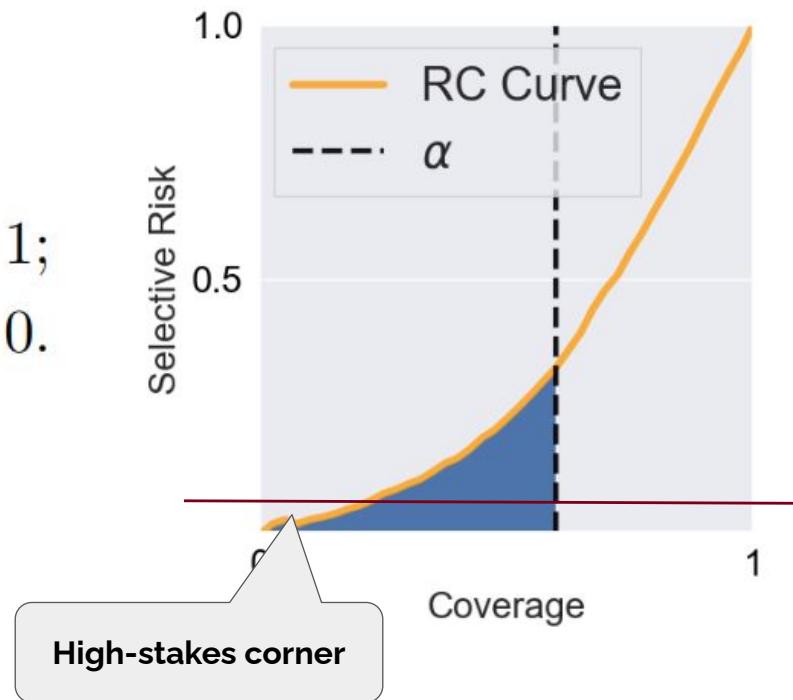
selector  $g : \mathcal{X} \rightarrow \{0, 1\}$

$$(f, g)(\mathbf{x}) \triangleq \begin{cases} f(\mathbf{x}) & \text{if } g(\mathbf{x}) = 1; \\ \text{abstain} & \text{if } g(\mathbf{x}) = 0. \end{cases}$$

No prediction on uncertain samples and defer them to humans

$$g_\gamma(\mathbf{x}) = \mathbb{1}[s(\mathbf{x}) > \gamma]$$

Typically, selection by thresholding prediction confidence



# Computation for Deep Learning with Nontrivial Constraints



NCVX PyGRANSO  
Documentation

Search the docs ...

Introduction

Installation

Sett

Exa

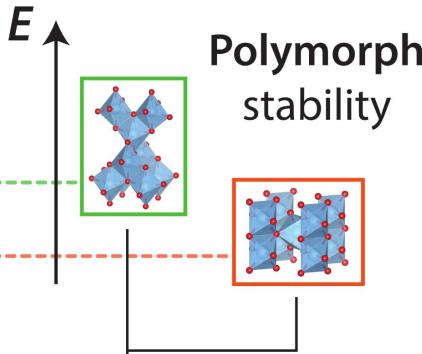
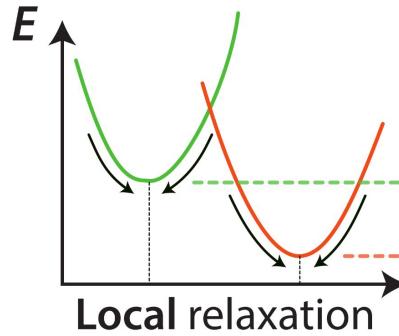
$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}), \text{ s.t. } c_i(\mathbf{x}) \leq 0, \forall i \in \mathcal{I}; \quad c_i(\mathbf{x}) = 0, \forall i \in \mathcal{E}$$

Home

<https://ncvx.org/>



- First general-purpose solver for hard-constrained deep learning problems
- Recently updated to be compatible with PyTorch 2.8



Stability  
respects  
phase sep

- **(W-CSTR-T) Weak explicit constraints informed by phase transition:** For each composition, the energy/atom of any non-ground-state polymorph is greater than that of the ground-state polymorph
- **(W-CSTR-S) Weak explicit constraints informed by phase separation:** For each chemical space, the energy/atom of any TUS material is above the lower convex envelope in the composition-energy space

Settings  
Examples

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{y}_i, f_{\theta}(\mathbf{x}_i)) + \Omega(\theta) \quad \text{s. t.} \quad f_{\theta}(\mathbf{x}_p) \geq f_{\theta}(\mathbf{x}_q) \quad \forall (p, q) \in \mathcal{O}_{\text{W-CSTR-T}}$$

## NCVX: A General-Purpose Optimization Solver for Constrained Machine and Deep Learning

Buyun Liang, Tim Mitchell, Ju Sun

## Enabling

### Robustness evaluation

$$\max_{\mathbf{x}'} \ell(\mathbf{y}, f_{\theta}(\mathbf{x}'))$$

$$\text{s. t. } \mathbf{x}' \in \Delta(\mathbf{x}) = \{\mathbf{x}' \in [0, 1]^n : d(\mathbf{x}, \mathbf{x}') \leq \varepsilon\}$$

$$\min_{\mathbf{x}' \in [0, 1]^n} d(\mathbf{x}, \mathbf{x}') \quad \text{s. t.} \quad \max_{\ell \neq y} f_{\theta}^{\ell}(\mathbf{x}') \geq f_{\theta}^y(\mathbf{x}')$$

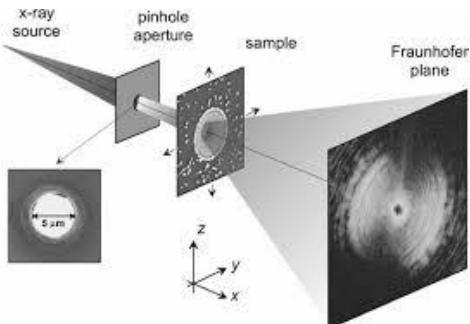
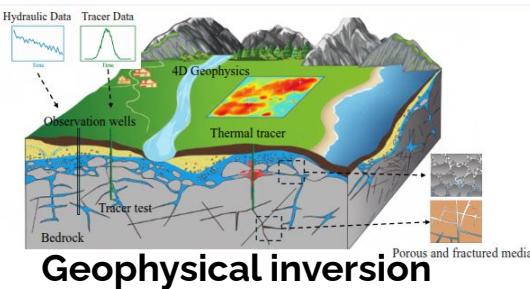
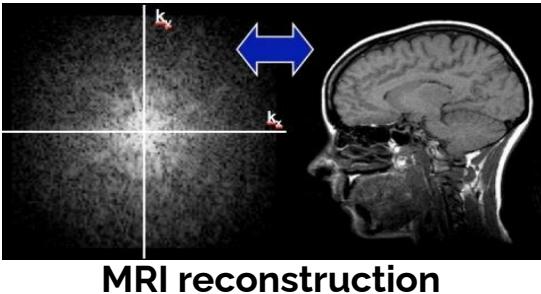
$$\angle_{i=1}^N \mathbb{1}\{f_{\theta}(\mathbf{x}_i) > t\}$$

$$\text{s. t.} \quad \frac{\sum_{i=1}^N \mathbb{1}\{y_i = +1\} \mathbb{1}\{f_{\theta}(\mathbf{x}_i) > t\}}{\sum_{i=1}^N \mathbb{1}\{y_i = +1\}} \geq \alpha$$

Yash, Le, Zhong, Buyun

# Generative Models for Scientific/Engineering Inverse Modeling

(inverse problems, reconstruction, data assimilation, inverse design/control, conditional generation, estimation ...)



Given  $\mathbf{y} \approx f(\mathbf{x})$ , recover  $\mathbf{x}$

$$\min_{\mathbf{x}} \ell(\mathbf{y}, f(\mathbf{x})) + \Omega(\mathbf{x})$$

## Challenges:

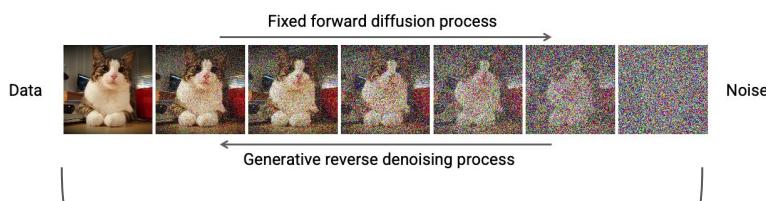
- Linear vs. **nonlinear**  $f$
- Unconstrained vs. **constrained**  
(e.g.,  $PDE(x, y) = 0$ )
- Explicit vs. **implicit**  $f$

A **plug-in principle** for leveraging **pretrained** deep generative models to solve IM

object-only datasets  $\{\mathbf{x}_i\}_{i=1,\dots,N}$



**Distribution learning** via  
deep generative models



$\mathcal{G}_\theta$   
**Plug in pretrained** deep  
generative priors

$$\min_z \mathcal{L}(z) \doteq \ell(\mathbf{y}, \mathcal{A} \circ \mathcal{G}_\theta(z)) + \Omega \circ \mathcal{G}_\theta(z)$$

**DMPlug: A Plug-in Method for Solving Inverse Problems with Diffusion Models**

Hengkang Wang, Xu Zhang, Taihui Li, Yuxiang Wan, Tiancong Chen, Ju Sun

NeurIPS'24, with **domain-specific** priors

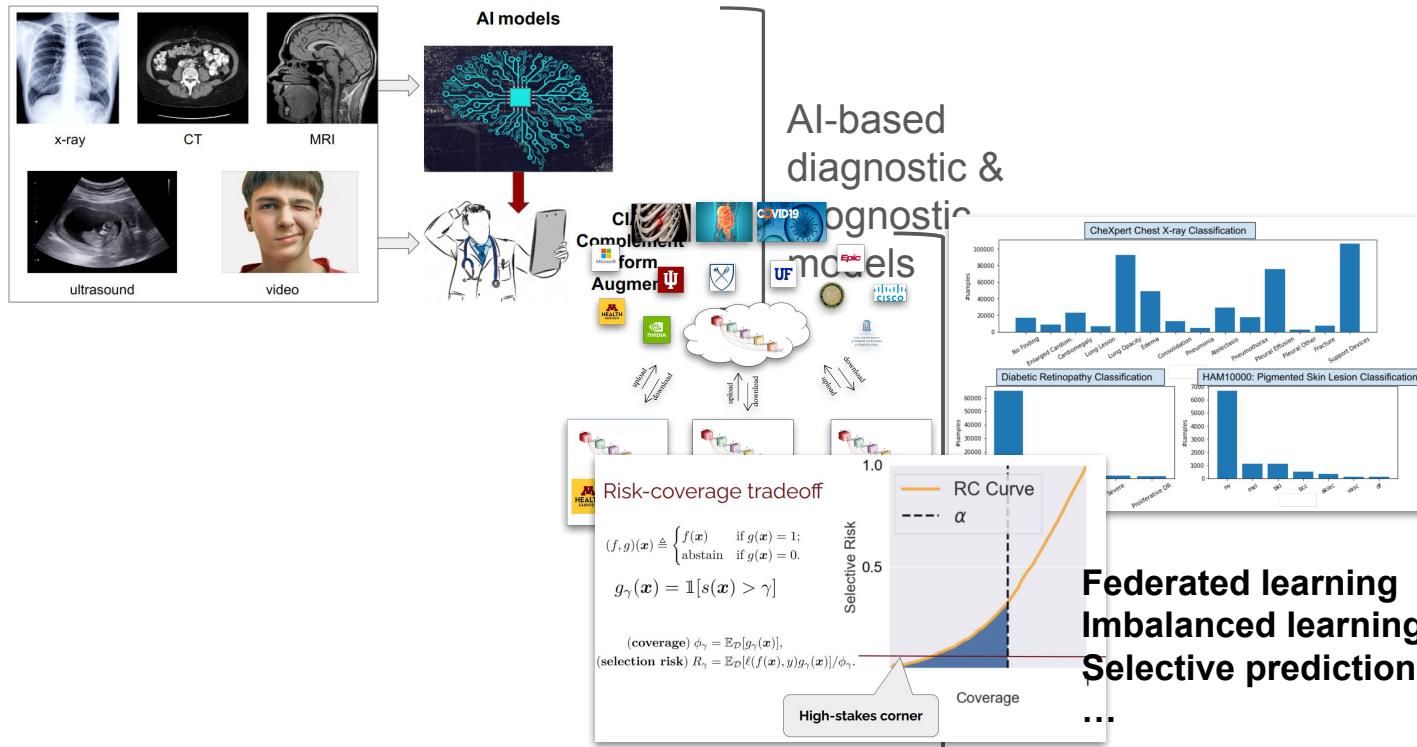
**Saving Foundation Flow-Matching Priors for Inverse Problems**

Yuxiang Wan, Ryan Devera, Wenjie Zhang, Ju Sun

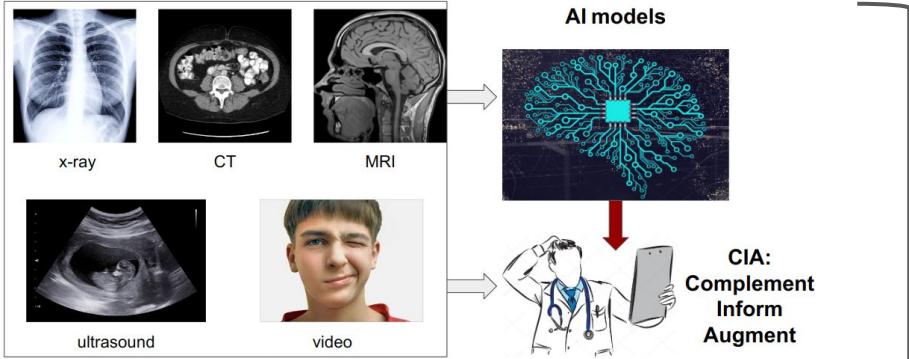
Submitted to ICLR'25, with **foundation** priors (e.g., stable diffusion models)

Challenge: Foundation generative models are powerful because **they're not specific**

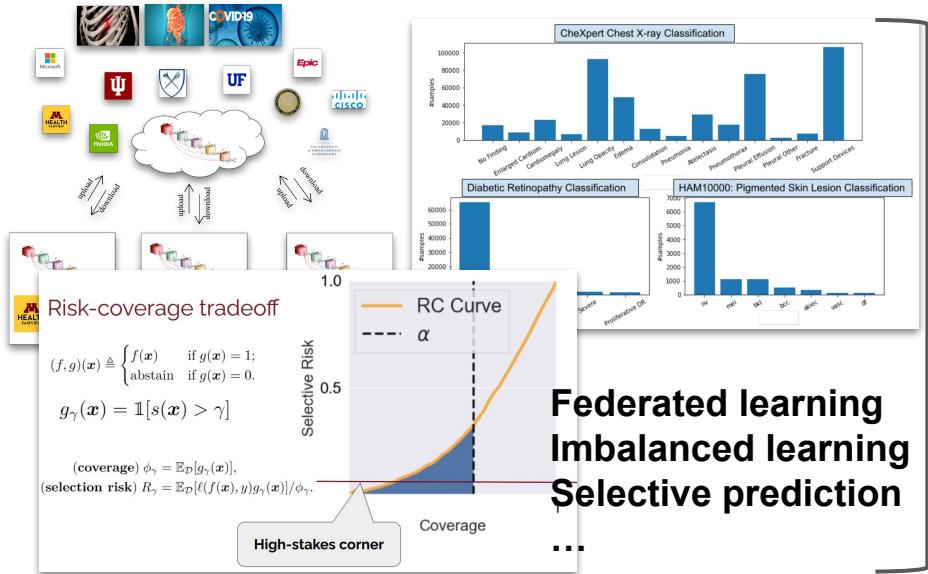
# **Applied and basic research in AI for Healthcare**



Basic research to address data sparsity, data imbalance, and AI safety in healthcare (& beyond)



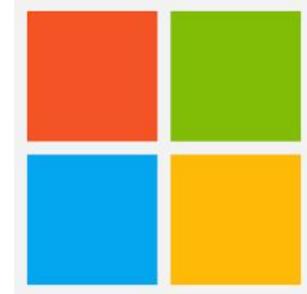
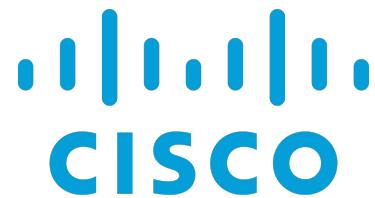
AI-based  
diagnostic  
models



Basic  
research to  
address data  
sparsity, data  
imbalance,  
and AI safety  
in healthcare  
(& beyond)

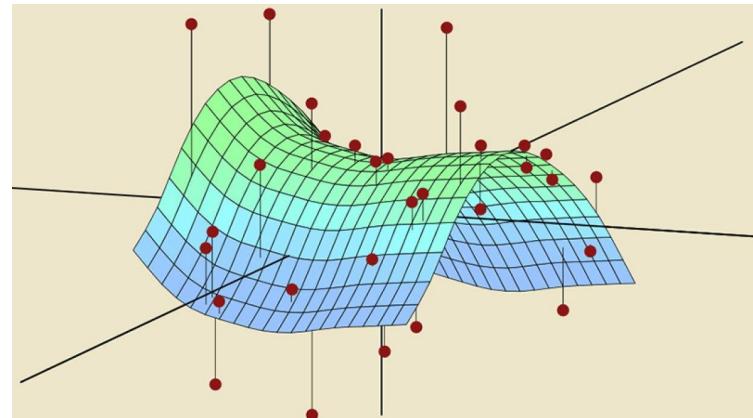
Supported by  
3+ (4th  
forthcoming)  
active NIH R01  
grants

Thanks to our funders



# Practical Optimization for Deep Learning with Nontrivial Constraints

# Three fundamental questions in DL



- **Approximation:** is it powerful, i.e., the  $\mathcal{H}$  large enough for all possible weights?
- **Optimization:** how to solve
$$\min_{\mathbf{w}'_i s, \mathbf{b}'_i s} \frac{1}{n} \sum_{i=1}^n \ell [\mathbf{y}_i, \{\text{NN}(\mathbf{w}_1, \dots, \mathbf{w}_k, b_1, \dots, b_k)\}(\mathbf{x}_i)]$$
- **Generalization:** does the learned NN work well on “similar” data?

# Algorithms

## Isn't it solved?

### Base class

CLASS `torch.optim.Optimizer(params, defaults)` [source]

Base class for all optimizers.

#### • WARNING

Parameters need to be specified as collections consistent between runs. Examples of objects and iterators over values of dictionaries.

#### Parameters:

- **params** (*iterable*) – an iterable of `Tensor`s

Tensors should be optimized.

- **defaults** – (*dict*): a dict containing default values for parameters

when a parameter group doesn't specify them.

`Adadelta`

Implements Adadelta algorithm.

`Adagrad`

Implements Adagrad algorithm.

`Adamax`

Implements Adamax algorithm (a variant of Adam based on infinity norm).

`ASGD`

Implements Averaged Stochastic Gradient Descent.

`LBFGS`

Implements L-BFGS algorithm, heavily inspired by `minFunc`.

`NAdam`

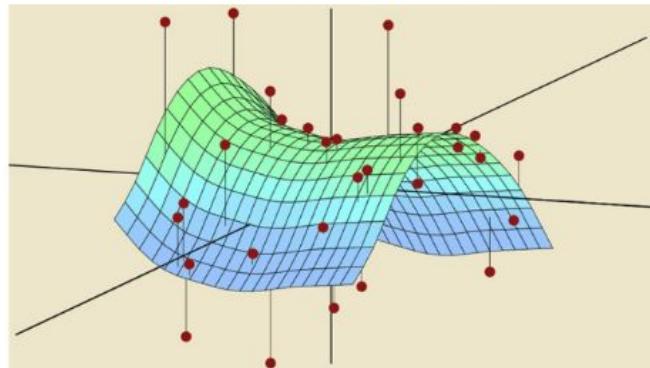
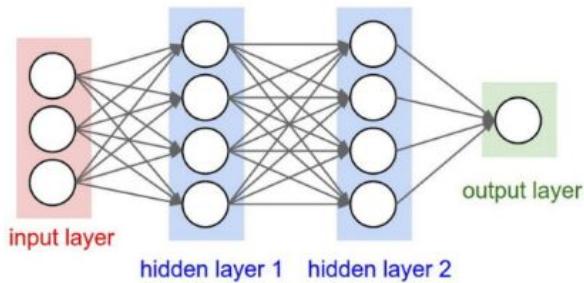
Implements NAdam algorithm.

`RAdam`

Implements RAdam algorithm.

# When DL meets constraints

Artificial neural networks



used to approximate nonlinear functions

## Unconstrained optimization

$$\min_{\boldsymbol{w}'_i s, \boldsymbol{b}'_i s} \frac{1}{n} \sum_{i=1}^n \ell [y_i, \{\text{NN}(\boldsymbol{w}_1, \dots, \boldsymbol{w}_k, \boldsymbol{b}_1, \dots, \boldsymbol{b}_k)\}(\boldsymbol{x}_i)]$$
$$\min_{\boldsymbol{x}} f(\boldsymbol{x})$$

**“Solved”**

## Constrained optimization

$$\min_{\boldsymbol{x}} f(\boldsymbol{x}) \quad \text{s. t. } g(\boldsymbol{x}) \leq \mathbf{0}$$

**largely “unsolved”**

# This talk is about GAPS

$$\min_{\mathbf{x}} f(\mathbf{x}) \quad \text{s. t. } g(\mathbf{x}) \leq \mathbf{0}$$

largely “unsolved”



An imaginary chat between a PhD student working in deep learning (**DLP**) and a PhD student working in optimization (**OP**)

DLP: Man, I've solved a constrained DL problem recently

OP: Oh, that's a hard problem

DLP: Really? I actually did it

OP: How?

DLP: My problem is  $\min_x f(x), \text{s.t. } g(x) \leq \mathbf{0}$ . I put  $g(x)$  as a penalty and then call ADAM

OP: Are you sure it works?

DLP: Yes, the performance is improved and my paper is published at ICML

OP: Why don't you try augmented Lagrangian methods?

DLP: No implementation in Pytorch. Is it possible we work out some theory about my method?

OP: I think it's hard. It's not convex

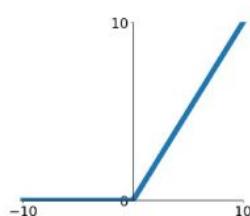
# Outline

## Constrained deep learning: CDL

- **What, how, and why for CDL**
- No good solvers for CDL yet
- Granso and PyGranso
- PyGranso in action
- Outlook

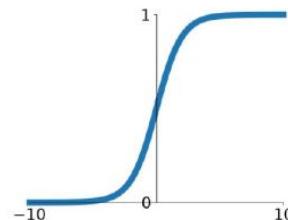
# DL with simple constraints

## Embedding constraints into DL models



**ReLU**  
(Rectified Linear Unit)

Nonnegativity



**Sigmoid**

[0, 1]

$$z \mapsto \left[ \frac{e^{z_1}}{\sum_j e^{z_j}}, \dots, \frac{e^{z_p}}{\sum_j e^{z_j}} \right]^\top$$

**Softmax**

Nonnegativity and summed to 1

# DL with nontrivial constraints

- Robustness evaluation
- Imbalanced learning
- Topology optimization
- Contrastive learning

Navigation: Home | About | Contact | Log In

## Deep Learning with Nontrivial Constraints: Methods and Applications

**Chuan He<sup>1</sup>, Ryan Devera<sup>1</sup>, Wenjie Zhang<sup>1</sup>, Ying Cui<sup>2</sup>, Zhaosong Lu<sup>3</sup> and Ju Sun<sup>1</sup>**

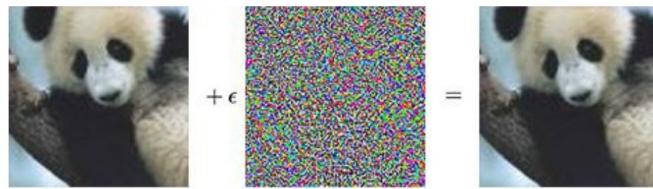
<sup>1</sup>Computer Science and Engineering, University of Minnesota

<sup>2</sup>Industrial Engineering and Operations Research, University of California, Berkeley

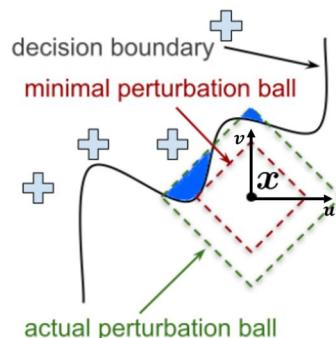
<sup>3</sup>Industrial and Systems Engineering, University of Minnesota

{he000233, dever120, zhan7867}@umn.edu, yingcui@berkeley.edu, {zhaosong, jusun}@umn.edu

# Robustness evaluation (RE)



$$\mathbf{x} + \delta = \mathbf{x}'$$



Maximize loss function

$$\max_{\mathbf{x}'} \ell(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}'))$$

s. t.  $d(\mathbf{x}, \mathbf{x}') \leq \varepsilon$ ,  $\mathbf{x}' \in [0, 1]^n$

Allowable perturbation      Valid image

Minimize robustness radius

$$\min_{\mathbf{x}'} d(\mathbf{x}, \mathbf{x}')$$

s. t.  $\max_{i \neq y} f_{\boldsymbol{\theta}}^i(\mathbf{x}') \geq f_{\boldsymbol{\theta}}^y(\mathbf{x}')$ ,  $\mathbf{x}' \in [0, 1]^n$

Change the predicted class      Valid image

# Projected gradient descent (PGD) for RE

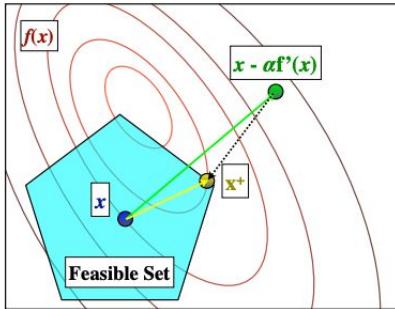
$$\min_{\mathbf{x} \in \mathcal{Q}} f(\mathbf{x})$$

**Step size**

$$\mathbf{x}_{k+1} = P_{\mathcal{Q}} \left( \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}_k) \right)$$

$$P_{\mathcal{Q}}(\mathbf{x}_0) = \arg \min_{\mathbf{x} \in \mathcal{Q}} \frac{1}{2} \|\mathbf{x} - \mathbf{x}_0\|_2^2$$

**Projection operator**



**Key hyperparameters:**  
 (1) step size  
 (2) iteration number

$$\max_{\mathbf{x}'} \ell(\mathbf{y}, f_{\theta}(\mathbf{x}'))$$

s. t.  $d(\mathbf{x}, \mathbf{x}') \leq \varepsilon$ ,  $\mathbf{x}' \in [0, 1]^n$

## Algorithm 1 APGD

```

1: Input:  $f, S, x^{(0)}, \eta, N_{\text{iter}}, W = \{w_0, \dots, w_n\}$ 
2: Output:  $x_{\max}, f_{\max}$ 
3:  $x^{(1)} \leftarrow P_S(x^{(0)} + \eta \nabla f(x^{(0)}))$ 
4:  $f_{\max} \leftarrow \max\{f(x^{(0)}), f(x^{(1)})\}$ 
5:  $x_{\max} \leftarrow x^{(0)}$  if  $f_{\max} \equiv f(x^{(0)})$  else  $x_{\max} \leftarrow x^{(1)}$ 
6: for  $k = 1$  to  $N_{\text{iter}} - 1$  do
7:    $z^{(k+1)} \leftarrow P_S(x^{(k)} + \eta \nabla f(x^{(k)}))$ 
8:    $x^{(k+1)} \leftarrow P_S \left( x^{(k)} + \alpha(z^{(k+1)} - x^{(k)}) \right.$ 
       $\quad \left. + (1 - \alpha)(x^{(k)} - x^{(k-1)}) \right)$ 
9:   if  $f(x^{(k+1)}) > f_{\max}$  then
10:     $x_{\max} \leftarrow x^{(k+1)}$  and  $f_{\max} \leftarrow f(x^{(k+1)})$ 
11:   end if
12:   if  $k \in W$  then
13:     if Condition 1 or Condition 2 then
14:        $\eta \leftarrow \eta/2$  and  $x^{(k+1)} \leftarrow x_{\max}$ 
15:     end if
16:   end if
17: end for

```

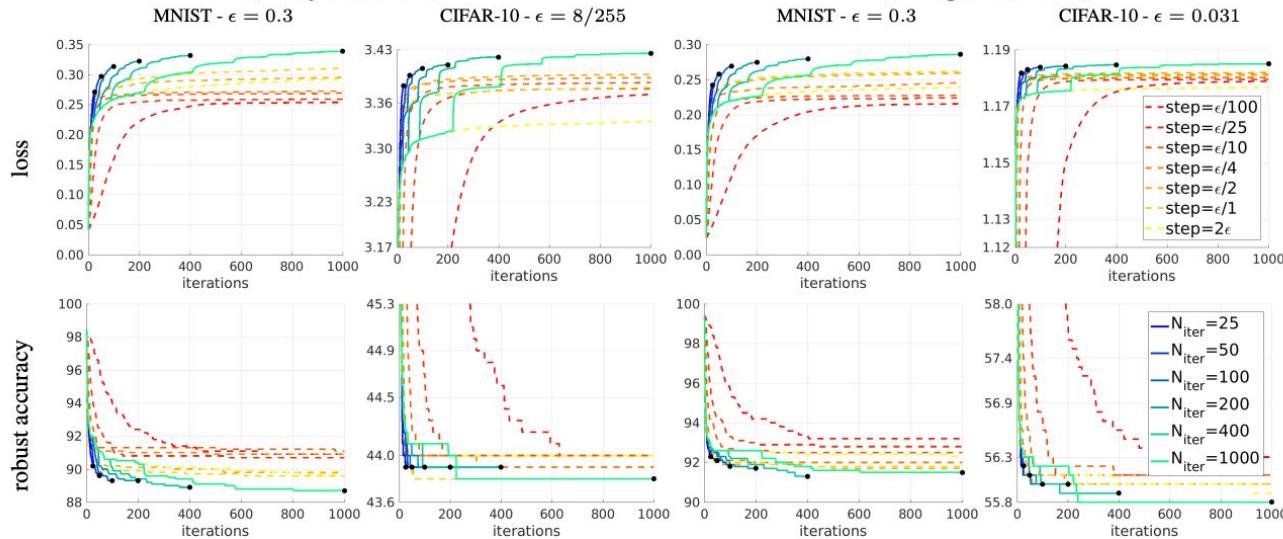
Ref [https://angms.science/doc/CVX/CVX\\_PGD.pdf](https://angms.science/doc/CVX/CVX_PGD.pdf)

<https://www.cs.ubc.ca/~schmidtm/Courses/5XX-S20/S5.pdf>

Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks. Croce, F., Hein, M., ICML 2020

<https://arxiv.org/pdf/2003.01690.pdf>

# Problem with PGD



$$\max_{\mathbf{x}'} \ell(\mathbf{y}, f_{\boldsymbol{\theta}}(\mathbf{x}'))$$

$$\text{s.t. } d(\mathbf{x}, \mathbf{x}') \leq \varepsilon, \quad \mathbf{x}' \in [0, 1]^n$$

**Tricky to set:  
iteration number & step size  
i.e., tricky to decide where to stop**

# RE: penalty methods for complicated d

$$\begin{aligned} & \max_{\mathbf{x}'} \ell(\mathbf{y}, f_{\theta}(\mathbf{x}')) \\ \text{s.t. } & d(\mathbf{x}, \mathbf{x}') \leq \varepsilon, \quad \mathbf{x}' \in [0, 1]^n \end{aligned}$$

$d(\mathbf{x}, \mathbf{x}') \doteq \|\phi(\mathbf{x}) - \phi(\mathbf{x}')\|_2$       **perceptual distance**  
where  $\phi(\mathbf{x}) \doteq [\hat{g}_1(\mathbf{x}), \dots, \hat{g}_L(\mathbf{x})]$

**Projection onto the constraint is complicated**

## Penalty methods

$$\max_{\tilde{\mathbf{x}}} \quad \mathcal{L}(f(\tilde{\mathbf{x}}), y) - \lambda \max \left( 0, \|\phi(\tilde{\mathbf{x}}) - \phi(\mathbf{x})\|_2 - \epsilon \right)$$

Solve it for each fixed  $\lambda$  and then increase  $\lambda$

---

### Algorithm 2 Lagrangian Perceptual Attack (LPA)

---

```
1: procedure LPA(classifier network  $f(\cdot)$ , LPIPS distance  $d(\cdot, \cdot)$ , input  $\mathbf{x}$ , label  $y$ , bound  $\epsilon$ )
2:    $\lambda \leftarrow 0.01$ 
3:    $\tilde{\mathbf{x}} \leftarrow \mathbf{x} + 0.01 * \mathcal{N}(0, 1)$             $\triangleright$  initialize perturbations with random Gaussian noise
4:   for  $i$  in  $1, \dots, S$  do            $\triangleright$  we use  $S = 5$  iterations to search for the best value of  $\lambda$ 
5:     for  $t$  in  $1, \dots, T$  do            $\triangleright T$  is the number of steps
6:        $\Delta \leftarrow \nabla_{\tilde{\mathbf{x}}} [\mathcal{L}(f(\tilde{\mathbf{x}}), y) - \lambda \max(0, d(\tilde{\mathbf{x}}, \mathbf{x}) - \epsilon)]$             $\triangleright$  take the gradient of (5)
7:        $\hat{\Delta} \leftarrow \Delta / \|\Delta\|_2$             $\triangleright$  normalize the gradient
8:        $\eta = \epsilon * (0.1)^{t/T}$             $\triangleright$  the step size  $\eta$  decays exponentially
9:        $m \leftarrow d(\tilde{\mathbf{x}}, \tilde{\mathbf{x}} + h\hat{\Delta})/h$             $\triangleright m \approx$  derivative of  $d(\tilde{\mathbf{x}}, \cdot)$  in the direction of  $\hat{\Delta}$ ;  $h = 0.1$ 
10:       $\tilde{\mathbf{x}} \leftarrow \tilde{\mathbf{x}} + (\eta/m)\hat{\Delta}$             $\triangleright$  take a step of size  $\eta$  in LPIPS distance
11:    end for
12:    if  $d(\tilde{\mathbf{x}}, \mathbf{x}) > \epsilon$  then
13:       $\lambda \leftarrow 10\lambda$             $\triangleright$  increase  $\lambda$  if the attack goes outside the bound
14:    end if
15:  end for
16:   $\tilde{\mathbf{x}} \leftarrow \text{PROJECT}(d, \tilde{\mathbf{x}}, \mathbf{x}, \epsilon)$ 
17:  return  $\tilde{\mathbf{x}}$ 
18: end procedure
```

---

# Problem with penalty methods

Method	cross-entropy loss		margin loss	
	Viol. (%) ↓	Att. Succ. (%) ↑	Viol. (%) ↓	Att. Succ. (%) ↑
Fast-LPA	73.8	3.54	41.6	56.8
LPA	<b>0.00</b>	80.5	<b>0.00</b>	97.0
PPGD	5.44	25.5	<b>0.00</b>	38.5
PWCF (ours)	0.62	<b>93.6</b>	<b>0.00</b>	<b>100</b>

**LPA, Fast-LPA:** penalty methods

**PPGD:** Projected gradient descent

Penalty methods tend to encounter

**large constraint violation** (i.e., infeasible solution, known in optimization theory) or **suboptimal solution**

$$\begin{aligned} & \max_{\mathbf{x}'} \ell(\mathbf{y}, f_{\theta}(\mathbf{x}')) \\ \text{s.t. } & d(\mathbf{x}, \mathbf{x}') \leq \varepsilon, \quad \mathbf{x}' \in [0, 1]^n \\ & d(\mathbf{x}, \mathbf{x}') \doteq \|\phi(\mathbf{x}) - \phi(\mathbf{x}')\|_2 \\ \text{where } & \phi(\mathbf{x}) \doteq [\hat{g}_1(\mathbf{x}), \dots, \hat{g}_L(\mathbf{x})] \end{aligned}$$

**PWCF**, an optimizer with a principled stopping criterion on **stationarity & feasibility**

**Unreliable optimization ==  
Unreliable RE**

# Outline

## Constrained deep learning: CDL

- What, how, and why for CDL
- **No good solvers for CDL yet**
- Granso and PyGranso
- PyGranso in action
- Outlook

# DL frameworks



## JAX: Autograd and XLA



For unconstrained DL problems

# Convex optimization solvers and frameworks



Modeling languages



**SDPT<sup>3</sup>** - a M<sub>ATLAB</sub> software package for  
semidefinite-quadratic-linear programming

[K. C. Toh](#), [R. H. Tütüncü](#), and [M. J. Todd](#).

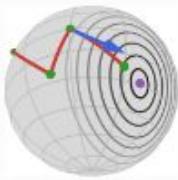
**TFOCS: Templates for First-Order Conic Solvers**

Solvers

**Not for DL**, which involves NCVX optimization

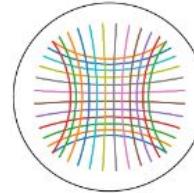
Note: Gurobi can handle certain NCVX problems

# Manifold optimization



Manopt.jl

Geomstats



$\mathcal{T}_p \mathcal{G}$   
geoopt

**McTorch Lib, a manifold optimization library for deep learning**

---

Only for **differentiable manifolds constraints**

# General constrained optimization

**KNITRO®**

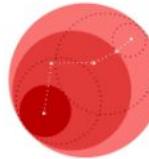


IPOPT

Interior-point methods



Cooper



**ensmallen**

flexible C++ library for efficient numerical optimization

## GENO

Augmented Lagrangian methods

Lagrangian-method-based constrained optimization

TensorFlow Constrained Optimization (TFCO)

# Specialized ML packages



Problem-specific solvers that **cannot be easily extended** to new formulations

# Outline

## Constrained deep learning: CDL

- What, how, and why for CDL
- No good solvers for CDL yet
- **Granso and PyGranso**
- PyGranso in action
- Outlook

# Issues with typical CDL methods

## projected gradient descent

$$\min_{\mathbf{x} \in \mathcal{Q}} f(\mathbf{x})$$

$$\mathbf{x}_{k+1} = P_{\mathcal{Q}} \left( \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}_k) \right)$$

Issue: no principled stopping criterion/step size rules

## Lagrangian method

$$\min_{\mathbf{x}} \max_{\boldsymbol{\lambda} \geq 0} \hat{f}(\mathbf{x}) + \boldsymbol{\lambda}^\top g(\mathbf{x})$$

Idea: alternating minimize  $\mathbf{x}$  and maximize  $\boldsymbol{\lambda}$  via gradient descent

## penalty methods

$$\min_{\mathbf{x}} f(\mathbf{x}) \quad \text{s. t. } g(\mathbf{x}) \leq \mathbf{0}$$

$$\min_{\mathbf{x}} f(\mathbf{x}) + \lambda \max(0, g(\mathbf{x}))$$

Solved with increasing  $\lambda$ : sequence

Issue: infeasible solution

## Issues

- Infeasible solution
- Slow convergence

## Want

- Feasible & stationary solution
- Reasonable speed

# Principled answers to these questions

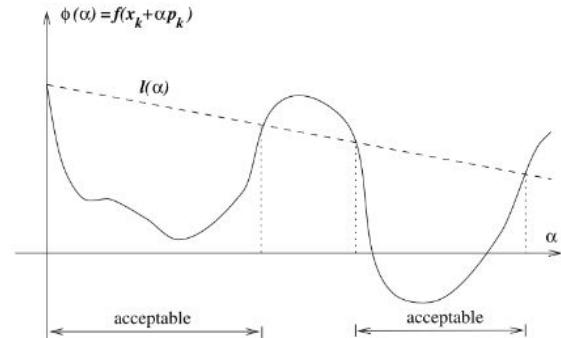
- **Feasible & stationary solution**

**Stationarity and feasibility check: KKT condition**

- **Reasonable speed**

**Line search**

- **A hidden problem: nonsmoothness**



Armijo (Sufficient Decrease) Condition



# Key algorithm

**Nonconvex, nonsmooth, constrained**

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}), \quad \text{s.t. } c_i(\mathbf{x}) \leq 0, \quad \forall i \in \mathcal{I}; \quad c_i(\mathbf{x}) = 0, \quad \forall i \in \mathcal{E}.$$

**Penalty sequential quadratic programming (P-SQP)**

$$\begin{aligned} \min_{d \in \mathbb{R}^n, s \in \mathbb{R}^p} \quad & \mu(f(x_k) + \nabla f(x_k)^T d) + e^T s + \frac{1}{2} d^T H_k d \\ \text{s.t.} \quad & c(x_k) + \nabla c(x_k)^T d \leq s, \quad s \geq 0, \end{aligned}$$

Ref: **Curtis, Frank E., Tim Mitchell, and Michael L. Overton.** "A BFGS-SQP method for nonsmooth, nonconvex, constrained optimization and its evaluation using relative minimization profiles." *Optimization Methods and Software* 32.1 (2017): 148-181.

# Algorithm highlights

## Steering strategy for the penalty parameter

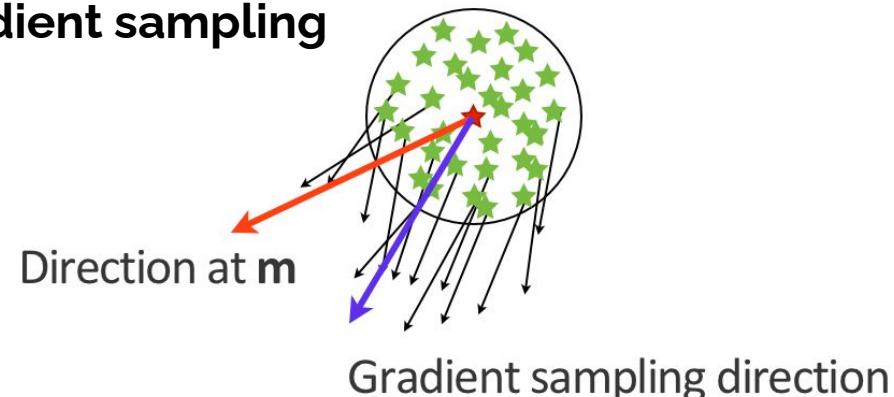
If feasibility improvement is insufficient :  $l_\delta(d_k; x_k) < c_v v(x_k)$

## Stationarity based on (approximate) gradient sampling

$$G_k := [\nabla f(x^k) \quad \nabla f(x^{k,1}) \quad \cdots \quad \nabla f(x^{k,m})]$$

$$\min_{\lambda \in \mathbb{R}^{m+1}} \frac{1}{2} \|G_k \lambda\|_2^2$$

$$\text{s.t. } \mathbf{1}^T \lambda = 1, \quad \lambda \geq 0$$



# Key take-away



- Principled stopping criterion and line search, to obtain a **solution with certificate** (stationarity & feasibility check)
- Quasi-newton style method for fast convergence, i.e.,  
**reasonable speed and high-precision solution**

**Ref** Curtis, Frank E., Tim Mitchell, and Michael L. Overton. "A BFGS-SQP method for nonsmooth, nonconvex, constrained optimization and its evaluation using relative minimization profiles." Optimization Methods and Software 32.1 (2017): 148-181.



# Limitations of GRANSO

```
% Gradient of inner product with respect to A  
f_grad      = imag((conj(Bty)*Cx.')/(y'*x));  
f_grad      = f_grad(:);  
  
% Gradient of inner product with respect to A  
ci_grad     = real((conj(Bty)*Cx.')/(y'*x));  
ci_grad     = ci_grad(:);
```

**analytical gradients required**

```
p           = size(B,2);  
m           = size(C,1);  
X           = reshape(x,p,m);
```

**vector variables only**

**Lack of Auto-Differentiation**

**Lack of GPU Support**

**No native support of tensor variables**

**⇒ impossible to do deep learning with GRANSO**

# GRANSO meets PyTorch



NCVX PyGRANSO  
Documentation

Search the docs ...

Introduction

Installation

Settings

Examples



NCVX Package

**NCVX: A General-Purpose Optimization Solver for Constrained Machine and Deep Learning**

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}), \text{ s.t. } c_i(\mathbf{x}) \leq 0, \forall i \in \mathcal{I}; c_i(\mathbf{x}) = 0, \forall i \in \mathcal{E}$$

**First general-purpose solver for constrained DL problems**

# Outline

## Constrained deep learning: CDL

- What and how for CDL
- Why CDL
- No good solvers for CDL yet
- Granso and PyGranso
- **PyGranso in action**
- Outlook

# Example 1: Support Vector Machine (SVM)

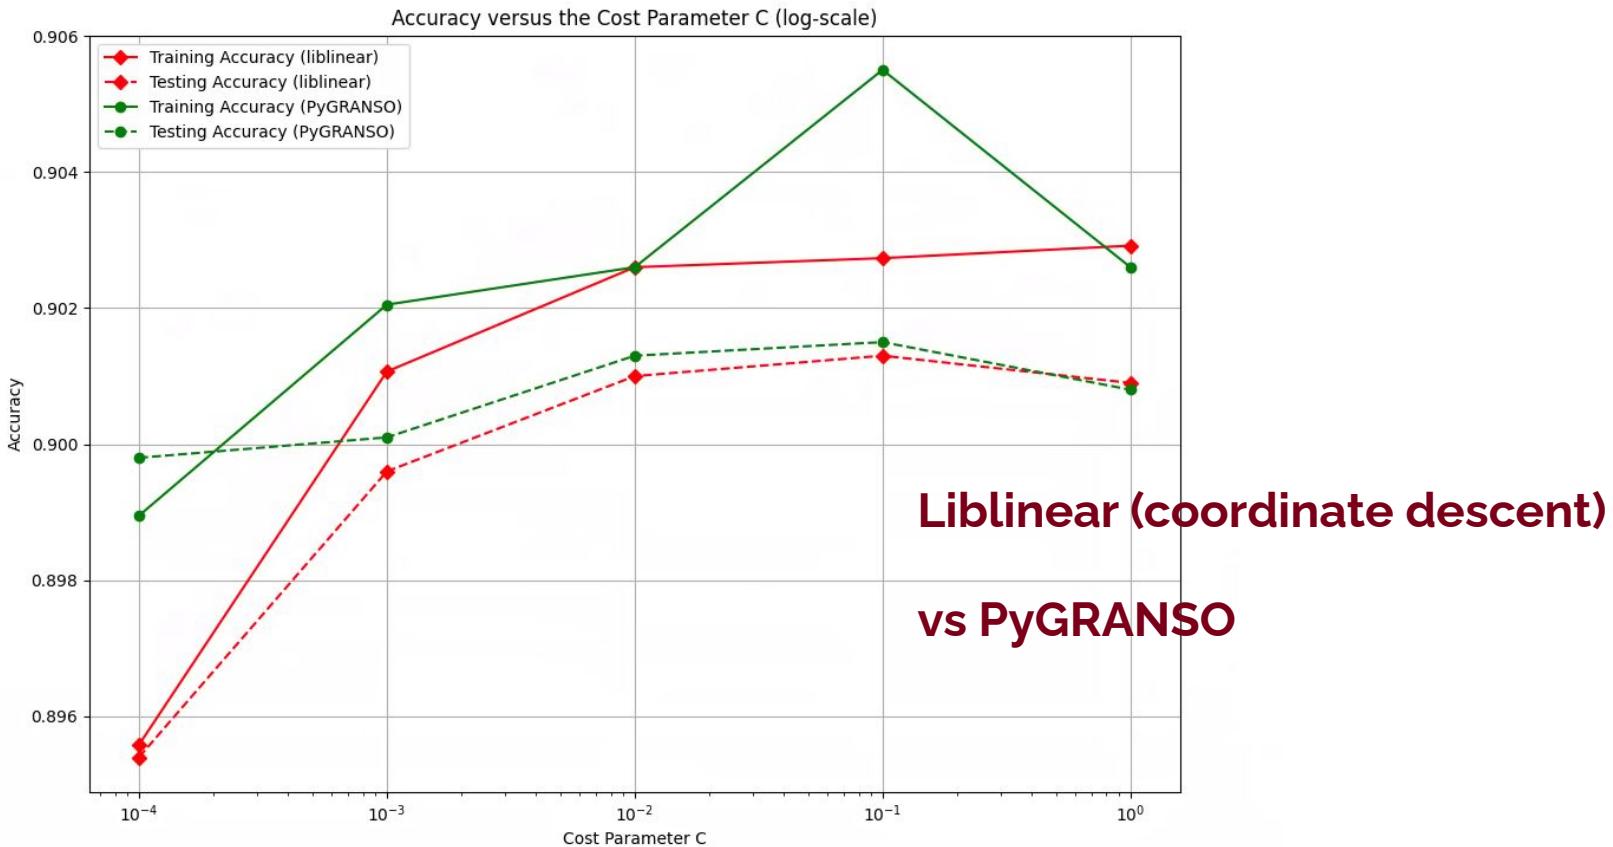
## Soft-margin SVM

$$\min_{\mathbf{w}, b, \zeta} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \zeta_i$$

$$\text{s.t. } y_i (\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \zeta_i, \quad \zeta_i \geq 0 \quad \forall i = 1, \dots, n$$

```
def comb_fn(X_struct):
    # obtain optimization variables
    w = X_struct.w
    b = X_struct.b
    zeta = X_struct.zeta
    # objective function
    f = 0.5*w.T@w + C*torch.sum(zeta)
    # inequality constraints
    ci = pygransoStruct()
    ci.c1 = 1 - zeta - y*(x@w+b)
    ci.c2 = -zeta
    # equality constraint
    ce = None
    return [f,ci,ce]
# specify optimization variables
var_in = {"w": [d,1], "b": [1,1], "zeta": [n,1]}
# pygranso main algorithm
soln = pygranso(var_in,comb_fn)
```

# Binary classification (odd vs even digits) on MNIST dataset



## Example 2: Robustness eval. —min formulation

$$\min_{\mathbf{x}'} d(\mathbf{x}, \mathbf{x}')$$

$$\text{s. t. } \max_{\ell \neq c} f_{\boldsymbol{\theta}}^{\ell}(\mathbf{x}') \geq f_{\boldsymbol{\theta}}^c(\mathbf{x}')$$

$$\mathbf{x}' \in [0, 1]^n$$

```
def comb_fn(X_struct):
    # obtain optimization variables
    x_prime = X_struct.x_prime
    # objective function
    f = d(x,x_prime)
    # inequality constraints
    ci = pygransoStruct()
    f_theta_all = f_theta(x_prime)
    fy = f_theta_all[:,y] # true class output
    # output except true class
    fi = torch.hstack((f_theta_all[:, :y], f_theta_all[:, y+1:]))
    ci.c1 = fy - torch.max(fi)
    ci.c2 = -x_prime
    ci.c3 = x_prime-1
    # equality constraint
    ce = None
    return [f,ci,ce]
# specify optimization variable (tensor)
var_in = {"x_prime": list(x.shape)}
# pygranso main algorithm
soln = pygranso(var_in,comb_fn)
```

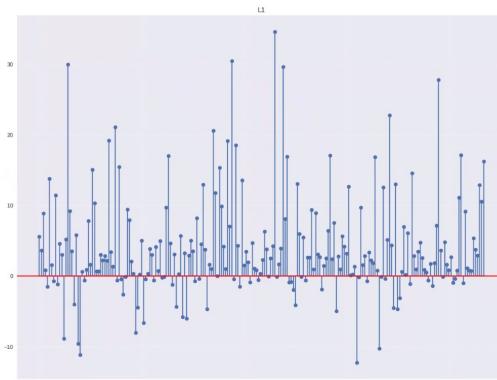
## CIFAR10 dataset

Compared with FAB [iterative constraint linearization + projected gradient]

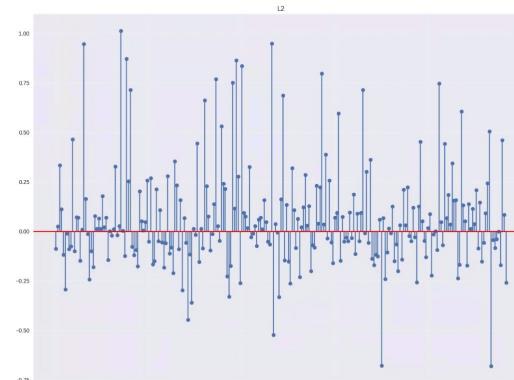
<https://github.com/fra31/auto-attack>

$$\begin{aligned} & \min_{\mathbf{x}'} d(\mathbf{x}, \mathbf{x}') \\ \text{s. t. } & \max_{\ell \neq c} f_{\boldsymbol{\theta}}^{\ell}(\mathbf{x}') \geq f_{\boldsymbol{\theta}}^c(\mathbf{x}') \\ & \mathbf{x}' \in [0, 1]^n \end{aligned}$$

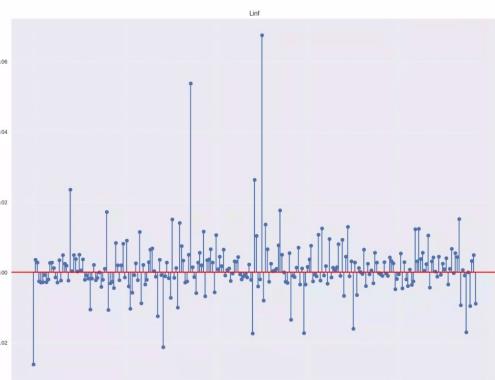
X-axis: image index; Y-axis: PyGRANSO radius - FAB radius



L1 attack



L2 attack



Linf attack

Many  
others

<https://ncvx.org/>

NCVX PyGRANSO Documentation

← ⌂ ⌂ ⌂

Search the docs ...

Introduction

Installation

Settings

Examples

Rosenbrock

Eigenvalue Optimization

Dictionary Learning

Nonlinear Feasibility Problem

Sphere Manifold

Trace Optimization

Robust PCA

Generalized LASSO

Logistic Regression

LeNet5

Perceptual Attack

Orthogonal RNN

Highlights

# Home



## NCVX Package

NCVX (**N**on**C**onVe**X**) is a user-friendly and scalable python software package targeting general nonsmooth NCVX problems with nonsmooth constraints. **NCVX** is being developed by **GLOVEX** at the Department of Computer Science & Engineering, University of Minnesota, Twin Cities.

The initial release of **NCVX** contains the solver **PyGRANSO**, a PyTorch-enabled port of **GRANSO** incorporating auto-differentiation, GPU acceleration, tensor input, and support for new QP solvers. As a highlight, **PyGRANSO** can solve general constrained deep learning problems, the first of its kind.



# Our current developments

- PyGrano: deterministic SQP solver
- More scalable deterministic solvers (exact penalty, single-loop ALM, etc)
- Algorithm foundations for stochastic objective with stochastic constraints/numerous constraints
- Scalable solvers for stochastic objective with stochastic constraints/numerous constraints

# Closing



## Deep Learning with Nontrivial Constraints: Methods and Applications

Chuan He<sup>1</sup>, Ryan Devera<sup>1</sup>, Wenjie Zhang<sup>1</sup>, Ying Cui<sup>2</sup>, Zhaosong Lu<sup>3</sup> and Ju Sun<sup>1</sup>

<sup>1</sup>Computer Science and Engineering, University of Minnesota

<sup>2</sup>Industrial Engineering and Operations Research, University of California, Berkeley

<sup>3</sup>Industrial and Systems Engineering, University of Minnesota

{he000233, dever120, zhan7867}@umn.edu, yingcui@berkeley.edu, {zhaosong, jusun}@umn.edu



PyGRANSO  
Documentation

Search the docs ...

Introduction  
Installation  
Settings  
Examples

Home



NCVX Package

NCVX: A General-Purpose Optimization Solver for  
Constrained Machine and Deep Learning

Biyun Liang, Tim Mitchell, Ju Sun

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}), \text{ s.t. } c_i(\mathbf{x}) \leq 0, \forall i \in \mathcal{I}; c_i(\mathbf{x}) = 0, \forall i \in \mathcal{E}$$

First general-purpose solver for constrained DL  
problems

# Thanks to all contributors



**Prof. Tim Mitchell**  
(CS, Queens Col.)



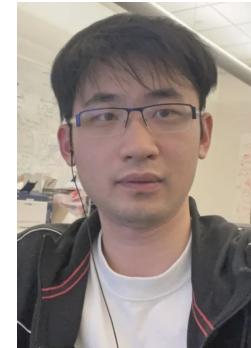
**Prof. Ying Cui**  
(IEOR, UC  
Berkeley)



**Prof. Qizhi He**  
(CSGE, UMN)



**Prof.  
Zhaosong Lu**  
(CSGE, UMN)



**Dr. Chuan He**  
(CS&E, UMN)

# Thanks to all contributors



**Buyun Liang**  
(CSE U Penn;  
Formally, CS&E, UMN)

**Ryan de Vera**  
(CS&E, UMN)



**Hengyue Liang**  
(ECE, UMN)



**Wenjie Zhang**  
(CS&E, UMN)



**Yash Travadi**  
(CS&E, UMN)



**Le Peng**  
(CS&E, UMN)