

Basics of Numerical Optimization: Computing Derivatives

Ju Sun

Computer Science & Engineering
University of Minnesota, Twin Cities

October 19, 2020


Logistics

- HW 2 out. Due: Oct 28
- Project grouping

Group ID (think of a cool name)	Group Members (Pis enter full name followed by email ID)	
Masked Preachers	Ju Sun, jusun	Hengkang Wang, wang9881
Abraca-data	Connor Theisen, theis417	Miguel Miguélez, miguel022
The Hyperclique	Mario Serraforo, serra082	Nathan Gessesse, gess0042
Anti-Virus	Ahmet Semi Asarkaya, asark001	Won Joon Choi, choi471
The ABC	Aaron Koenigsberg, koen296	Brandon Voigt, voigt227
Baseline	Qingyuan Jiang, jian0345	Jun-Jee Chao, chao0107
GISer	Chenxi Lin, lin00370	Jiayang Chen, chen6691
TBD	Elliot Orenstein, orens040	Sam Walczak, walcz076
NLP (No like Pandemic)	Tyler Wendland, wendl155	Risako Owan owan0002
Fashion Similarity Squad	Chris Thielsen, thiels391	Maryam Forootani Nia foroo002
Hakuna Matata	Xiang Li, lixx5000	Mingqian Duan, mduan
HYY_Gogogo	Haoyu Yang, yang6993	Shaoming Xu, xu000114
Puppybot	Matthew Tlachac, tlach007	Michael Lucke, lucke096
Heaven Canceller	Kangyu Zheng, zhen0196	Yanjun Cui, cui00022
	Changye Li lixx3013	
Professor Sun is very handsome	Myat Mo, mo000007	King Yiu Suen, suenx008

some cool ID's!

Derivatives for numerical optimization




Credit: aria42.com

- gradient descent
- Newton's method
- momentum methods
- quasi-Newton methods
- coordinate descent
- conjugate gradient methods
- trust-region methods

- Almost all methods entail low-order derivatives, i.e., gradient and/or Hessian, to proceed.
 - * 1st order methods: use $f(\mathbf{x})$ and $\nabla f(\mathbf{x})$
 - * 2nd order methods: use $f(\mathbf{x})$ and $\nabla f(\mathbf{x})$ and $\nabla^2 f(\mathbf{x})$
- **Numerical derivatives** (i.e., numbers) needed for the iterations

This lecture: how to compute the numerical derivatives

Four kinds of computing techniques



Outline

Analytic differentiation

Finite-difference approximation

Automatic differentiation

Differentiable programming

Suggested reading

Analytic derivatives

Idea: derive the analytic derivatives first, then make numerical substitution

To derive the analytic derivatives **by hand**:

- **Chain rule (vector version) method**

Let $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$ and $h : \mathbb{R}^n \rightarrow \mathbb{R}^k$, and f is differentiable at \mathbf{x} and $z = h(\mathbf{y})$ is differentiable at $\mathbf{y} = f(\mathbf{x})$. Then, $z = h \circ f(\mathbf{x}) : \mathbb{R}^m \rightarrow \mathbb{R}^k$ is differentiable at \mathbf{x} , and

$$\mathbf{J}_{[h \circ f]}(\mathbf{x}) = \mathbf{J}_h(f(\mathbf{x})) \mathbf{J}_f(\mathbf{x}), \text{ or } \frac{\partial z}{\partial \mathbf{x}} = \frac{\partial z}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{x}}$$

When $k = 1$,

$$\nabla [h \circ f](\mathbf{x}) = \mathbf{J}_f^\top(\mathbf{x}) \nabla h(f(\mathbf{x})).$$

- **Taylor expansion method**

Expand the perturbed function $f(\mathbf{x} + \boldsymbol{\delta})$ and then match it against Taylor expansions to read off the gradient and/or Hessian:

$$f(\mathbf{x} + \boldsymbol{\delta}) = f(\mathbf{x}) + \langle \nabla f(\mathbf{x}), \boldsymbol{\delta} \rangle + o(\|\boldsymbol{\delta}\|_2)$$

$$f(\mathbf{x} + \boldsymbol{\delta}) = f(\mathbf{x}) + \langle \nabla f(\mathbf{x}), \boldsymbol{\delta} \rangle + \frac{1}{2} \langle \boldsymbol{\delta}, \nabla^2 f(\mathbf{x}) \boldsymbol{\delta} \rangle + o(\|\boldsymbol{\delta}\|_2^2)$$

Symbolic differentiation

Idea: derive the analytic derivatives first, then make numerical substitution

To derive the analytic derivatives **by software**:

Differentiate Function

Find the derivative of the function $\sin(x^2)$.

```
syms f(x)
f(x) = sin(x^2);
df = diff(f,x)
```

```
df(x) =
2*x*cos(x^2)
```

Find the value of the derivative at $x = 2$. Convert the value to double.

```
df2 = df(2)
```

```
df2 =
4*cos(4)
```

- Matlab (Symbolic Math Toolbox, `diff`)
- Python (SymPy, `diff`)
- Mathematica (`D`)

Effective for functions with few variables only

Outline

Analytic differentiation


Finite-difference approximation

Automatic differentiation

Differentiable programming

Suggested reading

Limitation of analytic differentiation



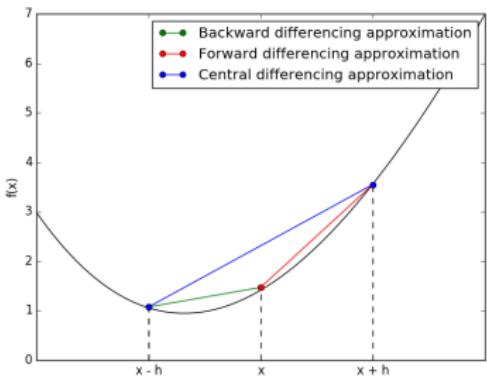
What is the gradient and/or Hessian of

$$f(\mathbf{W}) = \sum_i \| \mathbf{y}_i - \sigma(\mathbf{W}_k \sigma(\mathbf{W}_{k-1} \sigma \dots (\mathbf{W}_1 \mathbf{x}_i))) \|_F^2 ?$$

Applying the chain rule is boring and error-prone. Performing Taylor expansion is also tedious

Lesson we learn from tech history: leave boring jobs to computers

Approximate the gradient



(Credit: numex-blog.com)

Similarly, to approximate the Jacobian for $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}^m$:

$$\frac{\partial f_j}{\partial x_i} \approx \frac{f_j(\mathbf{x} + \delta e_i) - f_j(\mathbf{x})}{\delta} \quad (\text{one element each time})$$

$$\frac{\partial f}{\partial x_i} \approx \frac{f(\mathbf{x} + \delta e_i) - f(\mathbf{x})}{\delta} \quad (\text{one column each time})$$

$$J_f(\mathbf{x}) \mathbf{p} \approx \frac{f(\mathbf{x} + \delta p) - f(\mathbf{x})}{\delta} \quad (\text{directional})$$

central themes can also be derived

$$f'(\mathbf{x}) = \lim_{\delta \rightarrow 0} \frac{f(\mathbf{x} + \delta) - f(\mathbf{x})}{\delta} \approx \frac{f(\mathbf{x} + \delta) - f(\mathbf{x})}{\delta}$$

with δ sufficiently small

For $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$,

$$\frac{\partial f}{\partial x_i} \approx \frac{f(\mathbf{x} + \delta e_i) - f(\mathbf{x})}{\delta} \quad (\text{forward})$$

$$\frac{\partial f}{\partial x_i} \approx \frac{f(\mathbf{x}) - f(\mathbf{x} - \delta e_i)}{\delta} \quad (\text{backward})$$

$$\frac{\partial f}{\partial x_i} \approx \frac{f(\mathbf{x} + \delta e_i) - f(\mathbf{x} - \delta e_i)}{2\delta} \quad (\text{central})$$

Why central?

Stronger form of Taylor's theorems

- **1st order:** If $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$ is twice continuously differentiable,
$$f(\mathbf{x} + \delta) = f(\mathbf{x}) + \langle \nabla f(\mathbf{x}), \delta \rangle + O(\|\delta\|_2^2)$$
- **2nd order:** If $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$ is three-times continuously differentiable,
$$f(\mathbf{x} + \delta) = f(\mathbf{x}) + \langle \nabla f(\mathbf{x}), \delta \rangle + \frac{1}{2} \langle \delta, \nabla^2 f(\mathbf{x}) \delta \rangle + O(\|\delta\|_2^3)$$

Why the central theme is better?

- Forward: by 1st-order Taylor expansion
$$\frac{1}{\delta} (f(\mathbf{x} + \delta e_i) - f(\mathbf{x})) = \frac{1}{\delta} \left(\delta \frac{\partial f}{\partial x_i} + O(\delta^2) \right) = \frac{\partial f}{\partial x_i} + O(\delta)$$
- Central: by 2nd-order Taylor expansion
$$\frac{1}{\delta} (f(\mathbf{x} + \delta e_i) - f(\mathbf{x} - \delta e_i)) = \frac{1}{2\delta} \left(\delta \frac{\partial f}{\partial x_i} + \frac{1}{2} \delta^2 \frac{\partial^2 f}{\partial x_i^2} + \delta \frac{\partial f}{\partial x_i} - \frac{1}{2} \delta^2 \frac{\partial^2 f}{\partial x_i^2} + O(\delta^3) \right) = \frac{\partial f}{\partial x_i} + O(\delta^2)$$

Approximate the Hessian

- Recall that for $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$ that is 2nd-order differentiable,
 $\frac{\partial f}{\partial x_i}(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}$. So

$$\frac{\partial f^2}{\partial x_j \partial x_i}(\mathbf{x}) = \frac{\partial}{\partial x_j} \left(\frac{\partial f}{\partial x_i} \right)(\mathbf{x}) \approx \frac{\left(\frac{\partial f}{\partial x_i} \right)(\mathbf{x} + \delta \mathbf{e}_j) - \left(\frac{\partial f}{\partial x_i} \right)(\mathbf{x})}{\delta}$$

- We can also compute one row of Hessian each time by

$$\frac{\partial}{\partial x_j} \left(\frac{\partial f}{\partial \mathbf{x}} \right)(\mathbf{x}) \approx \frac{\left(\frac{\partial f}{\partial \mathbf{x}} \right)(\mathbf{x} + \delta \mathbf{e}_j) - \left(\frac{\partial f}{\partial \mathbf{x}} \right)(\mathbf{x})}{\delta},$$

obtaining $\widehat{\mathbf{H}}$, which might not be symmetric. Return $\frac{1}{2} (\widehat{\mathbf{H}} + \widehat{\mathbf{H}}^\top)$ instead

- Most times (e.g., in TRM, Newton-CG), only $\nabla^2 f(\mathbf{x}) \mathbf{v}$ for certain \mathbf{v} 's needed: (see, e.g., Manopt <https://www.manopt.org/>)

$$\nabla^2 f(\mathbf{x}) \mathbf{v} \approx \frac{\nabla f(\mathbf{x} + \delta \mathbf{v}) - \nabla f(\mathbf{x})}{\delta}$$

A few words

- Can be used for sanity check of correctness of analytic gradient
- Finite-difference approximation of higher (i.e., ≥ 2)-order derivatives combined with high-order iterative methods can be very efficient (e.g., Manopt
<https://www.manopt.org/tutorial.html#costdescription>)
- Numerical stability can be an issue: truncation and round off s (finite δ ; accurate evaluation of the nominators)

Outline

Analytic differentiation

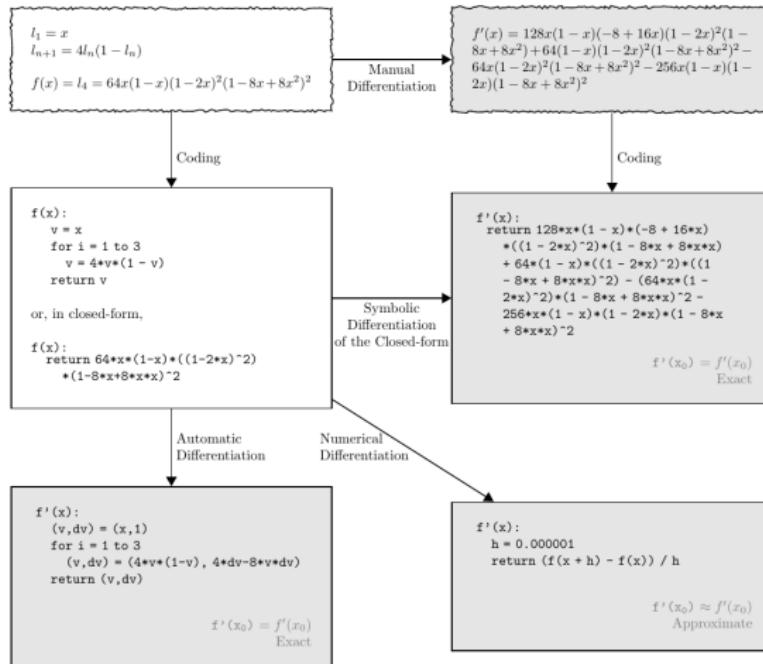
Finite-difference approximation

Automatic differentiation

Differentiable programming

Suggested reading

Four kinds of computing techniques



Credit: [Baydin et al., 2017]

Misnomer: should be **automatic numerical differentiation**

Auto differentiation in 1D

Consider a univariate function $f_k \circ f_{k-1} \circ \dots \circ f_2 \circ f_1(x) : \mathbb{R} \rightarrow \mathbb{R}$. Write $y_0 = x$, $y_1 = f_1(x)$, $y_2 = f_2(y_1)$, \dots , $y_k = f_k(y_{k-1})$, or in **computational graph** form:



Chain rule in Leibniz form:

$$\frac{\partial f}{\partial x} = \frac{\partial y_k}{\partial y_0} = \frac{\partial y_k}{\partial y_{k-1}} \frac{\partial y_{k-1}}{\partial y_{k-2}} \dots \frac{\partial y_2}{\partial y_1} \frac{\partial y_1}{\partial y_0}$$

How to evaluate the product?

- From left to right in the chain: **forward mode auto diff**
- From right to left in the chain: **backward/reverse mode auto diff**
- Hybrid: mixed mode

Forward mode in 1D



Chain rule: $\frac{df}{dx} = \frac{dy_k}{dy_0} = \left(\frac{dy_k}{dy_{k-1}} \left(\frac{dy_{k-1}}{dy_{k-2}} \left(\dots \left(\frac{dy_2}{dy_1} \left(\frac{dy_1}{dy_0} \right) \right) \right) \right) \right)$

Compute $\frac{df}{dx} \Big|_{x_0}$ in one pass, from inner to outer most parenthesis:

Input: x_0 , initialization $\frac{dy_0}{dy_0} \Big|_{x_0} = 1$

for $i = 1, \dots, k$ **do**

compute $y_i = f_i(y_{i-1})$

compute $\frac{dy_i}{dy_0} \Big|_{x_0} = \frac{dy_i}{dy_{i-1}} \Big|_{y_{i-1}} \cdot \frac{dy_{i-1}}{dy_0} \Big|_{x_0} = f'_i(y_{i-1}) \frac{dy_{i-1}}{dy_0} \Big|_{x_0}$

end for

Output: $\frac{dy_k}{dy_0} \Big|_{x_0}$

Example: For $f(x) = (x^2 + 1)^2$, calculate $\nabla f(1)$ (whiteboard)

Reverse mode in 1D



$$\text{Chain rule: } \frac{df}{dx} = \frac{df}{dy_0} = \left(\left(\left(\left(\frac{dy_k}{dy_{k-1}} \right) \frac{dy_{k-1}}{dy_{k-2}} \right) \dots \right) \frac{dy_2}{dy_1} \right) \frac{dy_1}{dy_0}$$


Compute $\frac{df}{dx} \Big|_{x_0}$ in **two** passes:

- Forward pass: calculate the y_i 's sequentially
- Backward pass: calculate the $\frac{dy_k}{dy_i} = \frac{dy_k}{dy_{i+1}} \frac{\partial dy_{i+1}}{\partial y_i}$ backward

Input: x_0 , $\frac{dy_k}{dy_k} = 1$
for $i = 1, \dots, k$ **do**
 compute $y_i = f_i(y_{i-1})$
end for // forward pass
for $i = k - 1, k - 2, \dots, 0$ **do**
 compute $\frac{dy_k}{dy_i} \Big|_{y_i} = \frac{dy_k}{dy_{i+1}} \Big|_{y_{i+1}} \cdot \frac{dy_{i+1}}{dy_i} \Big|_{y_i} = f'_{i+1}(y_i) \frac{dy_k}{dy_{i+1}} \Big|_{y_{i+1}}$
end for // backward pass
Output: $\frac{dy_k}{dy_0} \Big|_{x_0}$

Example: For $f(x) = (x^2 + 1)^2$, calculate $\nabla f(1)$ (whiteboard)

Forward vs reverse modes



- **forward mode AD:** one forward pass, compute y_i 's and $\frac{dy_i}{dy_0}$'s together
- **reverse mode AD:** one forward pass to compute y_i 's, one backward pass to compute $\frac{dy_k}{dy_i}$'s

Effectively, two different ways of grouping the multiplicative differential terms:

$$\frac{df}{dx} = \frac{df}{dy_0} = \left(\frac{dy_k}{dy_{k-1}} \left(\frac{dy_{k-1}}{dy_{k-2}} \left(\dots \left(\frac{dy_2}{dy_1} \left(\frac{dy_1}{dy_0} \right) \right) \right) \right) \right)$$

i.e., starting from the root: $\frac{dy_0}{dy_0} \mapsto \frac{dy_1}{dy_0} \mapsto \frac{dy_2}{dy_0} \mapsto \dots \mapsto \frac{dy_k}{dy_0}$

$$\frac{df}{dx} = \frac{df}{dy_0} = \left(\left(\left(\left(\left(\frac{dy_k}{dy_{k-1}} \right) \frac{dy_{k-1}}{dy_{k-2}} \right) \dots \right) \frac{dy_2}{dy_1} \right) \frac{dy_1}{dy_0} \right)$$


i.e., starting from the leaf: $\frac{dy_k}{dy_k} \mapsto \frac{dy_k}{dy_{k-1}} \mapsto \frac{dy_k}{dy_{k-2}} \mapsto \dots \mapsto \frac{dy_k}{dy_0}$

...mixed forward and reverse modes are indeed possible!

Auto differentiation in high-D

Chain Rule Let $f : \mathbb{R}^m \rightarrow \mathbb{R}^n$ and $h : \mathbb{R}^n \rightarrow \mathbb{R}^k$, and f is differentiable at \mathbf{x} and $\mathbf{z} = h(\mathbf{y})$ is differentiable at $\mathbf{y} = f(\mathbf{x})$. Then, $\mathbf{z} = h \circ f(\mathbf{x}) : \mathbb{R}^m \rightarrow \mathbb{R}^k$ is differentiable at \mathbf{x} , and

$$\mathbf{J}_{[h \circ f]}(\mathbf{x}) = \mathbf{J}_h(f(\mathbf{x})) \mathbf{J}_f(\mathbf{x}), \text{ or } \frac{\partial \mathbf{z}}{\partial \mathbf{x}} = \frac{\partial \mathbf{z}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{x}} \Leftrightarrow \frac{\partial z_j}{\partial x_i} = \sum_{\ell=1}^m \frac{\partial z_j}{\partial y_\ell} \frac{\partial y_\ell}{\partial x_i} \forall i, j$$




- Each node is a variable, as a function of all incoming variables
- If node B a child of node A , $\frac{\partial B}{\partial A}$ is the rate of change in B wrt change in A
- Traveling along a path, rates of changes should be multiplied
- Chain rule: summing up rates over all connecting paths! (e.g., x_2 to z_j as shown)

NB: this is a computational graph, not a NN

A multivariate example — forward mode

$$y = \left(\sin \frac{x_1}{x_2} + \frac{x_1}{x_2} - e^{x_2} \right) \left(\frac{x_1}{x_2} - e^{x_2} \right)$$

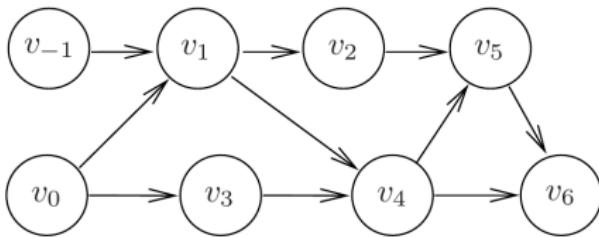


v_{-1}	$=$	x_1	$=$	1.5000
v_0	$=$	x_2	$=$	0.5000
v_1	$=$	v_{-1}/v_0	$=$	1.5000/0.5000 = 3.0000
v_2	$=$	$\sin(v_1)$	$=$	$\sin(3.0000)$ = 0.1411
v_3	$=$	$\exp(v_0)$	$=$	$\exp(0.5000)$ = 1.6487
v_4	$=$	$v_1 - v_3$	$=$	3.0000 - 1.6487 = 1.3513
v_5	$=$	$v_2 + v_4$	$=$	0.1411 + 1.3513 = 1.4924
v_6	$=$	$v_5 * v_4$	$=$	1.4924 * 1.3513 = 2.0167
y	$=$	v_6	$=$	2.0167

$v_{-1} = x_1$	$=$	1.5000
$\dot{v}_{-1} = \dot{x}_1$	$=$	1.0000
$v_0 = x_2$	$=$	0.5000
$\dot{v}_0 = \dot{x}_2$	$=$	0.0000
$v_1 = v_{-1}/v_0$	$=$	1.5000/0.5000 = 3.0000
$\dot{v}_1 = (\dot{v}_{-1} - v_1 * \dot{v}_0)/v_0 = 1.0000/0.5000$	$=$	2.0000
$v_2 = \sin(v_1)$	$=$	$\sin(3.0000)$ = 0.1411
$\dot{v}_2 = \cos(v_1) * \dot{v}_1$	$=$	-0.9900 * 2.0000 = -1.9800
$v_3 = \exp(v_0)$	$=$	$\exp(0.5000)$ = 1.6487
$\dot{v}_3 = v_3 * \dot{v}_0$	$=$	1.6487 * 0.0000 = 0.0000
$v_4 = v_1 - v_3$	$=$	3.0000 - 1.6487 = 1.3513
$\dot{v}_4 = \dot{v}_1 - \dot{v}_3$	$=$	2.0000 - 0.0000 = 2.0000
$v_5 = v_2 + v_4$	$=$	0.1411 + 1.3513 = 1.4924
$\dot{v}_5 = \dot{v}_2 + \dot{v}_4$	$=$	-1.9800 + 2.0000 = 0.0200
$v_6 = v_5 * v_4$	$=$	1.4924 * 1.3513 = 2.0167
$\dot{v}_6 = \dot{v}_5 * v_4 + v_5 * \dot{v}_4$	$=$	0.0200 * 1.3513 + 1.4924 * 2.0000 = 3.0118
$y = v_6$	$=$	2.0100
$\dot{y} = \dot{v}_6$	$=$	3.0110

- interested in $\frac{\partial}{\partial x_1}$; for each variable v_i , write $\dot{v}_i \doteq \frac{\partial v_i}{\partial x_1}$
- for each node, sum up partials over all incoming edges, e.g., $\dot{v}_4 = \frac{\partial v_4}{\partial v_1} \dot{v}_1 + \frac{\partial v_4}{\partial v_3} \dot{v}_3$
- complexity: $O(\#edges + \#nodes)$
- for $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, make n forward passes: $O(n(\#edges + \#nodes))$

A multivariate example — reverse mode



$$v_{-1} = x_1 = 1.5000$$

$$v_0 = x_2 = 0.5000$$

$$v_1 = v_{-1}/v_0 = 1.5000/0.5000 = 3.0000$$

$$v_2 = \sin(v_1) = \sin(3.0000) = 0.1411$$

$$v_3 = \exp(v_0) = \exp(0.5000) = 1.6487$$

$$v_4 = v_1 - v_3 = 3.0000 - 1.6487 = 1.3513$$

$$v_5 = v_2 + v_4 = 0.1411 + 1.3513 = 1.4924$$

$$v_6 = v_5 * v_4 = 1.4924 * 1.3513 = 2.0167$$

$$y = v_6 = 2.0167$$

$$\bar{v}_6 = \bar{y} = 1.0000$$

$$\bar{v}_5 = \bar{v}_6 * v_4 = 1.0000 * 1.3513 = 1.3513$$

$$\bar{v}_4 = \bar{v}_6 * v_5 = 1.0000 * 1.4924 = 1.4924$$

$$\bar{v}_4 = \bar{v}_4 + \bar{v}_5 = 1.4924 + 1.3513 = 2.8437$$

$$\bar{v}_2 = \bar{v}_5 = 1.3513$$

$$\bar{v}_3 = -\bar{v}_4 = -2.8437$$

$$\bar{v}_1 = \bar{v}_4 = 2.8437$$

$$\bar{v}_0 = \bar{v}_3 * v_3 = -2.8437 * 1.6487 = -4.6884$$

$$\bar{v}_1 = \bar{v}_1 + \bar{v}_2 * \cos(v_1) = 2.8437 + 1.3513 * (-0.9900) = 1.5059$$

$$\bar{v}_0 = \bar{v}_0 - \bar{v}_1 * v_1/v_0 = -4.6884 - 1.5059 * 3.000/0.5000 = -13.7239$$

$$\bar{v}_{-1} = \bar{v}_1/v_0 = 1.5059/0.5000 = 3.0118$$

$$\bar{x}_2 = \bar{v}_0 = -13.7239$$

$$\bar{x}_1 = \bar{v}_{-1} = 3.0118$$

- interested in $\frac{\partial y}{\partial}$; for each variable v_i , write $\bar{v}_i \doteq \frac{\partial y}{\partial v_i}$ (called **adjoint variable**)

- for each node, sum up partials over all outgoing edges, e.g.,

$$\bar{v}_4 = \frac{\partial v_5}{\partial v_4} \bar{v}_5 + \frac{\partial v_6}{\partial v_4} \bar{v}_6$$

- complexity:

$$O(\# \text{edges} + \# \text{nodes})$$

- for $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, make m backward passes:

$$O(m (\# \text{edges} + \# \text{nodes}))$$


example from Ch 1

of [Griewank and Walther, 2008]

Forward vs. reverse modes

For general function $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$, suppose there is no loop in the computational graph, i.e., **acyclic graph**.


Define E : set of edges ; V : set of nodes




	forward mode	reverse mode
start from	roots	leaves
end with	leaves	roots
invariants	$\dot{v}_i \doteq \frac{\partial v_i}{\partial x}$ (x —root of interest)	$\bar{v}_i \doteq \frac{\partial y}{\partial v_i}$ (y —leaf of interest)
rule	sum over incoming edges	sum over outgoing edges
complexity	$O(n E + n V)$	$O(m E + m V)$
better when	$m \gg n$	$n \gg m$

Implementation trick—tensor abstraction

Tensors: multi-dimensional arrays




Each node in the computational graph can be a tensor (scalar, vector, matrix, 3-D tensor, ...)




$$f(\mathbf{W}) = \|\mathbf{Y} - \sigma(\mathbf{W}_k \sigma(\mathbf{W}_{k-1} \sigma \dots (\mathbf{W}_1 \mathbf{X})))\|_F^2$$


computational graph for DNN



Implementation trick—tensor abstraction



computational graph for DNN



- neater computational graph
- tensor (i.e., vector) chain rules apply, often in tensor-free computation
 - * EX1: $\sigma(V_1)$ (whiteboard)
 - * EX2: V_{2k} (whiteboard)

Implementation trick—VJP

Interested in $\mathbf{J}_f(\mathbf{x})$ for $f : \mathbb{R}^n \mapsto \mathbb{R}^m$. Implement $\mathbf{v}^\top \mathbf{J}_f(\mathbf{x})$ for any $\mathbf{v} \in \mathbb{R}^m$

– Why?

- * set $\mathbf{v} = e_i$ for $i = 1, \dots, m$ to recover rows of $\mathbf{J}_f(\mathbf{x})$
- * special structures in $\mathbf{J}_f(\mathbf{x})$ (e.g., sparsity) can be exploited
- * often enough for application, e.g., calculate $\nabla(g \circ f) = (\nabla f^\top \mathbf{J}_f)^\top$ with known ∇f

– Why possible?

- * $\mathbf{v}^\top \mathbf{J}_f(\mathbf{x}) = \mathbf{J}_{\mathbf{v}^\top f}(\mathbf{x})$ so keep track of
$$\frac{\partial}{\partial v_i} (\mathbf{v}^\top f) = \sum_{k:\text{outgoing}} \frac{\partial v_k}{\partial v_i} \frac{\partial}{\partial v_k} (\mathbf{v}^\top f)$$
- * implemented in reverse-mode auto diff

```
torch.autograd.functional.vjp(func, inputs, v=None, create_graph=False, strict=False)
```

[SOURCE]

Function that computes the dot product between a vector `v` and the Jacobian of the given function at the point given by the inputs.

<https://pytorch.org/docs/stable/autograd.html>

Implementation trick—JVP

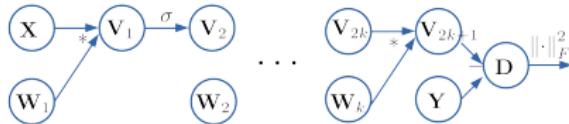
Interested in $\mathbf{J}_f(\mathbf{x})$ for $f : \mathbb{R}^n \mapsto \mathbb{R}^m$. Implement $\mathbf{J}_f(\mathbf{x})\mathbf{p}$ for any $\mathbf{p} \in \mathbb{R}^n$

- Why?
 - * set $\mathbf{p} = e_i$ for $i = 1, \dots, n$ to recover columns of $\mathbf{J}_f(\mathbf{x})$
 - * special structures in $\mathbf{J}_f(\mathbf{x})$ (e.g., sparsity) can be exploited
 - * often enough for application
- Why possible?
 - * (1) initialize partial derivatives for the input nodes as $D_{\mathbf{p}}v_{n-1} = p_1, \dots, D_{\mathbf{p}}v_0 = p_n$. (2) apply chain rule:

$$\nabla_{\mathbf{x}} v_i = \sum_{j:\text{incoming}} \frac{\partial v_i}{\partial v_j} \nabla_{\mathbf{x}} v_j \implies D_{\mathbf{p}} v_i = \sum_{j:\text{incoming}} \frac{\partial v_i}{\partial v_j} D_{\mathbf{p}} v_j$$

- * implemented in forward-mode auto diff

Putting tricks together



Basis of implementation for: Tensorflow, Pytorch, Jax, etc

<https://pytorch.org/docs/stable/autograd.html>

Jax: <https://github.com/google/jax> http://videolectures.net/deeplearning2017_johnson_automatic_differentiation/

Good to know:

- In practice, graphs are built automatically by software
- Higher-order derivatives can also be done, particularly Hessian-vector product $\nabla^2 f(x)v$ (Check out Jax!)
- Auto-diff in Tensorflow and Pytorch are specialized to DNNs , whereas Jax (in Python) is full fledged and more general
- General resources for autodiff: <http://www.autodiff.org/>, [Griewank and Walther, 2008]

Autodiff in Pytorch

Solve least squares $f(\mathbf{x}) = \frac{1}{2} \|\mathbf{y} - \mathbf{Ax}\|_2^2$ with $\nabla f(\mathbf{x}) = -\mathbf{A}^\top (\mathbf{y} - \mathbf{Ax})$

```
import torch
import matplotlib.pyplot as plt

dtype = torch.float
device = torch.device("cpu")

n, p = 500, 100

A = torch.randn(n, p, device=device, dtype=dtype)
y = torch.randn(n, device=device, dtype=dtype)

x = torch.randn(p, device=device, dtype=dtype, requires_grad=True)
step_size = 1e-4

num_step = 500
loss_vec = torch.zeros(500, device=device, dtype=dtype)

for t in range(500):
    pred = torch.matmul(A, x)
    loss = torch.pow(torch.norm(y - pred), 2)

    loss_vec[t] = loss.item()

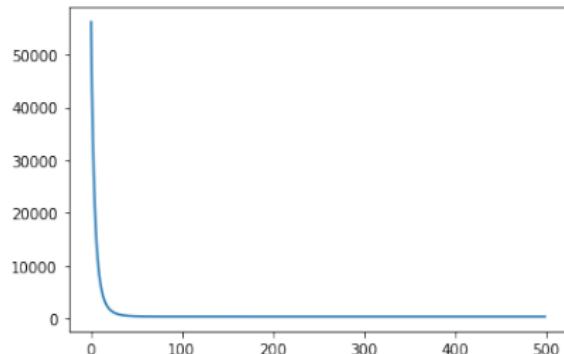
    # one line for computing the gradient
    loss.backward()

    # updates
    with torch.no_grad():
        x -= step_size*x.grad

    # zero the gradient after updating
    x.grad.zero_()

plt.plot(loss_vec.numpy())
```

loss vs. iterate



Autodiff in Pytorch

Train a shallow neural network

$$f(\mathbf{W}) = \sum_i \|\mathbf{y}_i - \mathbf{W}_2 \sigma(\mathbf{W}_1 \mathbf{x}_i)\|_2^2$$

where $\sigma(z) = \max(z, 0)$, i.e., ReLU

https://pytorch.org/tutorials/beginner/pytorch_with_examples.html

- `torch.mm`
- `torch.clamp`
- `torch.no_grad()`

Back propagation is reverse mode auto-differentiation!

Outline

Analytic differentiation

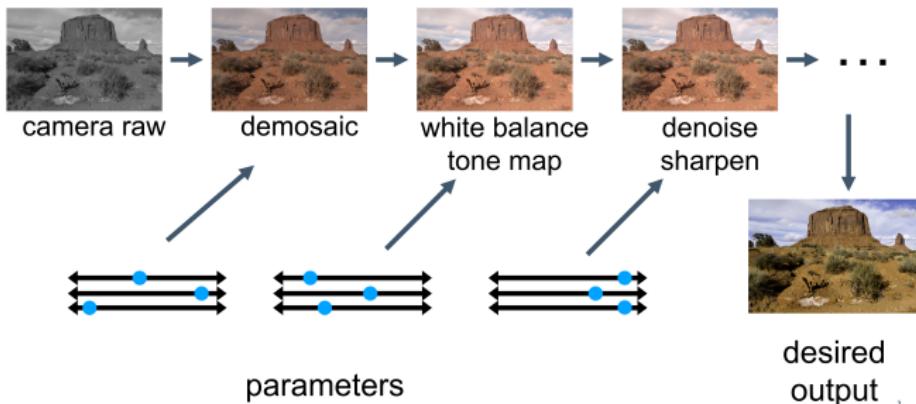
Finite-difference approximation

Automatic differentiation

Differentiable programming

Suggested reading

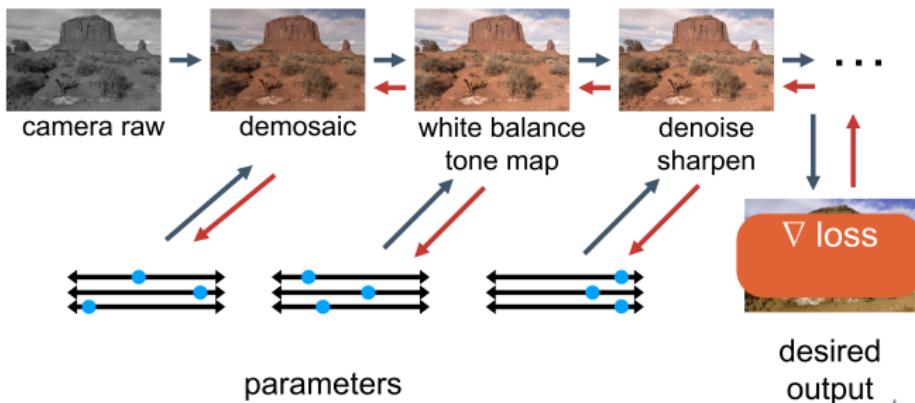
Example: image enhancement



- Each stage applies a parameterized function to the image, i.e.,
 $q_{w_k} \circ \dots \circ h_{w_3} \circ g_{w_2} \circ f_{w_1}(\mathbf{X})$ (\mathbf{X} is the camera raw)
- The parameterized functions may or may not be DNNs
- Each function may be analytic, or simply a chunk of codes dependent on the parameters
- w_i 's are the trainable parameters

Credit: https://people.csail.mit.edu/tzumao/gradient_halide/


Example: image enhancement



- the trainable parameters are learned by gradient descent based on auto-differentiation
- This is generalization of training DNNs with the classic feedforward structure to training general parameterized functions, using derivative-based methods

Credit: https://people.csail.mit.edu/tzumao/gradient_halide/

Example: control a trebuchet



<https://fluxml.ai/2019/03/05/dp-vs-rl.html>

- Given wind speed and target distance, the DNN predicts the **angle of release** and **mass of counterweight**
- Given the angle of release and mass of counterweight as initial conditions, the ODE solver calculates the actual distance by iterative methods
- We perform auto-differentiation through the iterative ODE solver and the DNN

Differential programming

Interesting resources

- Notable implementations: Swift for Tensorflow
<https://www.tensorflow.org/swift>, and Zygote in Julia
<https://github.com/FluxML/Zygote.jl>
- Flux: machine learning package based on Zygote
<https://fluxml.ai/>
- Taichi: differentiable programming language tailored to 3D computer graphics
<https://github.com/taichi-dev/taichi>

Outline

Analytic differentiation

Finite-difference approximation

Automatic differentiation

Differentiable programming

Suggested reading

Suggested reading

Autodiff in DNNs

- <http://neuralnetworksanddeeplearning.com/chap2.html>
- <https://colah.github.io/posts/2015-08-Backprop/>
- http://videolectures.net/deeplearning2017_johnson_automatic_differentiation/

Yes you should understand backprop

- <https://medium.com/@karpathy/yes-you-should-understand-backprop-e2f06eab496b>

Differentiable programming

- https://en.wikipedia.org/wiki/Differentiable_programming
- <https://fluxml.ai/2019/02/07/what-is-differentiable-programming.html>
- <https://fluxml.ai/2019/03/05/dp-vs-rl.html>

References i

- [Baydin et al., 2017] Baydin, A. G., Pearlmutter, B. A., Radul, A. A., and Siskind, J. M. (2017). **Automatic differentiation in machine learning: a survey.** *The Journal of Machine Learning Research*, 18(1):5595–5637.
- [Griewank and Walther, 2008] Griewank, A. and Walther, A. (2008). **Evaluating Derivatives: Principles and Techniques of Algorithmic Differentiation.** Society for Industrial and Applied Mathematics.