

Think Deep Learning: Overview

Ju Sun

Computer Science & Engineering

University of Minnesota, Twin Cities

September 14, 2020

Outline

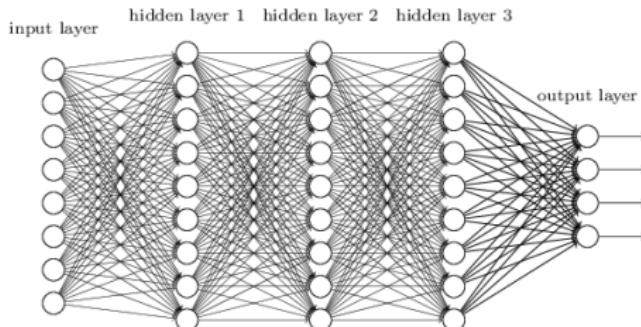
Why deep learning?

Why first principles?

Our topics

Course logistics

What is Deep Learning (DL)?



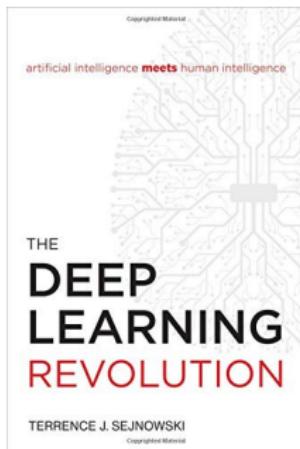
DL is about...

- Deep neural networks (DNNs)
- **Data** for training DNNs (e.g., images, videos, text sequences)
- **Methods** for training DNNs (e.g., AdaGrad, ADAM, RMSProp, Dropout)
- **Hardware** platforms for training DNNs (e.g., GPUs, TPUs, FPGAs)
- **Software** platforms for training DNNs (e.g., Tensorflow, PyTorch, MXNet)
- **Applications!** (e.g., vision, speech, NLP, imaging, physics, mathematics, finance)

Why DL?

DL leads to many things ...

Revolution: a great change in conditions, ways of working, beliefs, etc. that affects large numbers of people – *from the Oxford Dictionary*



Terrence Sejnowski (Salk Institute)

DL leads to hope

Academic breakthroughs

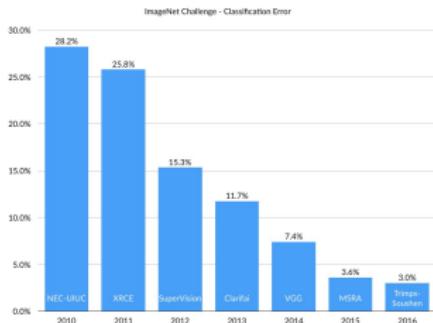
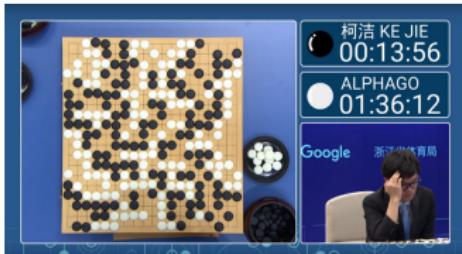
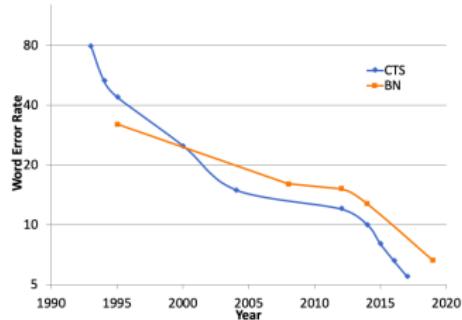


image classification



Go game (2017)



speech recognition credit: IBM



image generation credit: I. Goodfellow

DL leads to hope

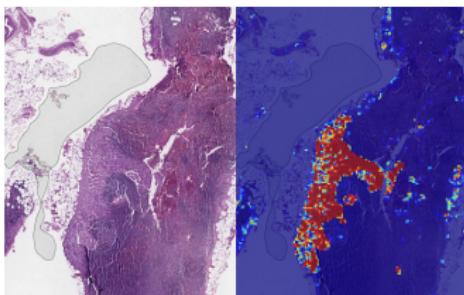
Commercial breakthroughs ...



self-driving vehicles credit: wired.com



smart-home devices credit: Amazon



healthcare credit: Google AI



robotics credit: Cornell U.

DL leads to productivity

Papers are produced at an **overwhelming** rate

Cornell University
arXiv.org > cs > cs.LG

Search... All fields Search Help | Advanced Search

Machine Learning

Authors and titles for recent submissions

- Tue, 18 Jun 2019
- Mon, 17 Jun 2019
- Fri, 14 Jun 2019
- Thu, 13 Jun 2019
- Wed, 12 Jun 2019

Total of 438 entries [438] showing 438 entries per page [newer] [older]

Tue, 18 Jun 2019

[1] arXiv:1906.07153 [pdf, other]
Adversarial attacks on Copyright Detection Systems
Panos Sachtouris, Ali Shafahi, Tom Goldstein
Subjects: Machine Learning (cs.LG); Cryptography and Security (cs.CR); Machine Learning (stat.ML)

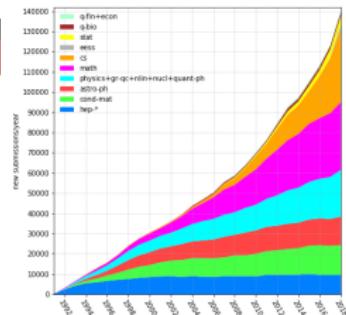
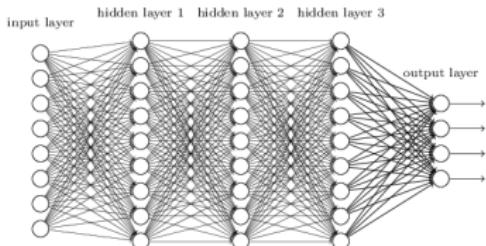


image credit: arxiv.org

$$400 \times 0.8 \times 52 / 140000 \approx 11.9\%$$

DL Supremacy!?

DL leads to fame



Turing Award 2018 credit: ACM.org

Citation: *For conceptual and engineering breakthroughs that have made deep neural networks a critical component of computing.*

DL leads to frustration

esp. for academic researchers ...

It's working amazingly well, but we don't understand why

 MORE AT SIAM

siam news

HOME | HAPPENING NOW | GET INVOLVED | RESEARCH

SIAM NEWS MAY 2017

 Research | May 01, 2017  Print

Deep, Deep Trouble

Deep Learning's Impact on Image Processing, Mathematics, and Humanity

By Michael Elad

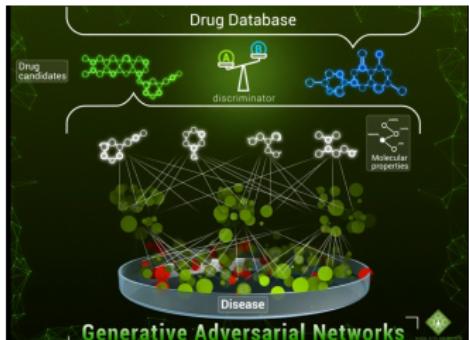
I am really confused. I keep changing my opinion on a daily basis, and I cannot seem to settle on one solid view of this puzzle. No, I am not talking about world politics or the current U.S. president, but rather something far more critical to humankind, and more specifically to our existence and work as engineers and researchers. I am talking about...deep learning.

While you might find the above statement rather bombastic and overstated, deep learning indeed raises several critical questions we must address. In the following paragraphs, I hope to expose one key conflict related to the emergence of this field, which is relevant to researchers in the image processing community.

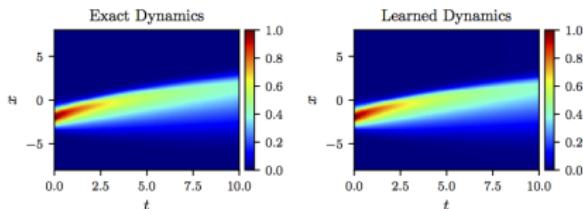
First, a few words about deep learning to put our discussion into perspective. Neural networks have been around for decades, proposing a universal learning mechanism that could, in principle, fit to any learnable data source. In



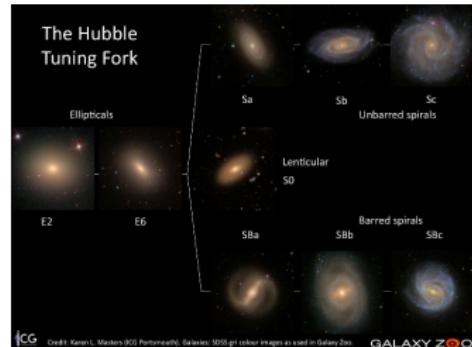
DL leads to new sciences



chemistry



applied math

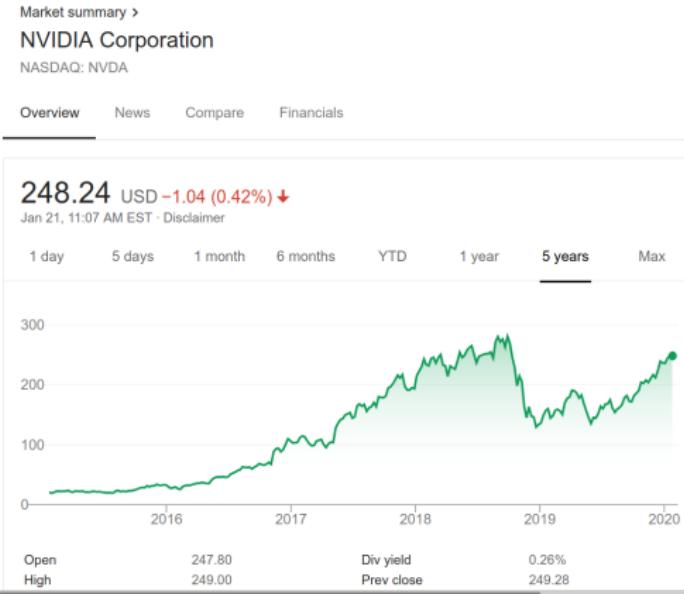


astronomy



social science

DL leads to money



- Funding
- Investment
- Job opportunities

Outline

Why deep learning?

Why first principles?

Our topics

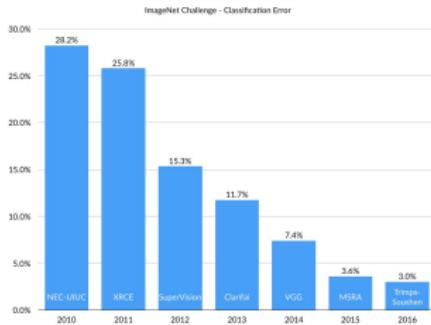
Course logistics

Why first principles?



- Tuning and optimizing for a task require basic intuitions
- **Historical lesson:** model structures in data
- **Current challenge:** move toward trustworthiness
- **Future world:** navigate uncertainties

Structures are crucial



Cat

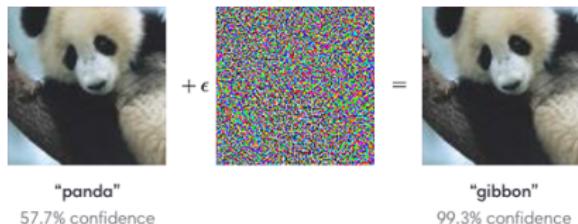


Cat

- Representation of images should ideally be **translation-invariant**.
- The 2012 breakthrough was based on modifying the classic DNNs setup to achieve translation-invariant.
- Similar success stories exist for sequences, graphs, 3D meshes.

Toward trustworthy AI

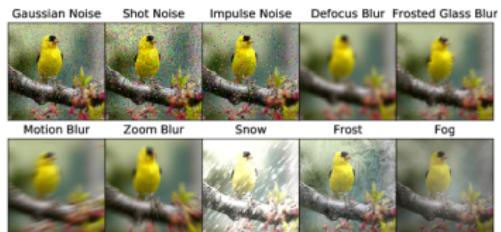
Super human-level vision?



credit: openai.com

Adversarial examples

- Trustworthiness: robustness, fairness, explainability, transparency
- We need to know first principles in order to improve and understand



credit: ImageNet-C

Natural corruptions

Future uncertainties

- New types of data (e.g., 6-D tensors)
- New hardware (e.g., better GPU memory)
- New model pipelines (e.g., network of networks, differential programming)
- New applications
- New techniques replacing DL

Outline

Why deep learning?

Why first principles?

Our topics

Course logistics

Outline of the course - I

Overview and history

Course overview (1)

Neural networks: old and new (1)

Fundamentals

Fundamental belief: universal approximation theorem (2)

Numerical optimization with math: optimization with gradient descent and beyond (2)

Numerical optimization without math: auto-differentiation and differential programming (2)

Outline of the course - II

Structured data: images and sequences

Work with images: convolutional neural networks (2)

Work with images: recognition, detection, segmentation (2)

Work with sequences: recurrent neural networks (2)

Deterministic DNN

To train or not? scattering transforms (2)

Other settings: generative/unsupervised/reinforcement learning

Learning probability distributions: generative adversarial networks (2)

Learning representation without labels: dictionary learning and autoencoders (1)

Gaming time: deep reinforcement learning (2)

Outline of tutorial/discussion sessions

Python, Numpy, and Google Cloud/Colab

Project ideas

Intro to Pytorch

Backpropagation and computational tricks

Research ideas

Outline

Why deep learning?

Why first principles?

Our topics

Course logistics

Who are we

- Instructor: **Professor Ju Sun** Email: jusun@umn.edu
Office hours: Tue/Thur 5–6pm
- TA: **Hengkang Wang** Email: wang9881@umn.edu
Office hours: Wed 4:30–6:30pm
- Guest lecturers (TBA)

Technology we use

- **Course Website:**

<https://sunju.org/teach/DL-Fall-2020/>

All course materials will be posted on the course website.

- **Communication:** **Canvas** is the preferred and most efficient way of communication. All questions and discussions go to Canvas. Send emails in exceptional situations.

For bookworms...

- **Deep Learning** by Ian Goodfellow and Yoshua Bengio and Aaron Courville. MIT Press, 2016. Online URL: <https://www.deeplearningbook.org/> (comprehensive coverage of recent developments)
- **Neural Networks and Deep Learning** by Charu Aggarwal. Springer, 2018. UMN library online access (login required): [Click here](#). (comprehensive coverage of recent developments)
- **The Deep Learning Revolution** by Terrence J. Sejnowski. MIT Press, 2018. UMN library online access (login required): [Click here](#). (account of historic developments and related fields)
- **Deep Learning with Python** by François Chollet. Online URL: <https://livebook.manning.com/book/deep-learning-with-python> (hands-on deep learning using Keras with the Tensorflow backend)
- **Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems** by Aurélien Géron (2ed). O'Reilly Media, 2019. UMN library online access (login required): [Click here](#). (hands-on machine learning, including deep learning, using Scikit-Learn and Keras)
- **Dive into Deep Learning** by Zhang et al. Live book: <https://d2l.ai/>. (comprehensive coverage & hand-ons)

How to get A(+)? 5980 Version

- 60 % homework + 40 % course project
- 5/7 homework counts. Submission to Canvas. Writing in \LaTeX (to PDF) and programming in Python 3 notebook.

Acknowledge your collaborators for each problem!

- Project based on team of 2 or 3. 5% proposal + 10% mid-term presentation + 25% final report
- Publish a paper \implies A!

How to get A(+)? 8980 Version

- 30 % homework + 40 % course project + 15% lecture teaching/scribing + 15% Short survey paper
- 3/7 homework counts. Submission to Canvas. Writing in \LaTeX (to PDF) and programming in Python 3 notebook.

Acknowledge your collaborators for each problem!

- Project based on team of 1 or 2. 5% proposal + 10% mid-term presentation + 25% final report
- Teach or scribe a 75 mins lecture session
- A short survey on topics not covered in class
- Publish a paper \implies A! test

Programming and Computing



≥ 3



≥ 2.0



≥ 1.0

Computing

- Local installation
- Google Colab: <https://colab.research.google.com/>
(Yes, it's free)
- Google Cloud (\$100 credits per student) (similarly AWS and Azure)
- Minnesota Supercomputing Institute (MSI) (class account; details forthcoming)

We're not alone

Related deep learning courses at UMN

- **Topics in Computational Vision: Deep networks** (Prof. Daniel Kersten, Department of Psychology. Focused on connection with computational neuroscience and vision)
- **Analytical Foundations of Deep Learning** (Prof. Jarvis Haupt, Department of Electrical and Computer Engineering. Focused on mathematical foundations and theories)

To learn more computational methods for large-scale optimization

- **IE5080: Optimization Models and Methods for Machine Learning** (Prof. Zhaosong Lu, Department of Industrial and Systems Engineering (ISyE))

Homework 0 later this week!

About basic **linear algebra** and **calculus** and **probability**, in
machine learning context

If you struggle too much with it

- Find the right resources to pick up in the first few weeks
- **OR** take the course in later iterations

Thank you!