

# Neural Networks: Old and New

---

**Ju Sun**

Computer Science & Engineering

University of Minnesota, Twin Cities

September 14, 2020


Start from neurons

Shallow to deep neural networks

A brief history of AI

Suggested reading


# Model of biological neurons



A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Credit: Stanford CS231N

# Model of biological neurons




Credit: Stanford CS231N

Biologically ...

- Each neuron receives signals from its **dendrites**

# Model of biological neurons




Credit: Stanford CS231N

Biologically ...

- Each neuron receives signals from its **dendrites**
- Each neuron outputs signals via its single **axon**

# Model of biological neurons




Credit: Stanford CS231N

## Biologically ...

- Each neuron receives signals from its **dendrites**
- Each neuron outputs signals via its single **axon**
- The axon branches out and connects via **synapse** to dendrites of other neurons


# Model of biological neurons



A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Credit: Stanford CS231N

# Model of biological neurons



A cartoon drawing of a biological neuron (left) and its mathematical model (right).


Credit: Stanford CS231N

Mathematically ...

- Each neuron receives  $x_i$ 's from its **dendrites**



# Model of biological neurons




A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Credit: Stanford CS231N

Mathematically ...

- Each neuron receives  $x_i$ 's from its **dendrites**
- $x_i$ 's weighted by  $w_i$ 's (synaptic strengths) and summed  $\sum_i w_i x_i$

# Model of biological neurons




A cartoon drawing of a biological neuron (left) and its mathematical model (right).

Credit: Stanford CS231N

Mathematically ...

- Each neuron receives  $x_i$ 's from its **dendrites**
- $x_i$ 's weighted by  $w_i$ 's (synaptic strengths) and summed  $\sum_i w_i x_i$
- The neuron fires only when the combined signal is above a certain threshold:  $\sum_i w_i x_i + b$

# Model of biological neurons




Credit: Stanford CS231N

Mathematically ...

- Each neuron receives  $x_i$ 's from its **dendrites**
- $x_i$ 's weighted by  $w_i$ 's (synaptic strengths) and summed  $\sum_i w_i x_i$
- The neuron fires only when the combined signal is above a certain threshold:  $\sum_i w_i x_i + b$
- Fire rate is modeled by an **activation function**  $f$ , i.e., outputting  $f(\sum_i w_i x_i + b)$


## Brain neural networks



Credit: Max Pixel


# Artificial neural networks

## Brain neural networks




Credit: Max Pixel

## Artificial neural networks




# Artificial neural networks

## Brain neural networks



Credit: Max Pixel


## Artificial neural networks



Why called **artificial**?


# Artificial neural networks

## Brain neural networks



Credit: Max Pixel

## Artificial neural networks




## Why called **artificial**?

- (Over-)simplification on neural level
- (Over-)simplification on connection level


# Artificial neural networks

## Brain neural networks



Credit: Max Pixel

## Artificial neural networks



## Why called **artificial**?

- (Over-)simplification on neural level
- (Over-)simplification on connection level

In this course, neural networks are always artificial.




Start from neurons

Shallow to deep neural networks


A brief history of AI

Suggested reading

# Artificial neurons




# Artificial neurons



$$f\left(\sum_i w_i x_i + b\right) = f(\mathbf{w}^\top \mathbf{x} + b)$$


# Artificial neurons



$$f\left(\sum_i w_i x_i + b\right) = f(\mathbf{w}^\top \mathbf{x} + b)$$

**We shall use  $\sigma$  instead of  $f$   
henceforth.**

# Artificial neurons




## Examples of activation function $\sigma$

$$f\left(\sum_i w_i x_i + b\right) = f(w^\top x + b)$$

We shall use  $\sigma$  instead of  $f$   
henceforth.


# Artificial neurons



$$f\left(\sum_i w_i x_i + b\right) = f(w^\top x + b)$$


We shall use  $\sigma$  instead of  $f$   
henceforth.

## Examples of activation function $\sigma$




Sigmoid

$$\phi(z) = \frac{1}{1 + e^{-z}}$$




Hyperbolic Tangent

$$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$



Rectified Linear

$$\phi(z) = \begin{cases} 0 & \text{if } z < 0 \\ z & \text{if } z \geq 0 \end{cases}$$



Radial Basis Function

$$\phi(z, c) = e^{-(c\|z - c\|)^2}$$


Credit: [Hughes and Correll, 2016]

One neuron:  $\sigma(\boldsymbol{w}^\top \boldsymbol{x} + b)$

# Neural networks

One neuron:  $\sigma(w^T x + b)$

Neural networks (NN): **structured** organization of artificial neurons






# Neural networks

One neuron:  $\sigma(w^T x + b)$

Neural networks (NN): **structured** organization of artificial neurons




$w$ 's and  $b$ 's are unknown and need to be learned

# Neural networks

One neuron:  $\sigma(w^T x + b)$

Neural networks (NN): **structured** organization of artificial neurons



$w$ 's and  $b$ 's are unknown and need to be learned

Many models in machine learning **are** neural networks

### Supervised Learning

- Gather training data  $(x_1, y_1), \dots, (x_n, y_n)$

## Supervised Learning

- Gather training data  $(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_n, \mathbf{y}_n)$
- Choose a family of functions, e.g.,  $\mathcal{H}$ , so that there is  $f \in \mathcal{H}$  to ensure  $\mathbf{y}_i \approx f(\mathbf{x}_i)$  for all  $i$

## Supervised Learning

- Gather training data  $(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_n, \mathbf{y}_n)$
- Choose a family of functions, e.g.,  $\mathcal{H}$ , so that there is  $f \in \mathcal{H}$  to ensure  $\mathbf{y}_i \approx f(\mathbf{x}_i)$  for all  $i$
- Set up a loss function  $\ell$  to measure the approximation quality

## Supervised Learning

- Gather training data  $(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_n, \mathbf{y}_n)$
- Choose a family of functions, e.g.,  $\mathcal{H}$ , so that there is  $f \in \mathcal{H}$  to ensure  $\mathbf{y}_i \approx f(\mathbf{x}_i)$  for all  $i$
- Set up a loss function  $\ell$  to measure the approximation quality
- Find an  $f \in \mathcal{H}$  to minimize the average loss

$$\min_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell(\mathbf{y}_i, f(\mathbf{x}_i))$$

## Supervised Learning

- Gather training data  $(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_n, \mathbf{y}_n)$
- Choose a family of functions, e.g.,  $\mathcal{H}$ , so that there is  $f \in \mathcal{H}$  to ensure  $\mathbf{y}_i \approx f(\mathbf{x}_i)$  for all  $i$
- Set up a loss function  $\ell$  to measure the approximation quality
- Find an  $f \in \mathcal{H}$  to minimize the average loss

$$\min_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell(\mathbf{y}_i, f(\mathbf{x}_i))$$

... known as **empirical risk minimization** (ERM) framework in learning theory

## A typical setup

### Supervised Learning from NN viewpoint

- Gather training data  $(x_1, y_1), \dots, (x_n, y_n)$



# A typical setup

## Supervised Learning from NN viewpoint

- Gather training data  $(x_1, y_1), \dots, (x_n, y_n)$
- Choose a NN with  $k$  neurons, so that there is a group of weights, e.g.,  $(w_1, \dots, w_k, b_1, \dots, b_k)$ , to ensure

$$y_i \approx \{\text{NN}(w_1, \dots, w_k, b_1, \dots, b_k)\}(x_i) \quad \forall i$$

# A typical setup

## Supervised Learning from NN viewpoint

- Gather training data  $(x_1, y_1), \dots, (x_n, y_n)$
- Choose a NN with  $k$  neurons, so that there is a group of weights, e.g.,  $(w_1, \dots, w_k, b_1, \dots, b_k)$ , to ensure

$$y_i \approx \{\text{NN}(w_1, \dots, w_k, b_1, \dots, b_k)\}(x_i) \quad \forall i$$

- Set up a loss function  $\ell$  to measure the approximation quality

# A typical setup

## Supervised Learning from NN viewpoint


- Gather training data  $(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_n, \mathbf{y}_n)$
- Choose a NN with  $k$  neurons, so that there is a group of weights, e.g.,  $(\mathbf{w}_1, \dots, \mathbf{w}_k, b_1, \dots, b_k)$ , to ensure

$$\mathbf{y}_i \approx \{\text{NN}(\mathbf{w}_1, \dots, \mathbf{w}_k, b_1, \dots, b_k)\}(\mathbf{x}_i) \quad \forall i$$

- Set up a loss function  $\ell$  to measure the approximation quality
- Find weights  $(\mathbf{w}_1, \dots, \mathbf{w}_k, b_1, \dots, b_k)$  to minimize the average loss


$$\min_{\mathbf{w}'_s, b'_s} \frac{1}{n} \sum_{i=1}^n \ell[\mathbf{y}_i, \{\text{NN}(\mathbf{w}_1, \dots, \mathbf{w}_k, b_1, \dots, b_k)\}(\mathbf{x}_i)]$$

# Linear regression



Credit: D2L


# Linear regression



Credit: D2L

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$


# Linear regression



Credit: D2L

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$
- Model:  $y_i \approx \mathbf{w}^\top \mathbf{x}_i + b$


# Linear regression



Credit: D2L

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$
- Model:  $y_i \approx \mathbf{w}^\top \mathbf{x}_i + b$
- Loss:  $\|y - \hat{y}\|_2^2$

# Linear regression




Credit: D2L

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$
- Model:  $y_i \approx \mathbf{w}^\top \mathbf{x}_i + b$
- Loss:  $\|y - \hat{y}\|_2^2$
- Optimization:

$$\min_{\mathbf{w}, b} \frac{1}{n} \sum_{i=1}^n \|y_i - (\mathbf{w}^\top \mathbf{x}_i + b)\|_2^2$$



# Linear regression




Credit: D2L

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n), \mathbf{x}_i \in \mathbb{R}^d$
- Model:  $y_i \approx \mathbf{w}^\top \mathbf{x}_i + b$
- Loss:  $\|y - \hat{y}\|_2^2$
- Optimization:

$$\min_{\mathbf{w}, b} \frac{1}{n} \sum_{i=1}^n \|y_i - (\mathbf{w}^\top \mathbf{x}_i + b)\|_2^2$$

Output layer

Input layer



Credit: D2L

$\sigma$  is the identity function



**Frank Rosenblatt**

(1928–1971)



**Frank Rosenblatt**

(1928–1971)

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  
 $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{+1, -1\}$


# Perceptron



**Frank Rosenblatt**

(1928–1971)

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  
 $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{+1, -1\}$
- Model:  $y_i \approx \sigma(\mathbf{w}^\top \mathbf{x}_i + b)$ ,  $\sigma$  sign function




# Perceptron



**Frank Rosenblatt**

(1928–1971)

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  
 $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{+1, -1\}$
- Model:  $y_i \approx \sigma(\mathbf{w}^\top \mathbf{x}_i + b)$ ,  $\sigma$  sign function



- Loss:  $\mathbf{1}\{y \neq \hat{y}\}$


# Perceptron



**Frank Rosenblatt**

(1928–1971)

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  
 $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{+1, -1\}$
- Model:  $y_i \approx \sigma(\mathbf{w}^\top \mathbf{x}_i + b)$ ,  $\sigma$  sign function




- Loss:  $\mathbf{1}\{y \neq \hat{y}\}$
- Optimization:

$$\min_{\mathbf{w}, b} \frac{1}{n} \sum_{i=1}^n \mathbf{1}\{y_i \neq \sigma(\mathbf{w}^\top \mathbf{x}_i + b)\}$$


# Perceptron

Perceptron is a single artificial neuron for **binary classification**



# Perceptron

Perceptron is a single artificial neuron for **binary classification**




dominated early AI (50's – 70's)




# Perceptron

Perceptron is a single artificial neuron for **binary classification**



dominated early AI (50's – 70's)

**Logistic regression** is similar but with **sigmoid** activation



## Softmax regression

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{L_1, \dots, L_p\}$ , i.e., multiclass classification problem

# Softmax regression

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{L_1, \dots, L_p\}$ , i.e., multiclass classification problem
- Data preprocessing: labels into vectors via **one-hot encoding**

$$L_k \implies \underbrace{[0, \dots, 0]}_{k-1 \text{ 0's}}, 1, \underbrace{[0, \dots, 0]}_{n-k \text{ 0's}}^\top$$

So:  $y_i \implies \mathbf{y}_i$

# Softmax regression

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{L_1, \dots, L_p\}$ , i.e., multiclass classification problem
- Data preprocessing: labels into vectors via **one-hot encoding**

$$L_k \implies [\underbrace{0, \dots, 0}_{k-1 \text{ 0's}}, 1, \underbrace{0, \dots, 0}_{n-k \text{ 0's}}]^\top$$

So:  $y_i \implies \mathbf{y}_i$

- Model:  $\mathbf{y}_i \approx \sigma(\mathbf{W}^\top \mathbf{x}_i + \mathbf{b})$ , here  $\sigma$  is the softmax function

# Softmax regression

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{L_1, \dots, L_p\}$ , i.e., multiclass classification problem
- Data preprocessing: labels into vectors via **one-hot encoding**

$$L_k \implies \underbrace{[0, \dots, 0]}_{k-1 \text{ 0's}}, \underbrace{[1, 0, \dots, 0]}_{n-k \text{ 0's}}^\top$$

So:  $y_i \implies \mathbf{y}_i$

- Model:  $\mathbf{y}_i \approx \sigma(\mathbf{W}^\top \mathbf{x}_i + \mathbf{b})$ , here  $\sigma$  is the softmax function (**maps vectors to vectors**): for  $\mathbf{z} \in \mathbb{R}^p$ ,

$$\mathbf{z} \mapsto \left[ \frac{e^{z_1}}{\sum_j e^{z_j}}, \dots, \frac{e^{z_p}}{\sum_j e^{z_j}} \right]^\top.$$

# Softmax regression

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{L_1, \dots, L_p\}$ , i.e., multiclass classification problem
- Data preprocessing: labels into vectors via **one-hot encoding**

$$L_k \implies \underbrace{[0, \dots, 0]}_{k-1 \text{ 0's}}, \underbrace{[1, 0, \dots, 0]}_{n-k \text{ 0's}}^\top$$

So:  $y_i \implies \mathbf{y}_i$

- Model:  $\mathbf{y}_i \approx \sigma(\mathbf{W}^\top \mathbf{x}_i + \mathbf{b})$ , here  $\sigma$  is the softmax function (**maps vectors to vectors**): for  $\mathbf{z} \in \mathbb{R}^p$ ,

$$\mathbf{z} \mapsto \left[ \frac{e^{z_1}}{\sum_j e^{z_j}}, \dots, \frac{e^{z_p}}{\sum_j e^{z_j}} \right]^\top.$$

- Loss: **cross-entropy loss**  $-\sum_j y_j \log \hat{y}_j$

# Softmax regression

- Data:  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$ ,  $\mathbf{x}_i \in \mathbb{R}^d$ ,  $y_i \in \{L_1, \dots, L_p\}$ , i.e., multiclass classification problem
- Data preprocessing: labels into vectors via **one-hot encoding**

$$L_k \implies \underbrace{[0, \dots, 0]}_{k-1 \text{ 0's}}, \underbrace{[1, 0, \dots, 0]}_{n-k \text{ 0's}}^\top$$

So:  $y_i \implies \mathbf{y}_i$


- Model:  $\mathbf{y}_i \approx \sigma(\mathbf{W}^\top \mathbf{x}_i + \mathbf{b})$ , here  $\sigma$  is the softmax function (**maps vectors to vectors**): for  $\mathbf{z} \in \mathbb{R}^p$ ,

$$\mathbf{z} \mapsto \left[ \frac{e^{z_1}}{\sum_j e^{z_j}}, \dots, \frac{e^{z_p}}{\sum_j e^{z_j}} \right]^\top.$$

- Loss: **cross-entropy loss**  $-\sum_j y_j \log \hat{y}_j$
- Optimization ...


# Softmax regression

... for multiclass classification






# Multilayer perceptrons



Credit: D2L

$$\text{Model: } y_i \approx \sigma_2(W_2^T \sigma_1(W_1^T x + b_1) + b_2)$$

# Multilayer perceptrons




Credit: D2L

$$\text{Model: } y_i \approx \sigma_2(W_2^T \sigma_1(W_1^T x + b_1) + b_2)$$

Also called **feedforward networks** or **fully-connected networks**

# Multilayer perceptrons



Credit: D2L

$$\text{Model: } y_i \approx \sigma_2(W_2^T \sigma_1(W_1^T x + b_1) + b_2)$$

Also called **feedforward networks** or **fully-connected networks**

Modern NNs: many hidden layers (deep), refined connection structure and/or activations

## They're all (shallow) NNs

- Linear regression
- Perception and Logistic regression
- Softmax regression
- Multilayer perceptron (feedforward NNs)

## They're all (shallow) NNs

- Linear regression
- Perception and Logistic regression
- Softmax regression
- Multilayer perceptron (feedforward NNs)
- Support vector machines (SVM)
- PCA (autoencoder)
- Matrix factorization

see, e.g., Chapter 2 of [[Aggarwal, 2018](#)].

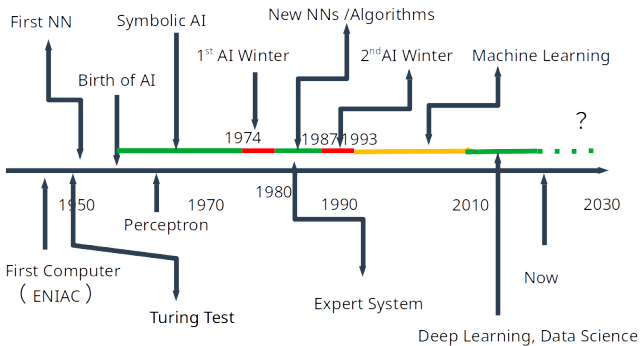
Start from neurons

Shallow to deep neural networks


A brief history of AI

Suggested reading

# Birth of AI




# Birth of AI



- Crucial precursors: first computer, Turing test




# Birth of AI




- Crucial precursors: first computer, Turing test
- 1956: Dartmouth Artificial Intelligence Summer Research Project — Birth of AI

# Turing test




Turing Test




Alan Turing (1912–1954)

# First golden age




# First golden age




Symbolic AI: based on rules and logic


# First golden age




Symbolic AI: based on rules and logic



# First golden age




## Symbolic AI: based on rules and logic





rules for recognizing dogs?

# First AI winter




# First AI winter



Gartner hype cycle




# Perceptron




invented 1962

# Perceptron



invented 1962



written in 1969, end of  
Perceptron era



Marvin Minsky (1927–2016)

# Birth of computer vision

MASSACHUSETTS INSTITUTE OF TECHNOLOGY  
PROJECT MAC

Artificial Intelligence Group  
Vision Memo. No. 100.


July 7, 1966

## THE SUMMER VISION PROJECT

Seymour Papert


The summer vision project is an attempt to use our summer workers effectively in the construction of a significant part of a visual system. The particular task was chosen partly because it can be segmented into sub-problems which will allow individuals to work independently and yet participate in the construction of a system complex enough to be a real landmark in the development of "pattern recognition".

1966




around 1980


# Second golden age



# Second golden age



## expert system




Can we build comprehensive knowledge bases and know all rules?


Key ingredients of DL have been in place for 25-30 years:

Landmark	Emblem	Epoch
Neocognitron	Fukushima	1980
CNN	Le Cun	mid 1980s'
Backprop	Hinton	mid 1980's
SGD	Le Cun, Bengio etc	mid 1990's
Various	Schmidhuber	mid 1980's
<i>CTF</i>	<i>DARPA etc</i>	<i>mid 1980's</i>


# After 2nd AI winter



# After 2nd AI winter



Machine learning takes over ...





Starting 1990's

Support vector machines (SVM)

Adaboost

Decision trees and random forests

Deep learning

...

Start from neurons

Shallow to deep neural networks

A brief history of AI

Suggested reading

## Suggested reading

- Chap 2, Neural Networks and Deep Learning.
- Chap 3–4, Dive into Deep Learning.
- Chap 1, Deep Learning with Python.

- [Aggarwal, 2018] Aggarwal, C. C. (2018). **Neural Networks and Deep Learning**. Springer International Publishing.
- [Hughes and Correll, 2016] Hughes, D. and Correll, N. (2016). **Distributed machine learning in materials that couple sensing, actuation, computation and communication**. *arXiv:1606.03508*.