

# Transformers, Large Language Models (LLMs), and Foundation Models

Ju Sun  
Computer Science & Engineering

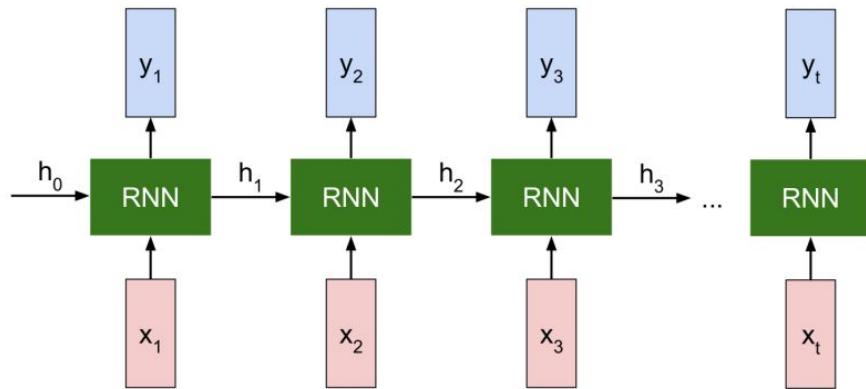
Apr 08, 2025



UNIVERSITY OF MINNESOTA  
Driven to Discover<sup>SM</sup>

# Quick recap

## RNN: model sequences



(Credit: Stanford CS231N)

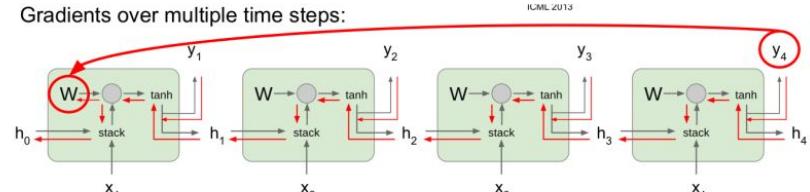
$$\mathbf{h}_t = \tanh(\mathbf{W}_h \mathbf{h}_{t-1} + \mathbf{W}_x \mathbf{x}_t)$$

$$\mathbf{y}_t = \mathbf{V}_y \mathbf{h}_t$$

$\mathbf{W}_h$ ,  $\mathbf{W}_x$  and  $\mathbf{V}_y$  are shared across the sequence

## Vanishing/exploding gradient issue

Gradients over multiple time steps:



(Credit: Stanford CS231N)

$$\begin{aligned}\frac{\partial L_t}{\partial W} &= \frac{\partial L_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_{t-1}} \dots \frac{\partial \mathbf{h}_1}{\partial \mathbf{W}} = \frac{\partial L_t}{\partial \mathbf{h}_t} \left( \prod_{k=2}^t \frac{\partial \mathbf{h}_k}{\partial \mathbf{h}_{k-1}} \right) \frac{\partial \mathbf{h}_1}{\partial \mathbf{W}} \\ &= \frac{\partial L_t}{\partial \mathbf{h}_t} \left( \prod_{k=2}^t \text{diag}(\tanh'(\mathbf{W}_h \mathbf{h}_{k-1} + \mathbf{W}_x \mathbf{x}_k)) \mathbf{W}_h \right) \frac{\partial \mathbf{h}_1}{\partial \mathbf{W}}\end{aligned}$$

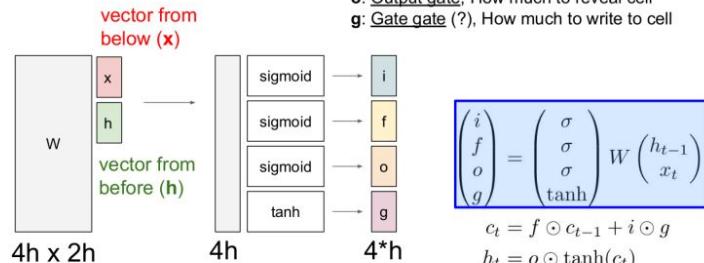
- \* when  $\|\mathbf{W}_h\| > 1$ , gradient **explodes** if  $t$  large
- \* when  $\|\mathbf{W}_h\| < 1$ , gradient **vaniishes** if  $t$  large

$$\begin{aligned}&\left\| \prod_{k=2}^t \text{diag}(\tanh'(\mathbf{W}_h \mathbf{h}_{k-1} + \mathbf{W}_x \mathbf{x}_k)) \mathbf{W}_h \right\| \\ &\leq \prod_{k=2}^t \|\text{diag}(\tanh'(\mathbf{W}_h \mathbf{h}_{k-1} + \mathbf{W}_x \mathbf{x}_k))\| \|\mathbf{W}_h\| \\ &\leq \prod_{k=2}^t \|\text{diag}(\tanh'(\mathbf{W}_h \mathbf{h}_{k-1} + \mathbf{W}_x \mathbf{x}_k))\| \|\mathbf{W}_h\|^{t-1}\end{aligned}$$

# Gated RNNs

Long Short Term Memory (LSTM)

[Hochreiter et al., 1997]



(Credit: Stanford CS231N)

$u$ : update gate, control state update

$r$ : reset gate, control how previous state affects new content

$g$ : new content

**Gated recurrent unit (GRU)**

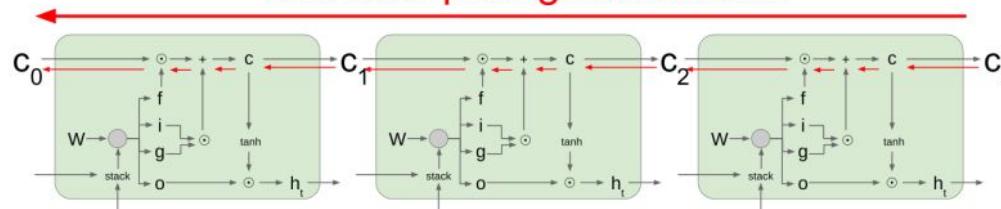
$$\begin{bmatrix} u \\ r \\ g \end{bmatrix} = \begin{bmatrix} \sigma \\ \sigma \\ \sigma \end{bmatrix} \left( \mathbf{W} \begin{bmatrix} h_{t-1} \\ x_t \end{bmatrix} \right)$$

$$g = \tanh(\mathbf{W}_h(r \odot h_{t-1}) + \mathbf{W}_x x_t + b_g)$$

$$h_t = u \odot h_{t-1} + (1 - u) \odot g$$

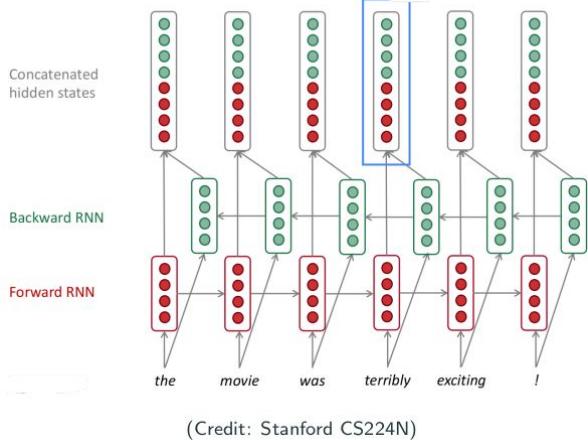
$f, i, o$  are merged

Uninterrupted gradient flow!

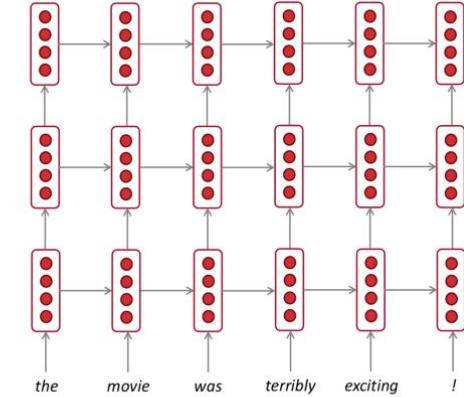


(Credit: Stanford CS231N)

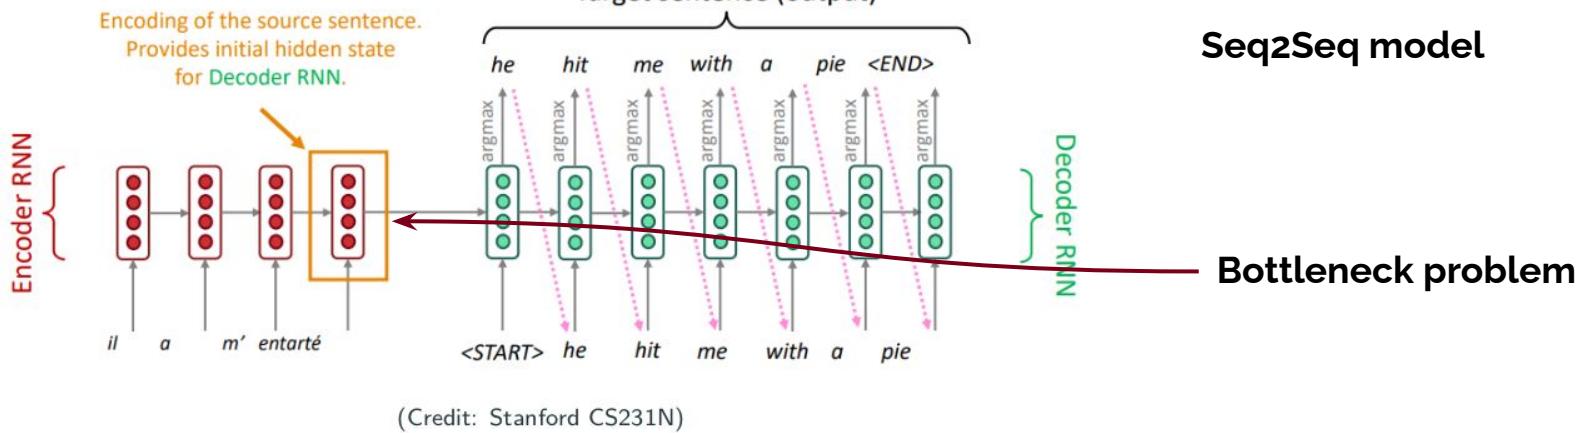
# Modern RNNs



## Bidirectional RNN

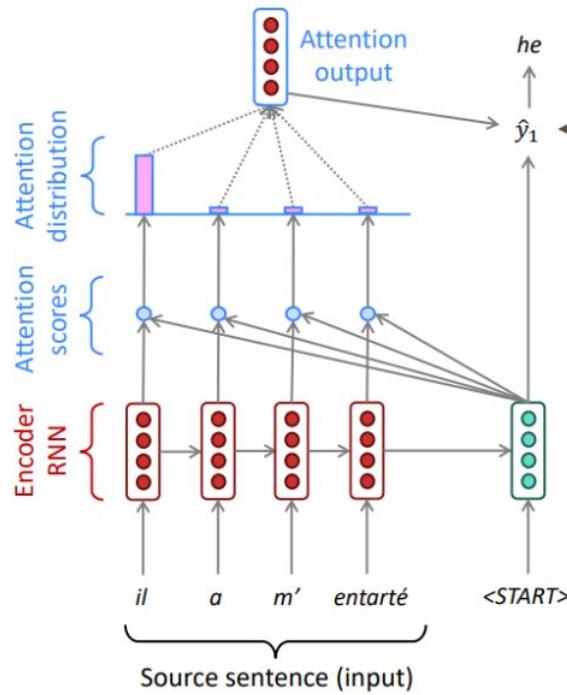


## Deep RNN



## Bottleneck problem

# Attention mechanism



(Credit: Stanford CS231N)

**Input:** source vectors  $s_1, \dots, s_N \in \mathbb{R}^h$ , and target vector  $t$

**Output:** weighted summation

$$\sum_{j=1}^N w_j s_j \quad \text{where } w_j = \text{similarity}(s_j, t)$$

Many possibilities:

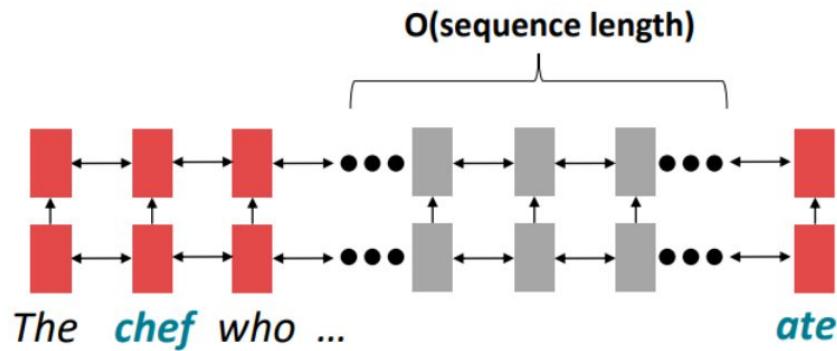
## Attention scores

- dot-product attention:  $\widehat{w}_j = \langle s_j, t \rangle$  (Is it better to normalize this or rescale it by the dimension factor? )
- multiplicative attention:  $\widehat{w}_j = \langle s_j, Wt \rangle$
- “additive attention”:  $\widehat{w}_j = v^\top \sigma(W_1 s_j + W_2 t)$

The actual weights are attention scores passed through **softmax**

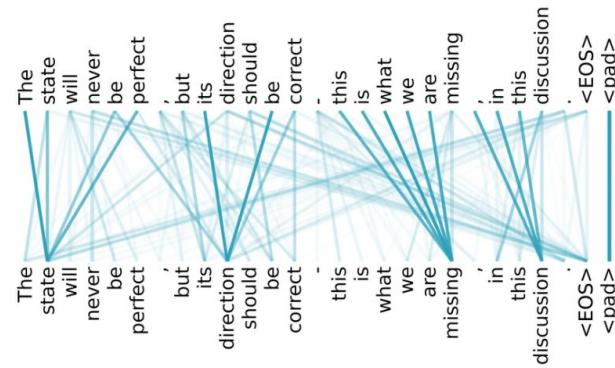
$$w_j = \frac{\exp(\widehat{w}_j)}{\sum_k \exp(\widehat{w}_k)}$$

# Self-attention



## RNN

- Long interaction distance
- Resistant to parallelization



## Self-attention

- $O(1)$  interaction distance
- Highly parallelizable

**Each token gets a selective summary of information from all others**

# Self-attention

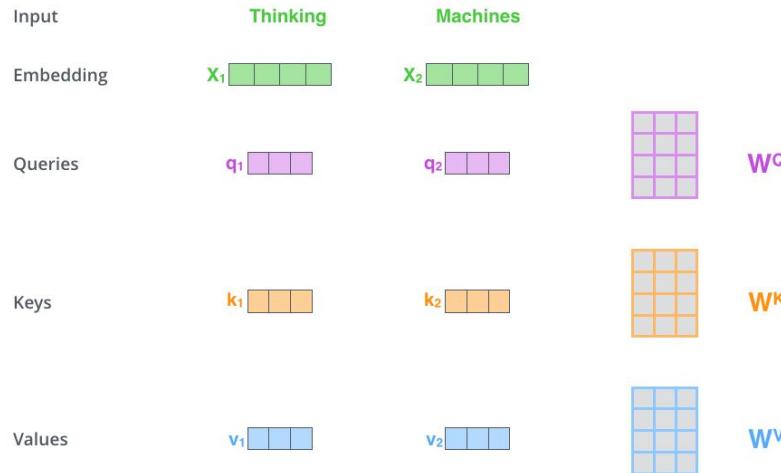


Image credit: <https://jalammar.github.io/illustrated-transformer/>

- Each word now encoded as (query, key, value) triple
- For an input  $x_i$ , we have:

$$q_i = (W^Q)^\top x_i, \quad k_i = (W^K)^\top x_i, \quad v_i = (W^V)^\top x_i$$

- Calculate attention scores between query and all keys:  $e_{ij} = \langle q_i, k_j \rangle$
- softmax normalization  $w_{ij} = \exp(e_{ij}) / \sum_k \exp(e_{ik})$
- output the weighted sum of values  $\sum_j w_{ij} v_j$

In matrix notation

- Compute queries, keys, and values

$$Q = XW^Q, \quad K = XW^K, \quad V = XW^V$$

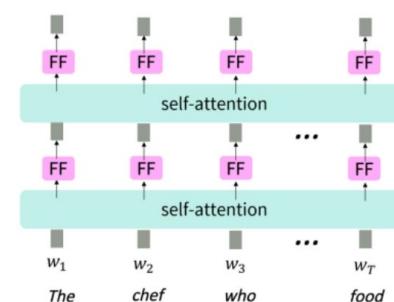
- Calculate attention scores between query and all keys:  $E = QK^\top$
- softmax normalization  $A = \text{softmax}(E)$
- output the weighted sum of values  $AV$

$$\text{output} = \text{softmax}(QK^\top)V$$

**Question:** why we need both query and key?

Equation for Feed Forward Layer

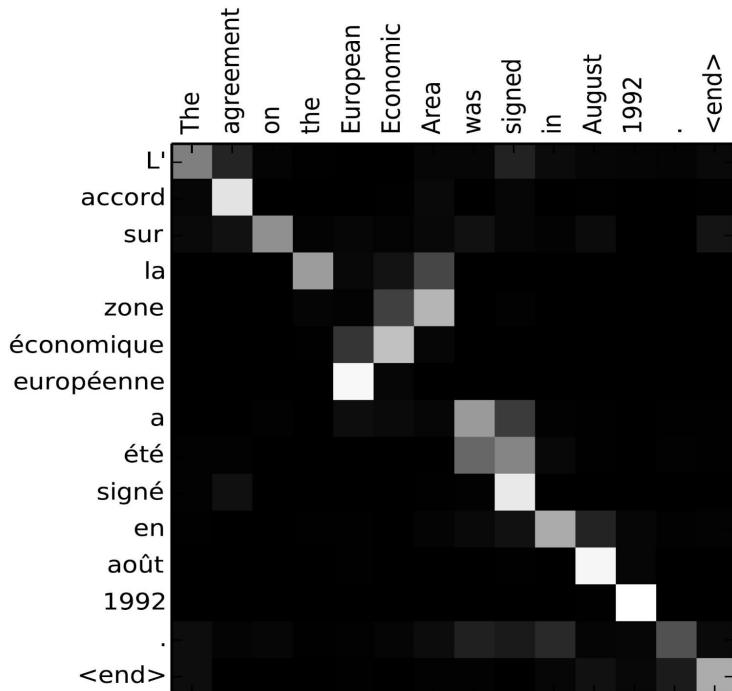
$$\begin{aligned} m_i &= \text{MLP}(\text{output}_i) \\ &= W_2 * \text{ReLU}(W_1 \times \text{output}_i + b_1) + b_2 \end{aligned}$$



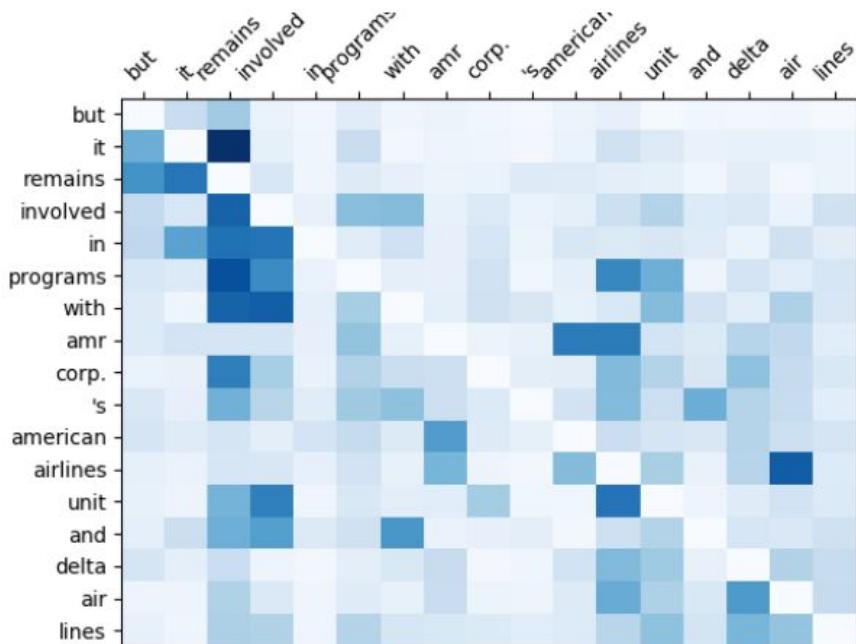
Adding in nonlinearity!

First step toward Transformers!

# Attention matrices—visualizing correlations



General attention



Self-attention

# Transformers

# Transformers

## Attention Is All You Need

**Ashish Vaswani\***  
Google Brain  
avaswani@google.com

**Noam Shazeer\***  
Google Brain  
noam@google.com

**Niki Parmar\***  
Google Research  
nikip@google.com

**Jakob Uszkoreit\***  
Google Research  
usz@google.com

**Llion Jones\***  
Google Research  
llion@google.com

**Aidan N. Gomez\*** †  
University of Toronto  
aidan@cs.toronto.edu

**Lukasz Kaiser\***  
Google Brain  
lukaszkaiser@google.com

**Illia Polosukhin\*** ‡  
illia.polosukhin@gmail.com

NIPS 2017; <https://arxiv.org/abs/1706.03762>

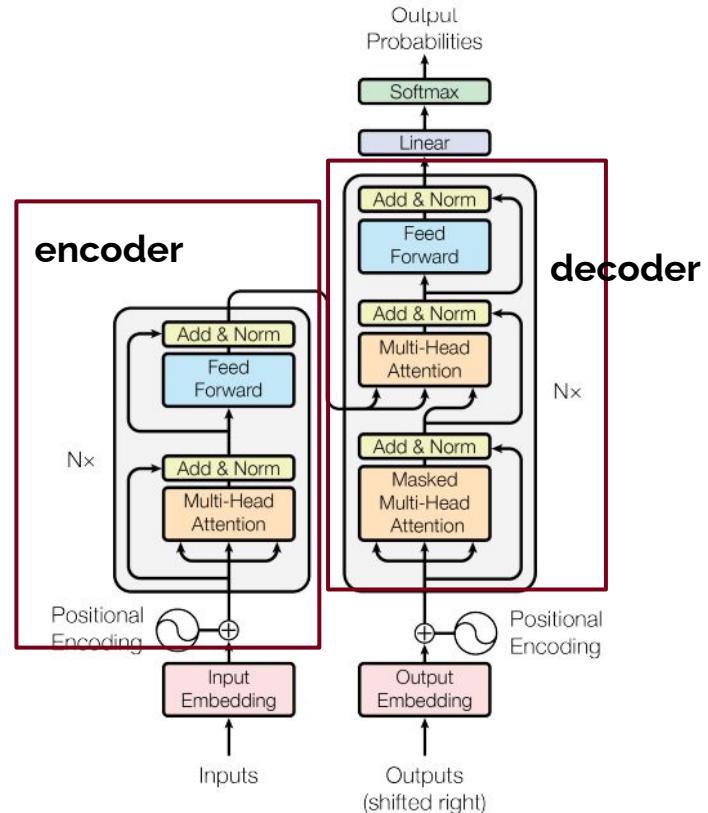
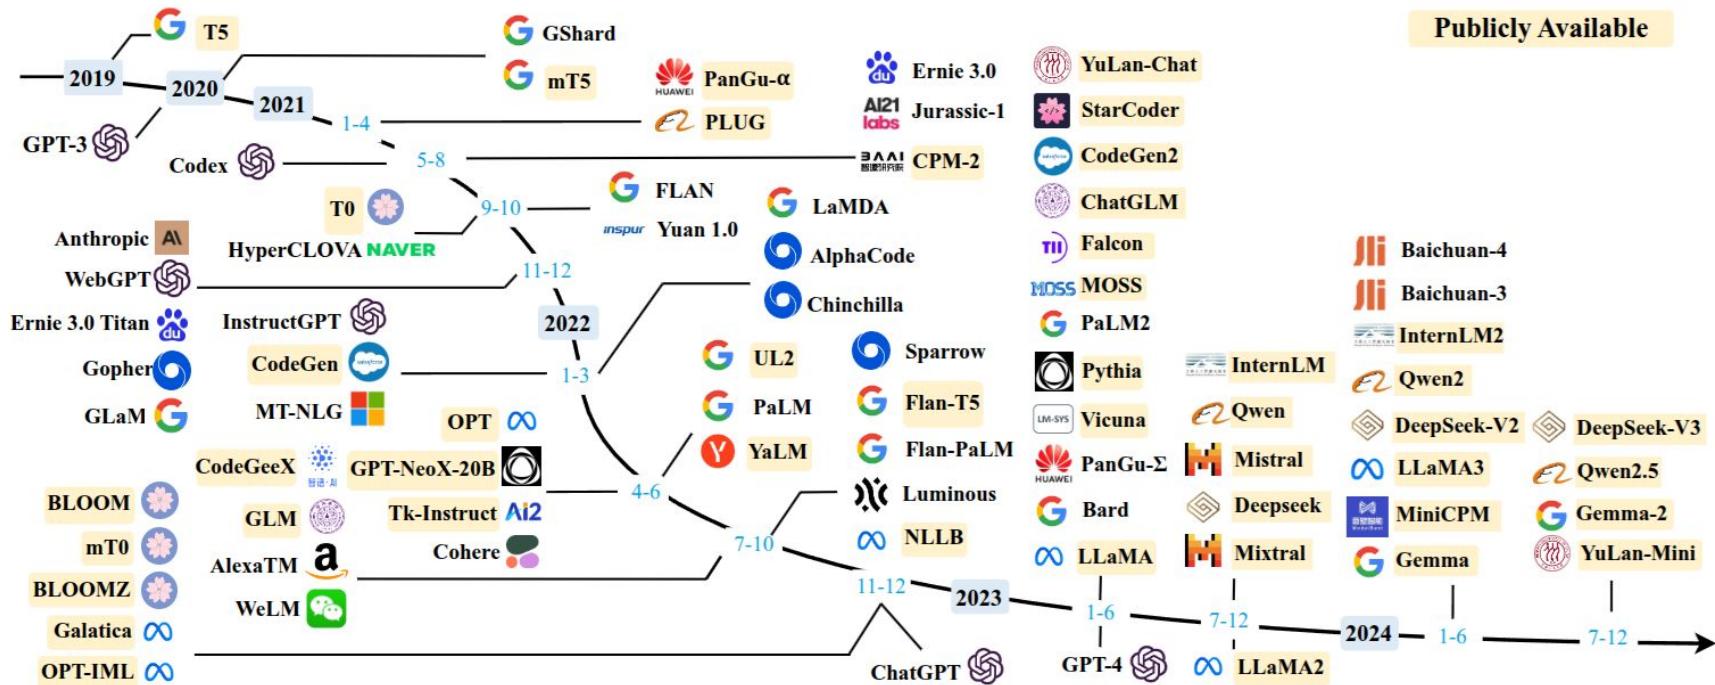
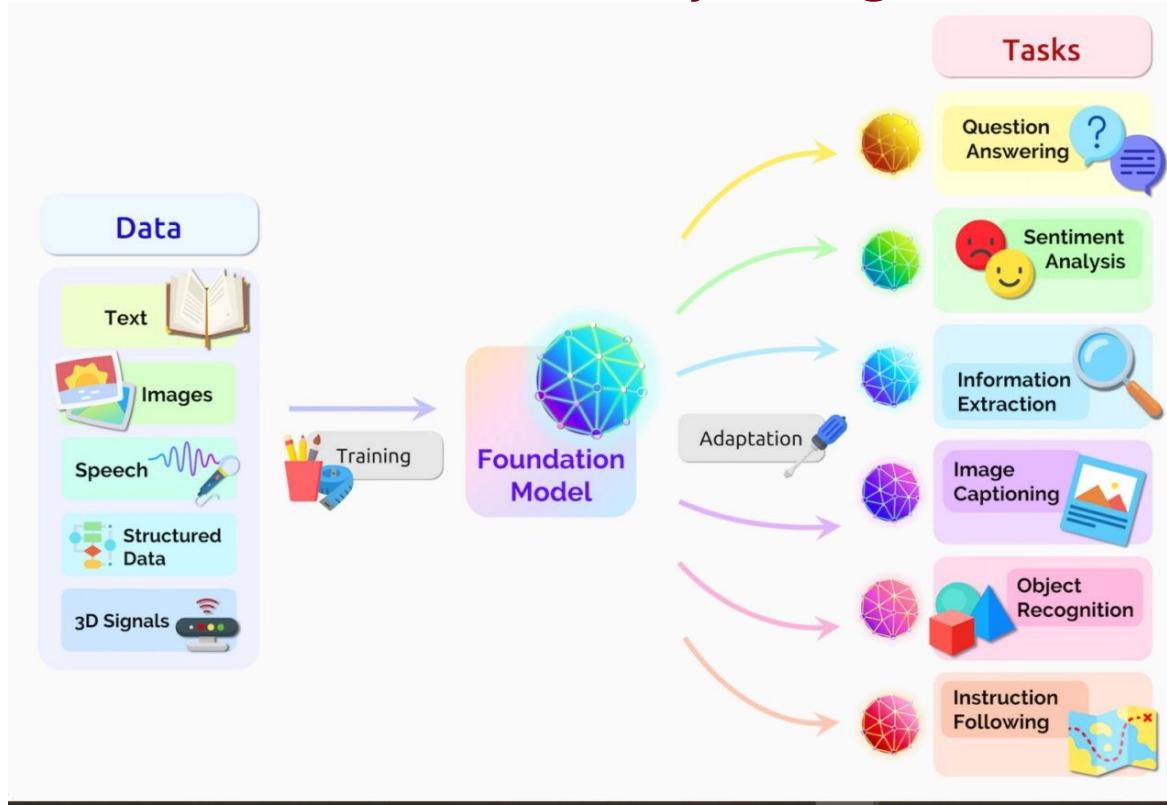


Figure 1: The Transformer - model architecture.

# Transformers reign in NLP!



# Transformers for everything!

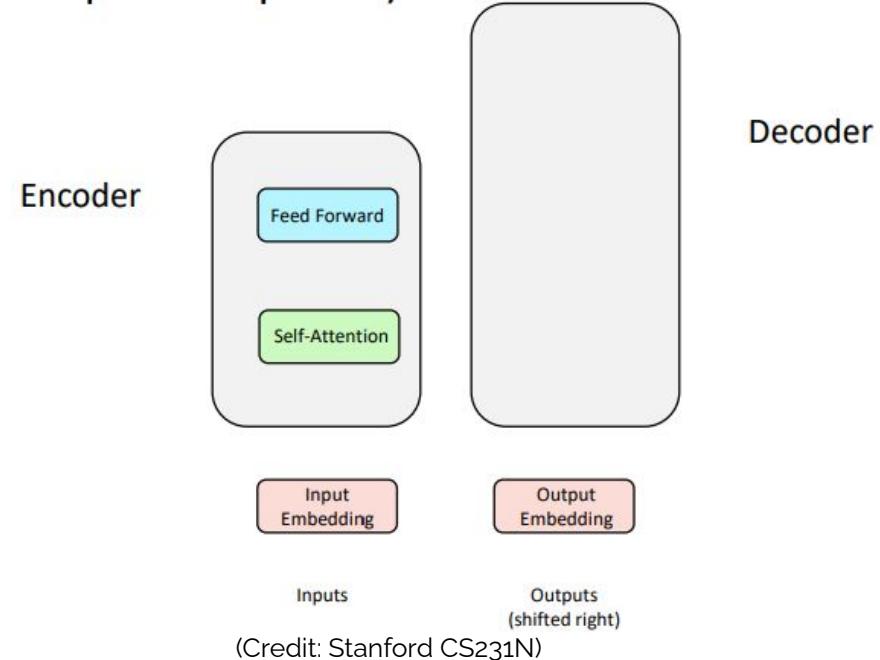
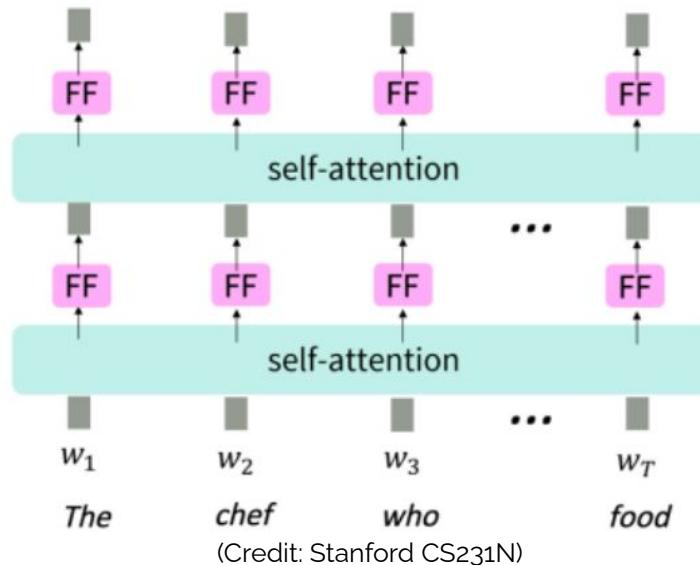


- Transformers have been modified to deal with **almost all** kinds of structured and **unstructured** data
- Enable multimodal data integration and interaction

# Starting from self-attention

Equation for Feed Forward Layer

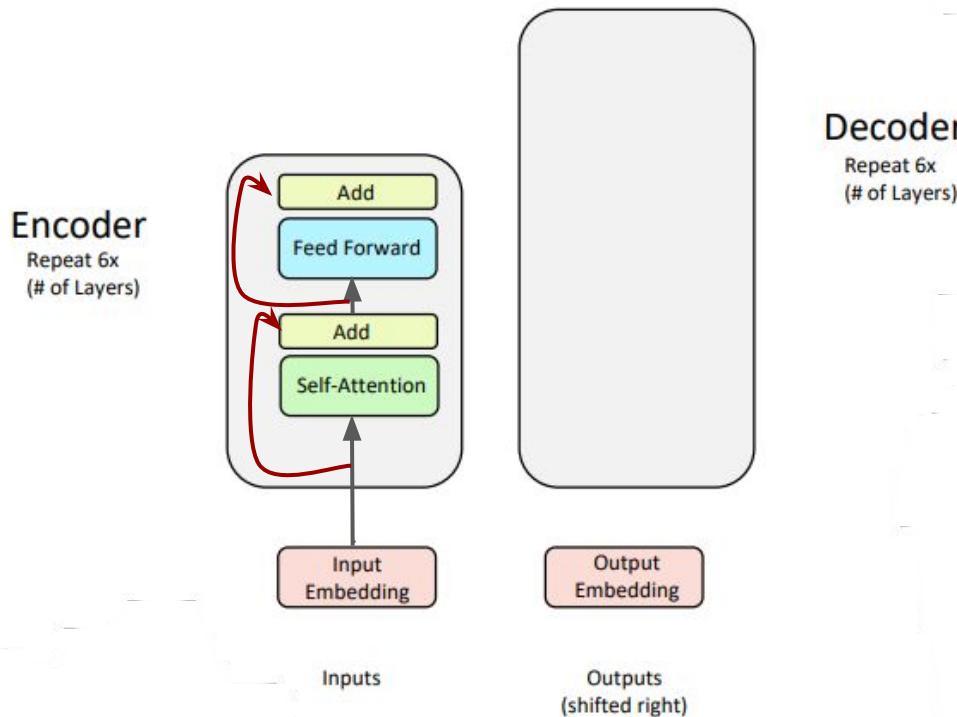
$$\begin{aligned}m_i &= \text{MLP}(\text{output}_i) \\&= W_2 * \text{ReLU}(W_1 \times \text{output}_i + b_1) + b_2\end{aligned}$$



Three tricks to build in depth:

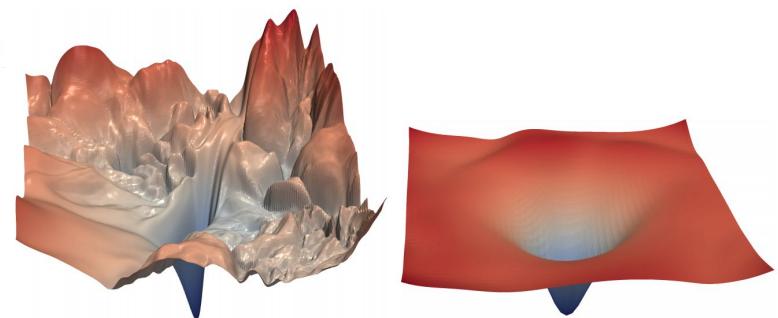
- Residual connection
- Layer normalization
- Scaled inner product attention

# Trick 1: Residual connection



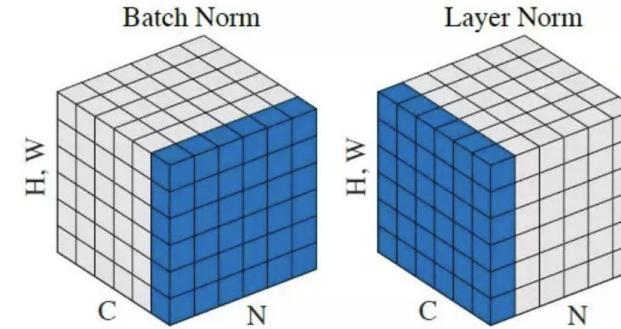
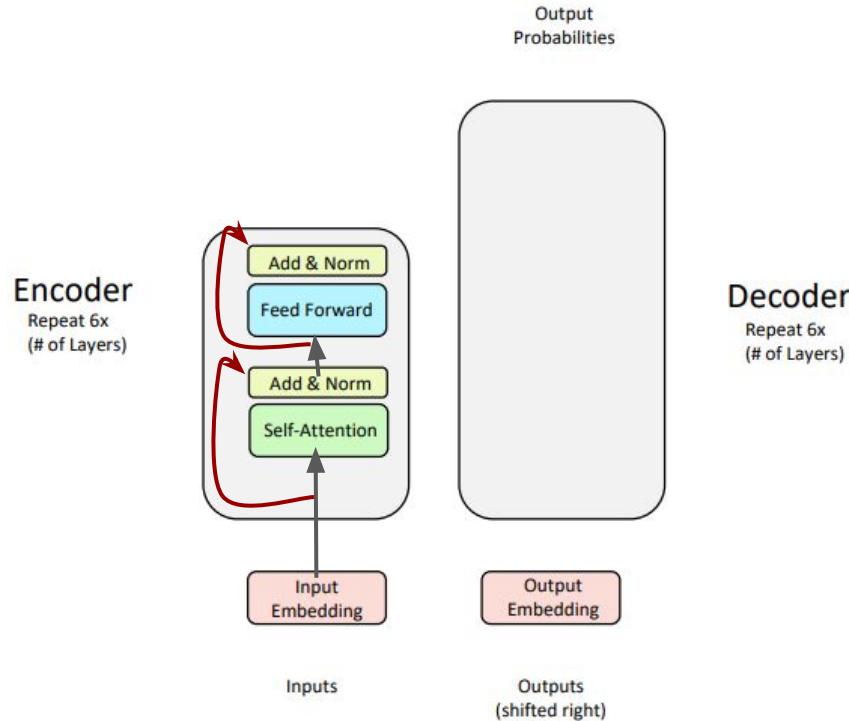
$$\mathbf{x}_k = F(\mathbf{x}_{k-1}) + \mathbf{x}_{k-1}$$

- Mitigating vanishing gradient
- Smoothing out landscape



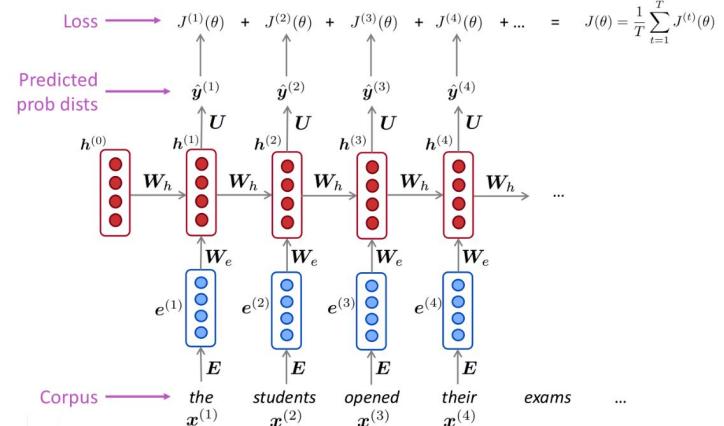
<https://arxiv.org/abs/1712.09913>

# Trick 2: Layer normalization

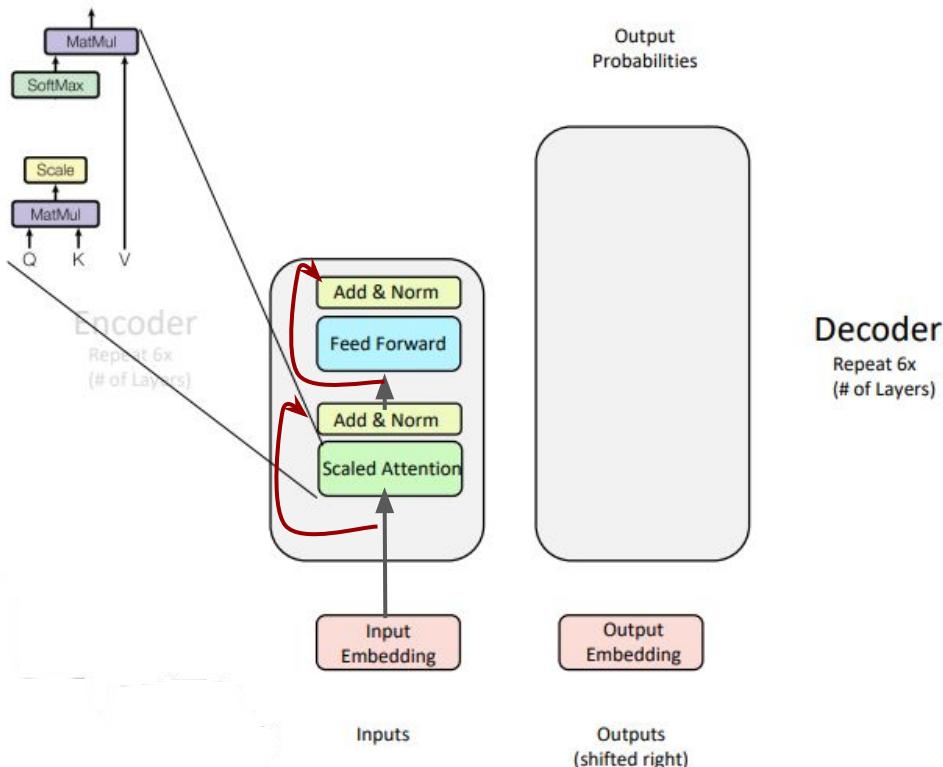


$$x^{\ell'} = \frac{x^{\ell} - \mu^{\ell}}{\sigma^{\ell} + \epsilon}$$

Why not batchnorm?



# Trick 3: Scaled inner product attention



$$\text{output} = \text{softmax}(\mathbf{Q}\mathbf{K}^\top)\mathbf{V}$$

- Suppose that entries of  $\mathbf{Q}$  and  $\mathbf{K}$  behaves like IID zero-mean, unit variance
- $\mathbb{E}\langle \mathbf{q}^i, \mathbf{k}^j \rangle = 0$  but

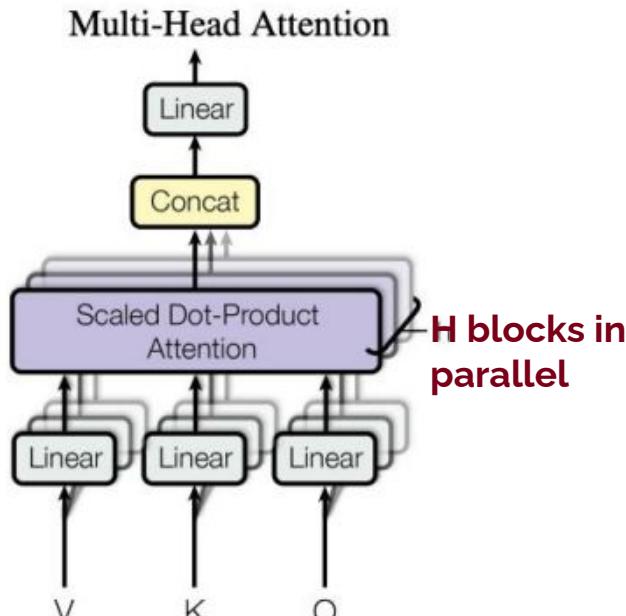
$$\text{Var}\langle \mathbf{q}^i, \mathbf{k}^j \rangle = d_k$$

This can blow up exp computation in the softmax normalization for large  $d_k$ !

Solution: normalize by standard deviation

$$\text{output} = \text{softmax}(\mathbf{Q}\mathbf{K}^\top/\sqrt{d_k})\mathbf{V}$$

# Multi-head attention



[Vaswani et al. 2017]

**Multiple, independent** self-attention blocks in parallel

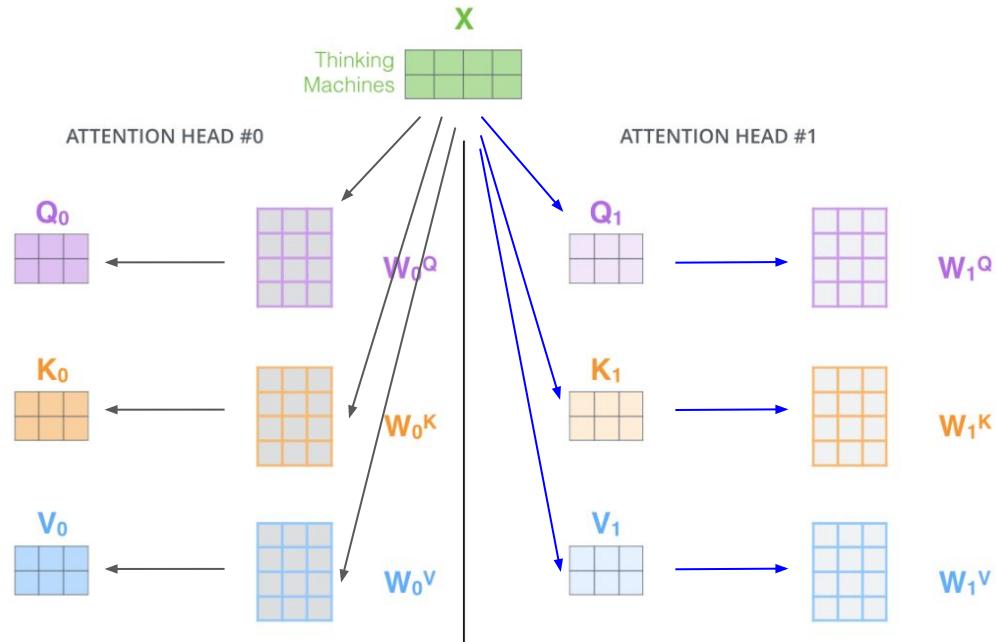


Image credit: <https://jalammar.github.io/illustrated-transformer/>

**Intuition:** allow the flexibility of capturing different kinds of “relevance”/correlations

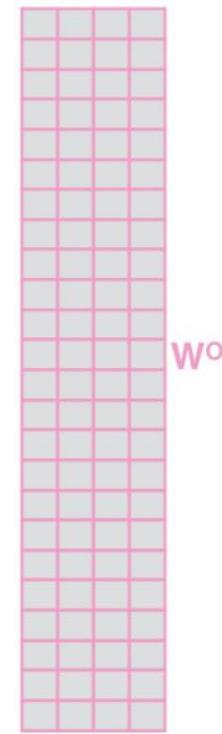
# Multi-head attention

1) Concatenate all the attention heads



2) Multiply with a weight matrix  $W^o$  that was trained jointly with the model

$x$



**Concatenate**



**Multiply**

3) The result would be the  $Z$  matrix that captures information from all the attention heads. We can send this forward to the FFNN



**Output**



Image credit: <https://jalammar.github.io/illustrated-transformer/>

# Multi-head attention

- 1) This is our input sentence\*  
Thinking Machines
- 2) We embed each word\*  
 $X$
- 3) Split into 8 heads.  
We multiply  $X$  or  $R$  with weight matrices
- 4) Calculate attention using the resulting  $Q/K/V$  matrices
- 5) Concatenate the resulting  $Z$  matrices, then multiply with weight matrix  $W^o$  to produce the output of the layer

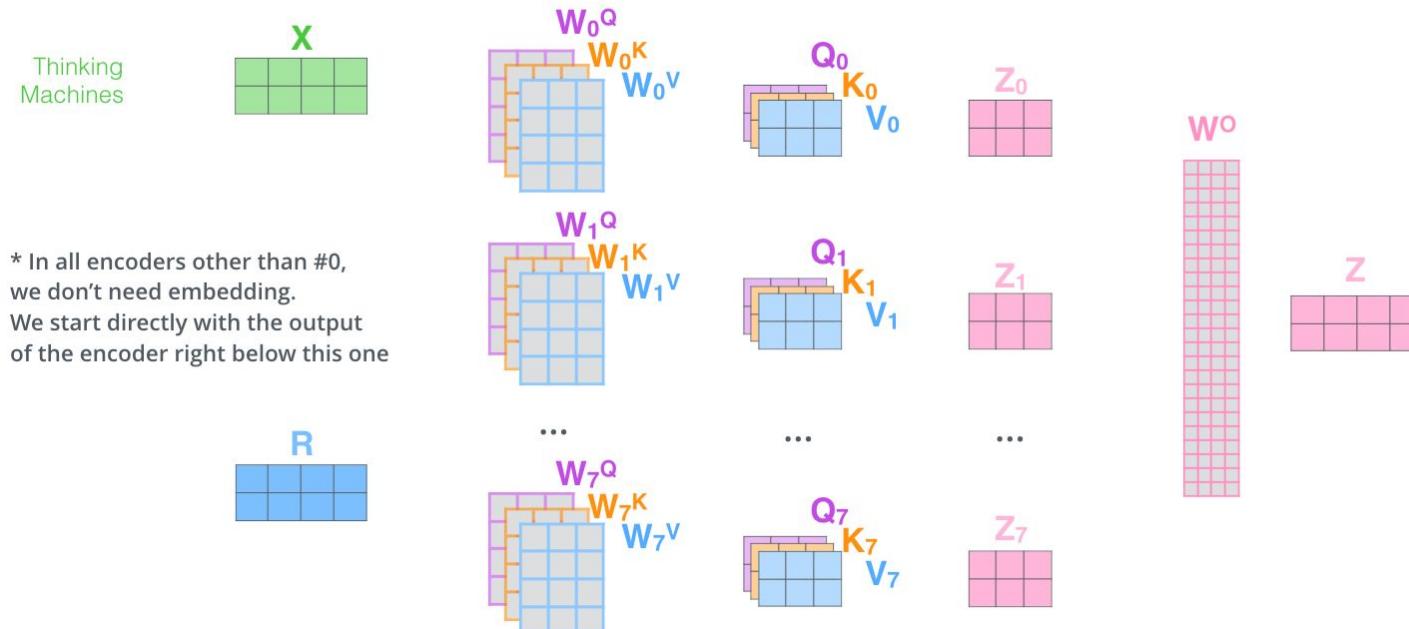
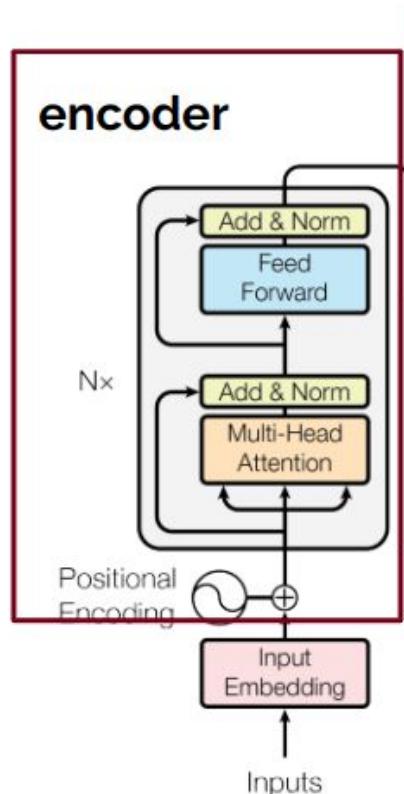


Image credit: <https://jalammar.github.io/illustrated-transformer/>

# Positional encoding



Does the input order matter or not?

$$\mathbf{Q} = \mathbf{X}\mathbf{W}^Q, \quad \mathbf{K} = \mathbf{X}\mathbf{W}^K, \quad \mathbf{V} = \mathbf{X}\mathbf{W}^V$$

$$\text{output} = \text{softmax}(\mathbf{Q}\mathbf{K}^\top / \sqrt{d_k})\mathbf{V}$$

**Positional encoding** to break the **order invariance**

- Idea: a positional vector to (hopefully) encode the position information
  - E.g.,  $\mathbf{X}_p = \mathbf{X} + \mathbf{P}$ , or  $\mathbf{X}_p = [\mathbf{X}, \mathbf{P}]$
- $\mathbf{P}$  can be pre-defined, or made learnable

# Sinusoidal positional encoding

$$\text{PE}(i, \delta) = \begin{cases} \sin\left(\frac{i}{10000^{2\delta'/d}}\right) & \text{if } \delta = 2\delta' \\ \cos\left(\frac{i}{10000^{2\delta'/d}}\right) & \text{if } \delta = 2\delta' + 1 \end{cases}$$

L: sequence length  
d: embedding dimension

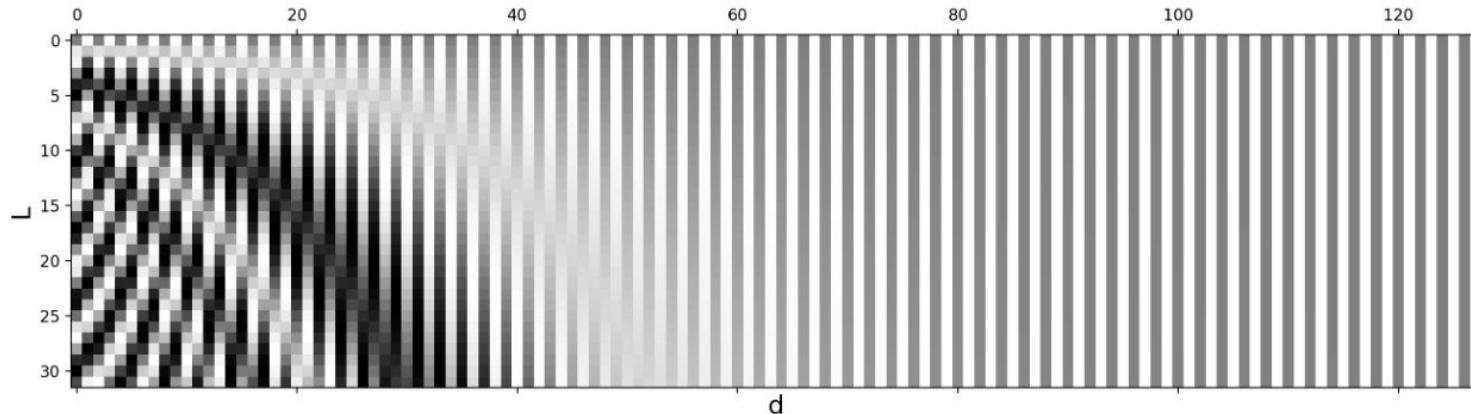
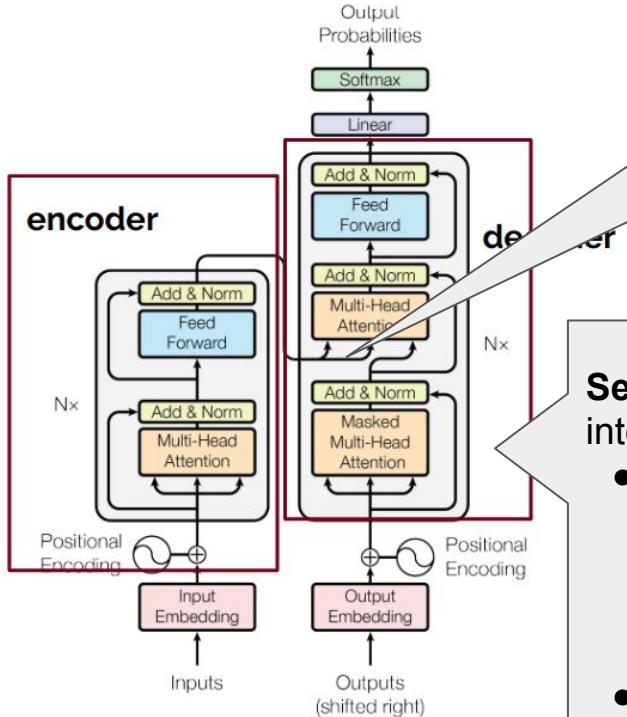


Image credit: <https://lilianweng.github.io/posts/2020-04-07-the-transformer-family/>

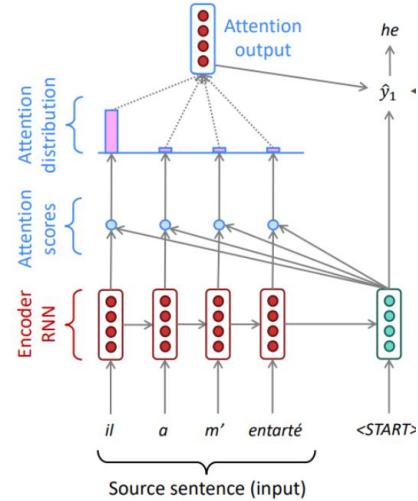
# Decoder



**Cross-attention** (to model the interaction between the encoder key-values and the current decoder query)

**Self-attention** (to model the interaction within itself)

- Respect the sequential nature (e.g., language modeling, assuming access to the future is cheating! )
- Masked out future tokens



(Credit: Stanford CS231N)  
we can look at these (not greyed out) words

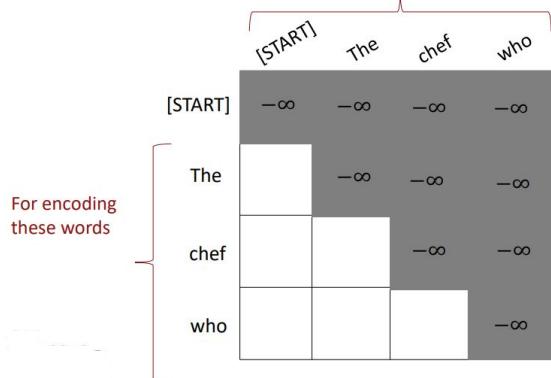


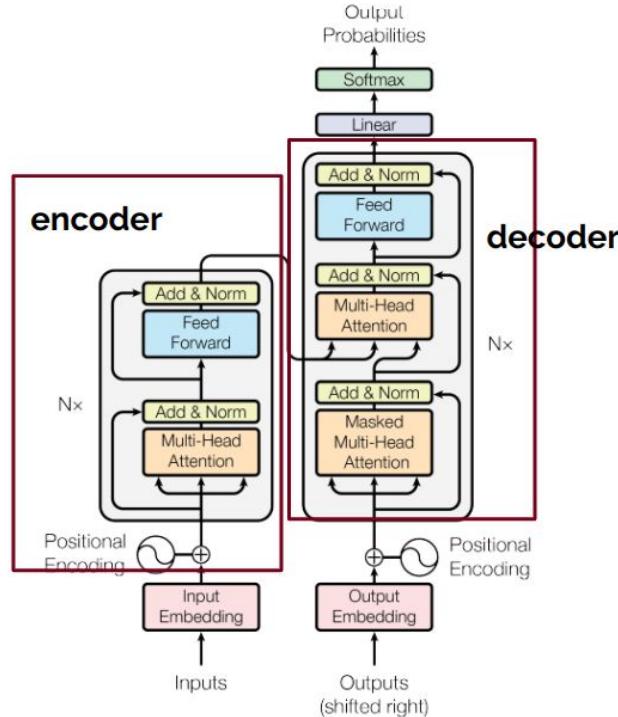
Figure 1: The Transformer - model architecture.

# Strong performance in machine translation

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	<b>41.29</b>	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1		<b><math>3.3 \cdot 10^{18}</math></b>
Transformer (big)	<b>28.4</b>	<b>41.8</b>		$2.3 \cdot 10^{19}$

# Computation



What's the total computation?

$$Q = \mathbf{X}\mathbf{W}^Q, \quad K = \mathbf{X}\mathbf{W}^K, \quad V = \mathbf{X}\mathbf{W}^V$$

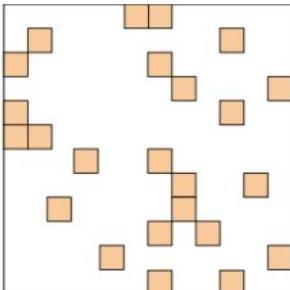
$$\text{output} = \text{softmax}(\mathbf{Q}\mathbf{K}^\top / \sqrt{d_k})\mathbf{V}$$

$$O(T^2d)$$

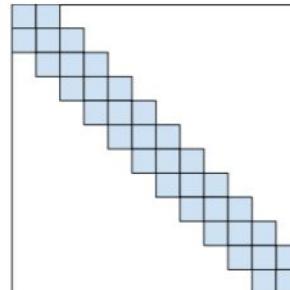
Quadratic computation vs. linear computation in RNNs (**T** is the length of each input sequence, **d** is the embedding dimension)

Figure 1: The Transformer - model architecture.

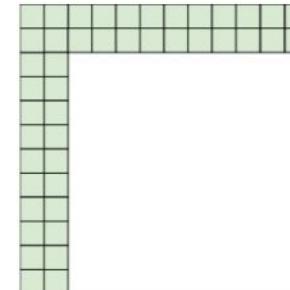
# Computation



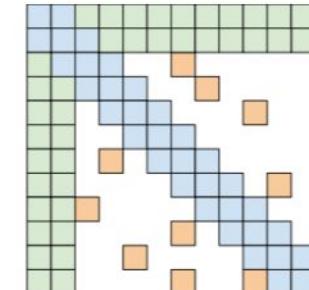
(a) Random attention



(b) Window attention



(c) Global Attention



(d) BIGBIRD

Idea; building in sparsity <https://arxiv.org/abs/2007.14062>

## Do Transformer Modifications Transfer Across Implementations and Applications?

Sharan Narang\* Hyung Won Chung

Yi Tay

William Fedus

Thibault Fevry<sup>†</sup>

Michael Matena<sup>†</sup>

Karishma Malkan<sup>†</sup>

Noah Fiedel

Noam Shazeer

Zhenzhong Lan<sup>†</sup>

Yanqi Zhou

Wei Li

Nan Ding

Jake Marcus

Adam Roberts

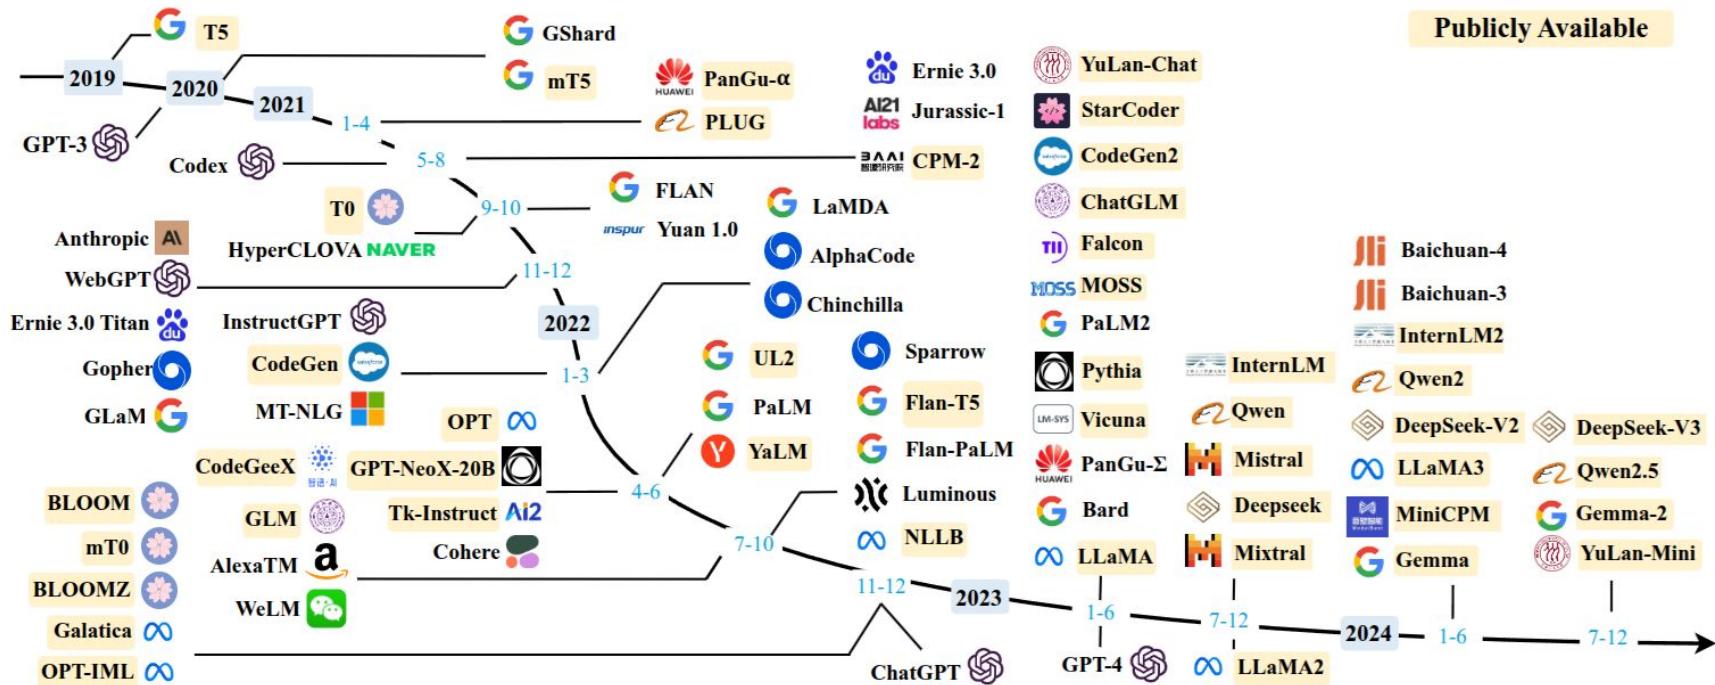
Colin Raffel<sup>†</sup>

But not much consistent improvement so far

<https://arxiv.org/abs/2102.11972>

# Large language models (LLMs)

# Transformers reign in NLP!



# LLMs: large models trained on large datasets

	<b>Model</b>	<b>Release Time</b>	<b>Size (B)</b>	
Publicly Available	T5 [73]	Oct-2019	11	GPT-3 [55]
	mT5 [74]	Oct-2020	13	GShard [91]
	PanGu- $\alpha$ [75]	Apr-2021	13*	Codex [92]
	CPM-2 [76]	Jun-2021	198	ERNIE 3.0 [93]
	T0 [28]	Oct-2021	11	Jurassic-1 [94]
	CodeGen [77]	Mar-2022	16	HyperCLOVA [95]
	GPT-NeoX-20B [78]	Apr-2022	20	FLAN [62]
	Tk-Instruct [79]	Apr-2022	11	Yuan 1.0 [96]
	UL2 [80]	May-2022	20	Anthropic [97]
	OPT [81]	May-2022	175	WebGPT [72]
	NLLB [82]	Jul-2022	54.5	Gopher [59]
	CodeGeeX [83]	Sep-2022	13	ERNIE 3.0 Titan [98]
	GLM [84]	Oct-2022	130	GLaM [99]
	Flan-T5 [64]	Oct-2022	11	LaMDA [63]
	BLOOM [69]	Nov-2022	176	MT-NLG [100]
	mT0 [85]	Nov-2022	13	AlphaCode [101]
	Galactica [35]	Nov-2022	120	InstructGPT [61]
	BLOOMZ [85]	Nov-2022	176	Chinchilla [34]
	OPT-IML [86]	Dec-2022	175	PaLM [56]
	LLaMA [57]	Feb-2023	65	AlexaTM [102]
	Pythia [87]	Apr-2023	12	Sparrow [103]
	CodeGen2 [88]	May-2023	16	WeLM [104]
	StarCoder [89]	May-2023	15.5	U-PaLM [105]
	LLaMA2 [90]	Jul-2023	70	Flan-PaLM [64]
				Flan-U-PaLM [64]
				GPT-4 [46]
				PanGu- $\Sigma$ [106]
				PaLM2 [107]

Image credit: A Survey of Large Language Models <https://arxiv.org/abs/2303.18223>

# LLMs: large models trained on large datasets

TABLE 2: Statistics of commonly-used data sources.

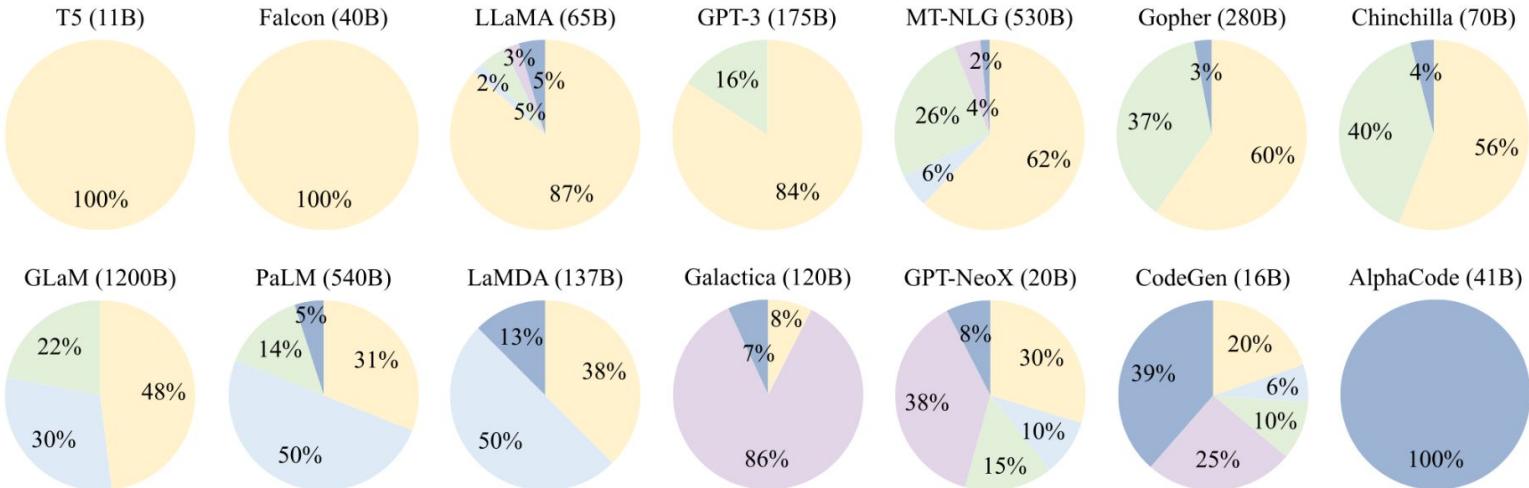
Corpora	Size	Source	Latest Update Time
BookCorpus [138]	5GB	Books	Dec-2015
Gutenberg [139]	-	Books	Dec-2021
C4 [73]	800GB	CommonCrawl	Apr-2019
CC-Stories-R [140]	31GB	CommonCrawl	Sep-2019
CC-NEWS [27]	78GB	CommonCrawl	Feb-2019
REALNEWS [141]	120GB	CommonCrawl	Apr-2019
OpenWebText [142]	38GB	Reddit links	Mar-2023
Pushift.io [143]	2TB	Reddit links	Mar-2023
Wikipedia [144]	21GB	Wikipedia	Mar-2023
BigQuery [145]	-	Codes	Mar-2023
the Pile [146]	800GB	Other	Dec-2020
ROOTS [147]	1.6TB	Other	Jun-2022

# Two crucial technical steps toward LLMs

- **Pretraining**
- **Finetuning (Adaptation)**

Recall transfer learning?

# Pretraining: data collection



- Webpages
- Conversation Data
- Books & News
- Scientific Data
- Code
- C4 (800G, 2019), ■ OpenWebText (38G, 2023), ■ Wikipedia (21G, 2023)
- ❑ the Pile - StackExchange (41G, 2020)
- ❑ BookCorpus (5G, 2015), ❑ Gutenberg (-, 2021), ❑ CC-Stories-R (31G, 2019), ❑ CC-NEWES (78G, 2019), ❑ REALNEWS (120G, 2019)
- ❑ the Pile - ArXiv (72G, 2020), ❑ the Pile - PubMed Abstracts (25G, 2020)
- ❑ BigQuery (-, 2023), the Pile - GitHub (61G, 2020)

# Pretraining: data collection

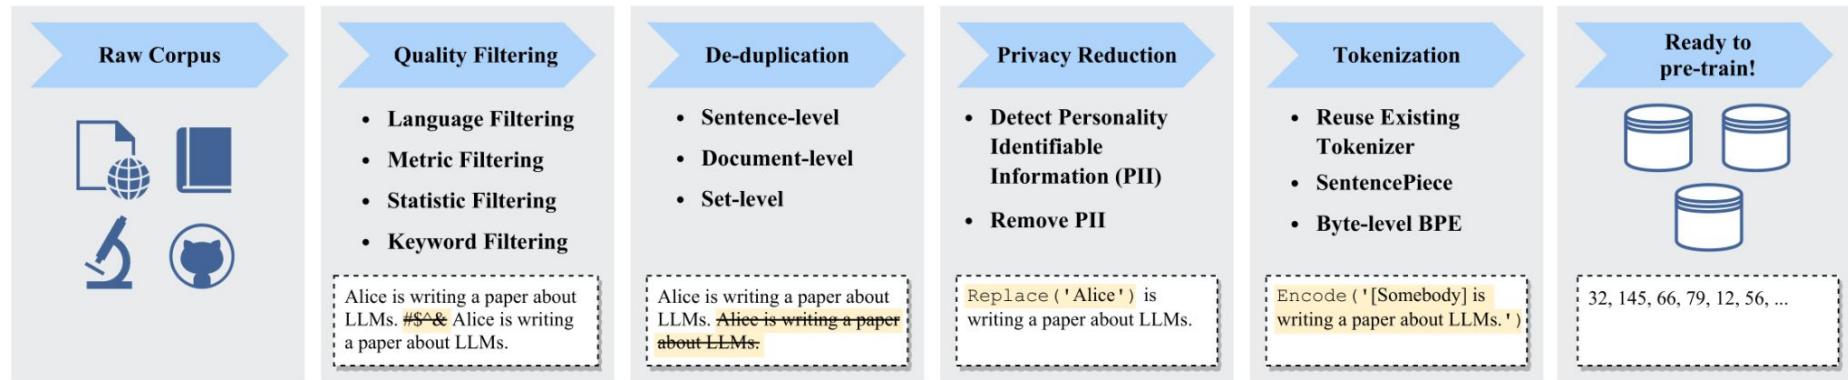
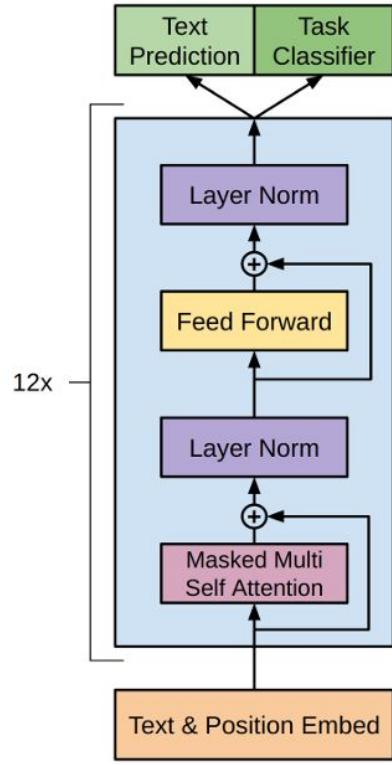
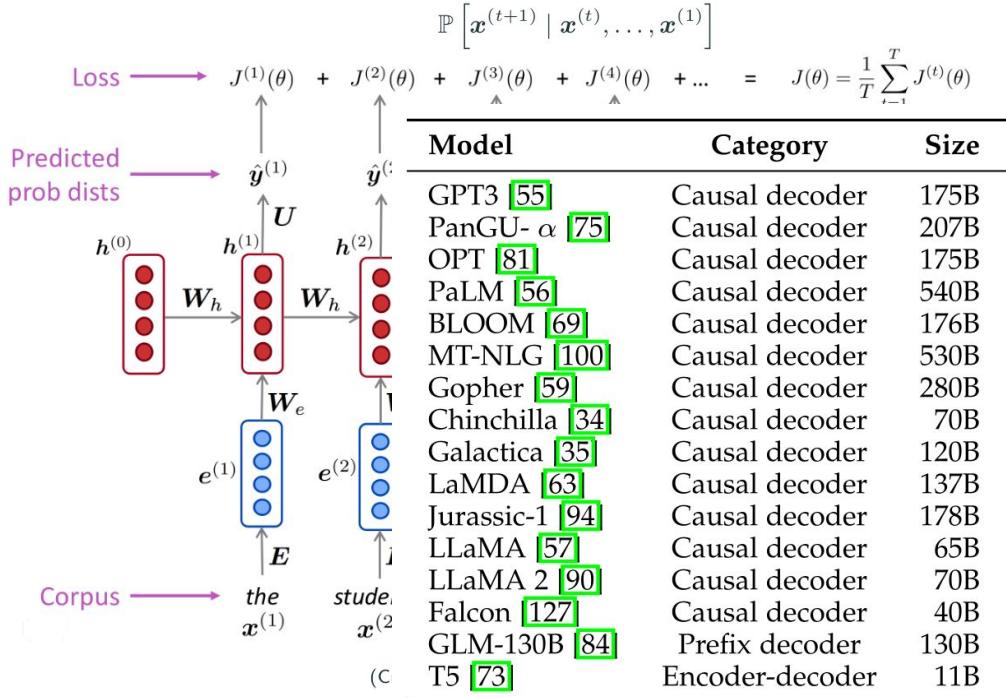


Fig. 6: An illustration of a typical data preprocessing pipeline for pre-training large language models.

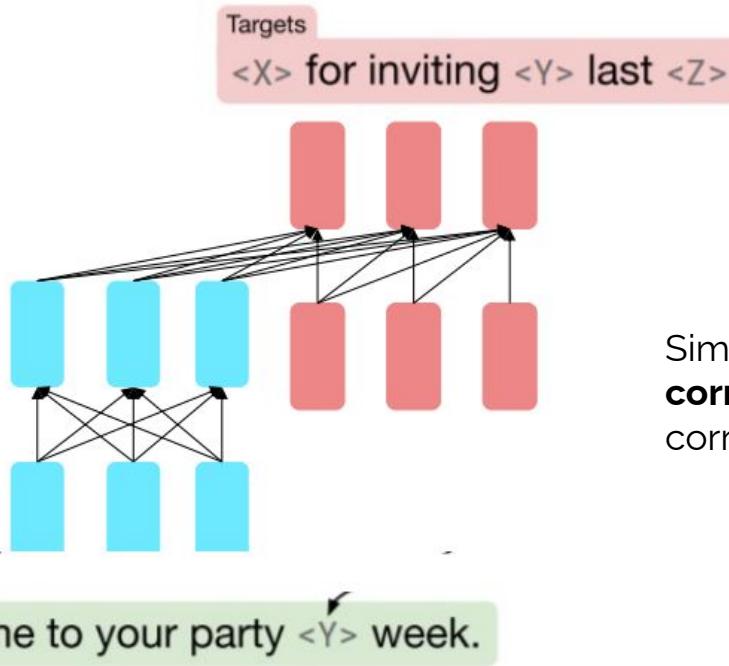
# Pretraining: architecture & task



Most popular: (transformer-based) **decoder-only** architectures pretrained on **language modeling**, i.e. model



# Pretraining: architecture & task – alternative



Similar to pretraining encoder,  
**corruption removal!** (called span  
corruption)

# Pretraining: architecture details

Configuration	Method	Equation
Normalization position	Post Norm [22]	$\text{Norm}(\mathbf{x} + \text{Sublayer}(\mathbf{x}))$
	Pre Norm [26]	$\mathbf{x} + \text{Sublayer}(\text{Norm}(\mathbf{x}))$
	Sandwich Norm [201]	$\mathbf{x} + \text{Norm}(\text{Sublayer}(\text{Norm}(\mathbf{x})))$
Normalization method	LayerNorm [202]	$\frac{\mathbf{x} - \mu}{\sqrt{\sigma}} \cdot \gamma + \beta, \quad \mu = \frac{1}{d} \sum_{i=1}^d x_i, \quad \sigma = \sqrt{\frac{1}{d} \sum_{i=1}^d (x_i - \mu)^2}$
	RMSNorm [203]	$\frac{\mathbf{x}}{\text{RMS}(\mathbf{x})} \cdot \gamma, \quad \text{RMS}(\mathbf{x}) = \sqrt{\frac{1}{d} \sum_{i=1}^d x_i^2}$
	DeepNorm [204]	$\text{LayerNorm}(\alpha \cdot \mathbf{x} + \text{Sublayer}(\mathbf{x}))$
Activation function	ReLU [205]	$\text{ReLU}(\mathbf{x}) = \max(\mathbf{x}, \mathbf{0})$
	GeLU [206]	$\text{GeLU}(\mathbf{x}) = 0.5\mathbf{x} \otimes [1 + \text{erf}(\mathbf{x}/\sqrt{2})], \quad \text{erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$
	Swish [207]	$\text{Swish}(\mathbf{x}) = \mathbf{x} \otimes \text{sigmoid}(\mathbf{x})$
	SwiGLU [208]	$\text{SwiGLU}(\mathbf{x}_1, \mathbf{x}_2) = \text{Swish}(\mathbf{x}_1) \otimes \mathbf{x}_2$
	GeGLU [208]	$\text{GeGLU}(\mathbf{x}_1, \mathbf{x}_2) = \text{GeLU}(\mathbf{x}_1) \otimes \mathbf{x}_2$
Position embedding	Absolute [22]	$\mathbf{x}_i = \mathbf{x}_i + \mathbf{p}_i$
	Relative [73]	$A_{ij} = \mathbf{W}_q \mathbf{x}_i \mathbf{x}_j^T \mathbf{W}_k^T + r_{i-j}$
	RoPE [209]	$A_{ij} = \mathbf{W}_q \mathbf{x}_i \mathbf{R}_{\theta, i-j} \mathbf{x}_j^T \mathbf{W}_k^T$
	Alibi [210]	$A_{ij} = \mathbf{W}_q \mathbf{x}_i \mathbf{R}_{\theta, i-j} \mathbf{x}_j^T \mathbf{W}_k^T \quad A_{ij} = \mathbf{W}_q \mathbf{x}_i \mathbf{x}_j^T \mathbf{W}_k^T - m(i - j)$

# Pretraining: optimization details

TABLE 5: Detailed optimization settings of several existing LLMs.

Model	Batch Size (#tokens)	Learning Rate	Warmup	Decay Method	Optimizer	Precision Type	Weight Decay	Grad Clip	Dropout
GPT3 (175B)	32K→3.2M	$6 \times 10^{-5}$	yes	cosine decay to 10%	Adam	FP16	0.1	1.0	-
PanGu- $\alpha$ (200B)	-	$2 \times 10^{-5}$	-	-	Adam	-	0.1	-	-
OPT (175B)	2M	$1.2 \times 10^{-4}$	yes	manual decay	AdamW	FP16	0.1	-	0.1
PaLM (540B)	1M→4M	$1 \times 10^{-2}$	no	inverse square root	Adafactor	BF16	$lr^2$	1.0	0.1
BLOOM (176B)	4M	$6 \times 10^{-5}$	yes	cosine decay to 10%	Adam	BF16	0.1	1.0	0.0
MT-NLG (530B)	64 K→3.75M	$5 \times 10^{-5}$	yes	cosine decay to 10%	Adam	BF16	0.1	1.0	-
Gopher (280B)	3M→6M	$4 \times 10^{-5}$	yes	cosine decay to 10%	Adam	BF16	-	1.0	-
Chinchilla (70B)	1.5M→3M	$1 \times 10^{-4}$	yes	cosine decay to 10%	AdamW	BF16	-	-	-
Galactica (120B)	2M	$7 \times 10^{-6}$	yes	linear decay to 10%	AdamW	-	0.1	1.0	0.1
LaMDA (137B)	256K	-	-	-	-	BF16	-	-	-
Jurassic-1 (178B)	32 K→3.2M	$6 \times 10^{-5}$	yes	-	-	-	-	-	-
LLaMA (65B)	4M	$1.5 \times 10^{-4}$	yes	cosine decay to 10%	AdamW	-	0.1	1.0	-
LLaMA 2 (70B)	4M	$1.5 \times 10^{-4}$	yes	cosine decay to 10%	AdamW	-	0.1	1.0	-
Falcon (40B)	2M	$1.85 \times 10^{-4}$	yes	cosine decay to 10%	AdamW	BF16	0.1	-	-
GLM (130B)	0.4M→8.25M	$8 \times 10^{-5}$	yes	cosine decay to 10%	AdamW	FP16	0.1	1.0	0.1
T5 (11B)	64K	$1 \times 10^{-2}$	no	inverse square root	AdaFactor	-	-	-	0.1
ERNIE 3.0 Titan (260B)	-	$1 \times 10^{-4}$	-	-	Adam	FP16	0.1	1.0	-
PanGu- $\Sigma$ (1.085T)	0.5M	$2 \times 10^{-5}$	yes	-	Adam	FP16	-	-	-

# Supervised adaptation—instruction tuning

TABLE 6: A detailed list of available collections for instruction tuning.

Categories	Collections	Time	#Examples
Task	Nat. Inst. [264]	Apr-2021	193K
	FLAN [62]	Sep-2021	4.4M
	P3 [265]	Oct-2021	12.1M
	Super Nat. Inst. [79]	Apr-2022	5M
	MVPCorpus [266]	Jun-2022	41M
	xP3 [85]	Nov-2022	81M
	OIC <sup>22</sup>	Mar-2023	43M
Chat	HH-RLHF [267]	Apr-2022	160K
	HC3 [268]	Jan-2023	87K
	ShareGPT <sup>23</sup>	Mar-2023	90K
	Dolly <sup>24</sup>	Apr-2023	15K
	OpenAssistant [269]	Apr-2023	161K
Synthetic	Self-Instruct [129]	Dec-2022	82K
	Alpaca [123]	Mar-2023	52K
	Guanaco <sup>25</sup>	Mar-2023	535K
	Baize [270]	Apr-2023	158K
	BELLE [271]	Apr-2023	1.5M

Image credit: A Survey of Large Language Models  
<https://arxiv.org/abs/2303.18223>

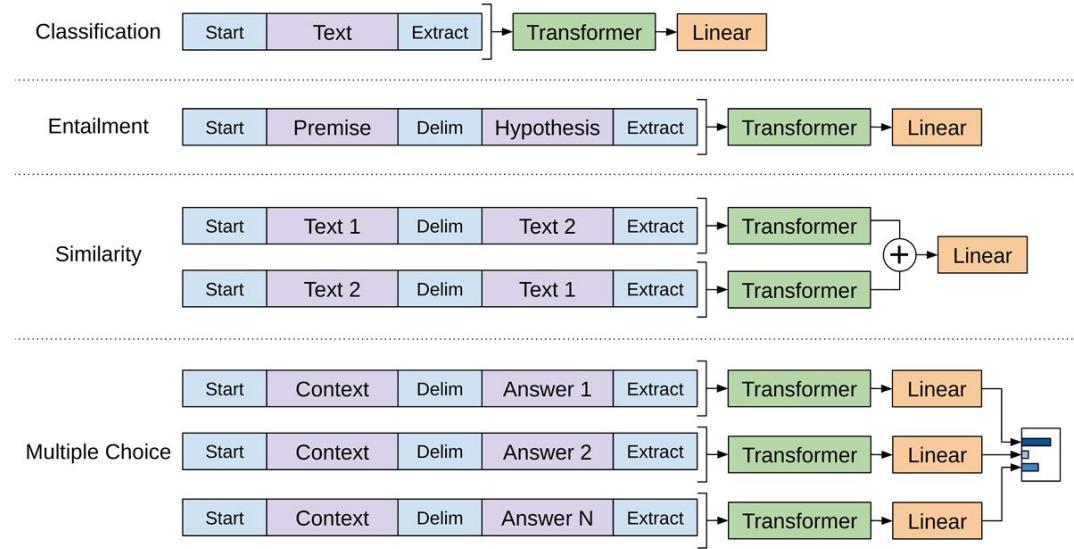


Image credit: Improving Language Understanding by Generative Pre-Training  
<https://gwern.net/doc/www/s3-us-west-2.amazonaws.com/d73fdc5ffa8627bce4.pdf>

# Constructing the instruction sets

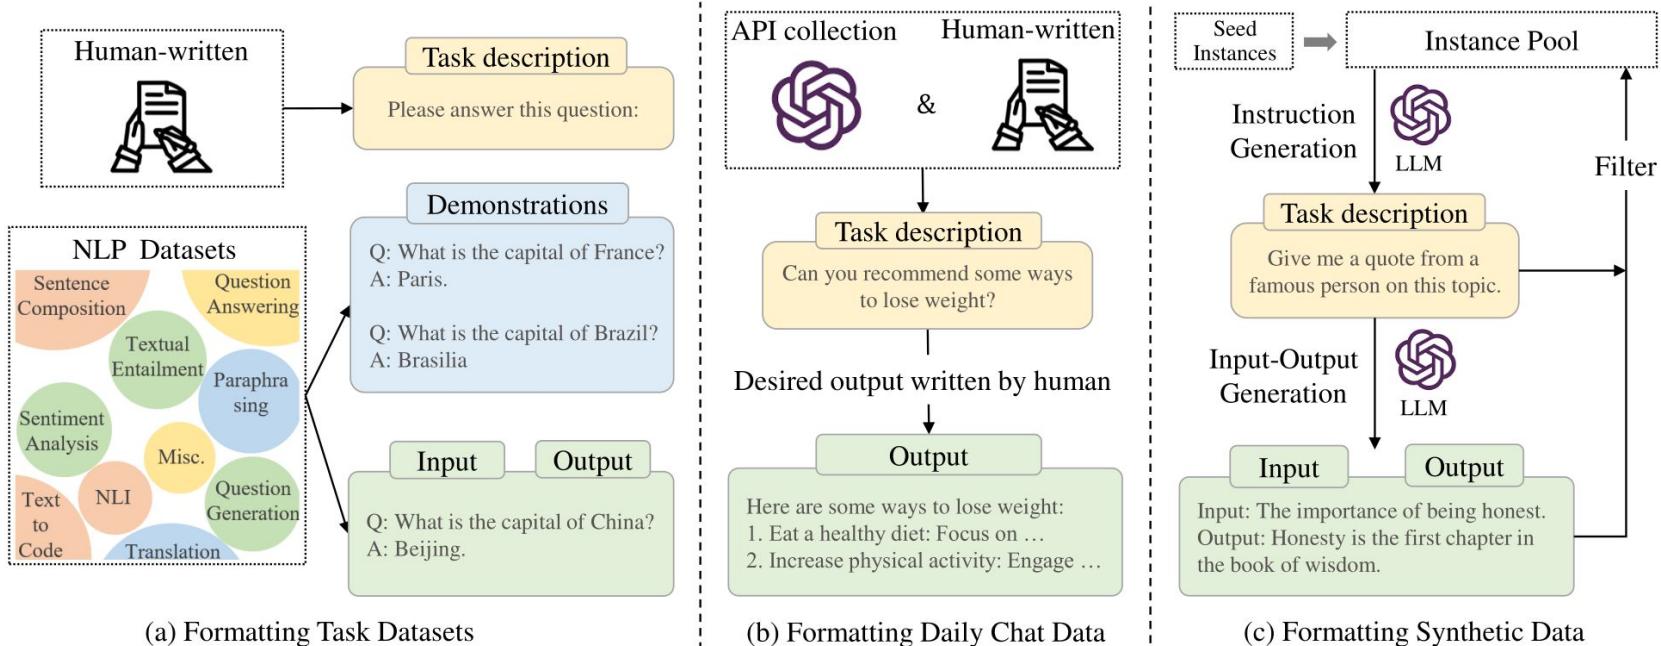
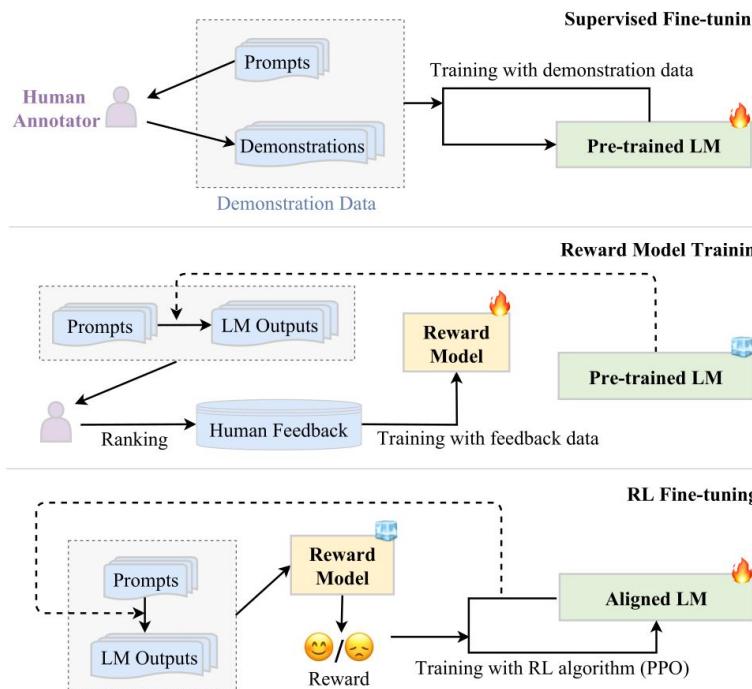


Fig. 9: An illustration of instance formatting and three different methods for constructing the instruction-formatted instances.

# Supervised adaptation—alignment tuning



Make sure the output is aligned with human values and not harmful

## Reinforcement learning with human feedback (RLHF)

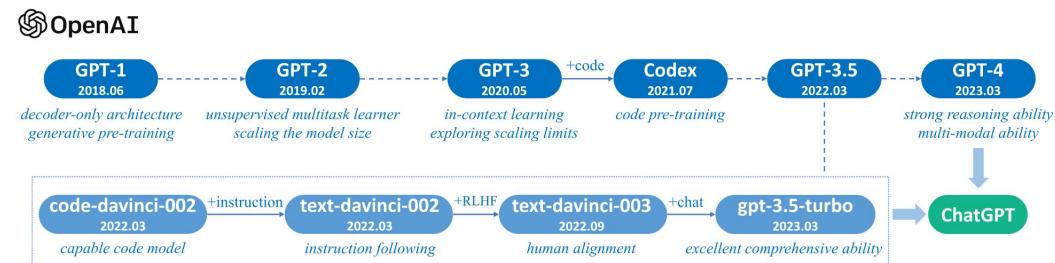


Fig. 3: A brief illustration for the technical evolution of GPT-series models. We plot this figure mainly based on the papers, blog articles and official APIs from OpenAI. Here, solid lines denote that there exists an explicit evidence (e.g., the official statement that a new model is developed based on a base model) on the evolution path between two models, while dashed lines denote a relatively weaker evolution relation.

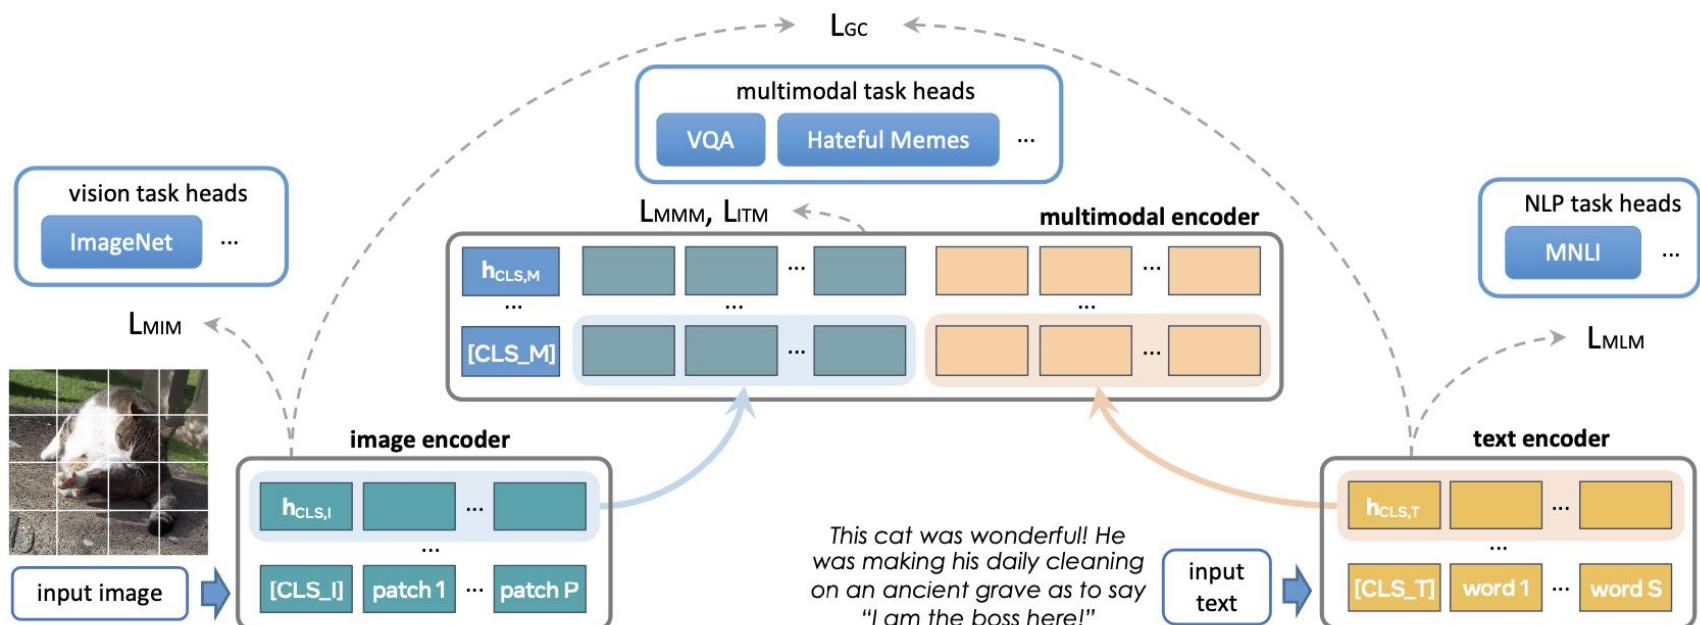
# Transformers for other domains

# Vision Transformers

*Transformers | Davide Cacomin | 2021*

[https://en.wikipedia.org/wiki/Vision\\_transformer](https://en.wikipedia.org/wiki/Vision_transformer)

# Multimodal foundation models



<https://pytorch.org/blog/scaling-multimodal-foundation-models-in-torchmultimodal-with-pytorch-distributed/>